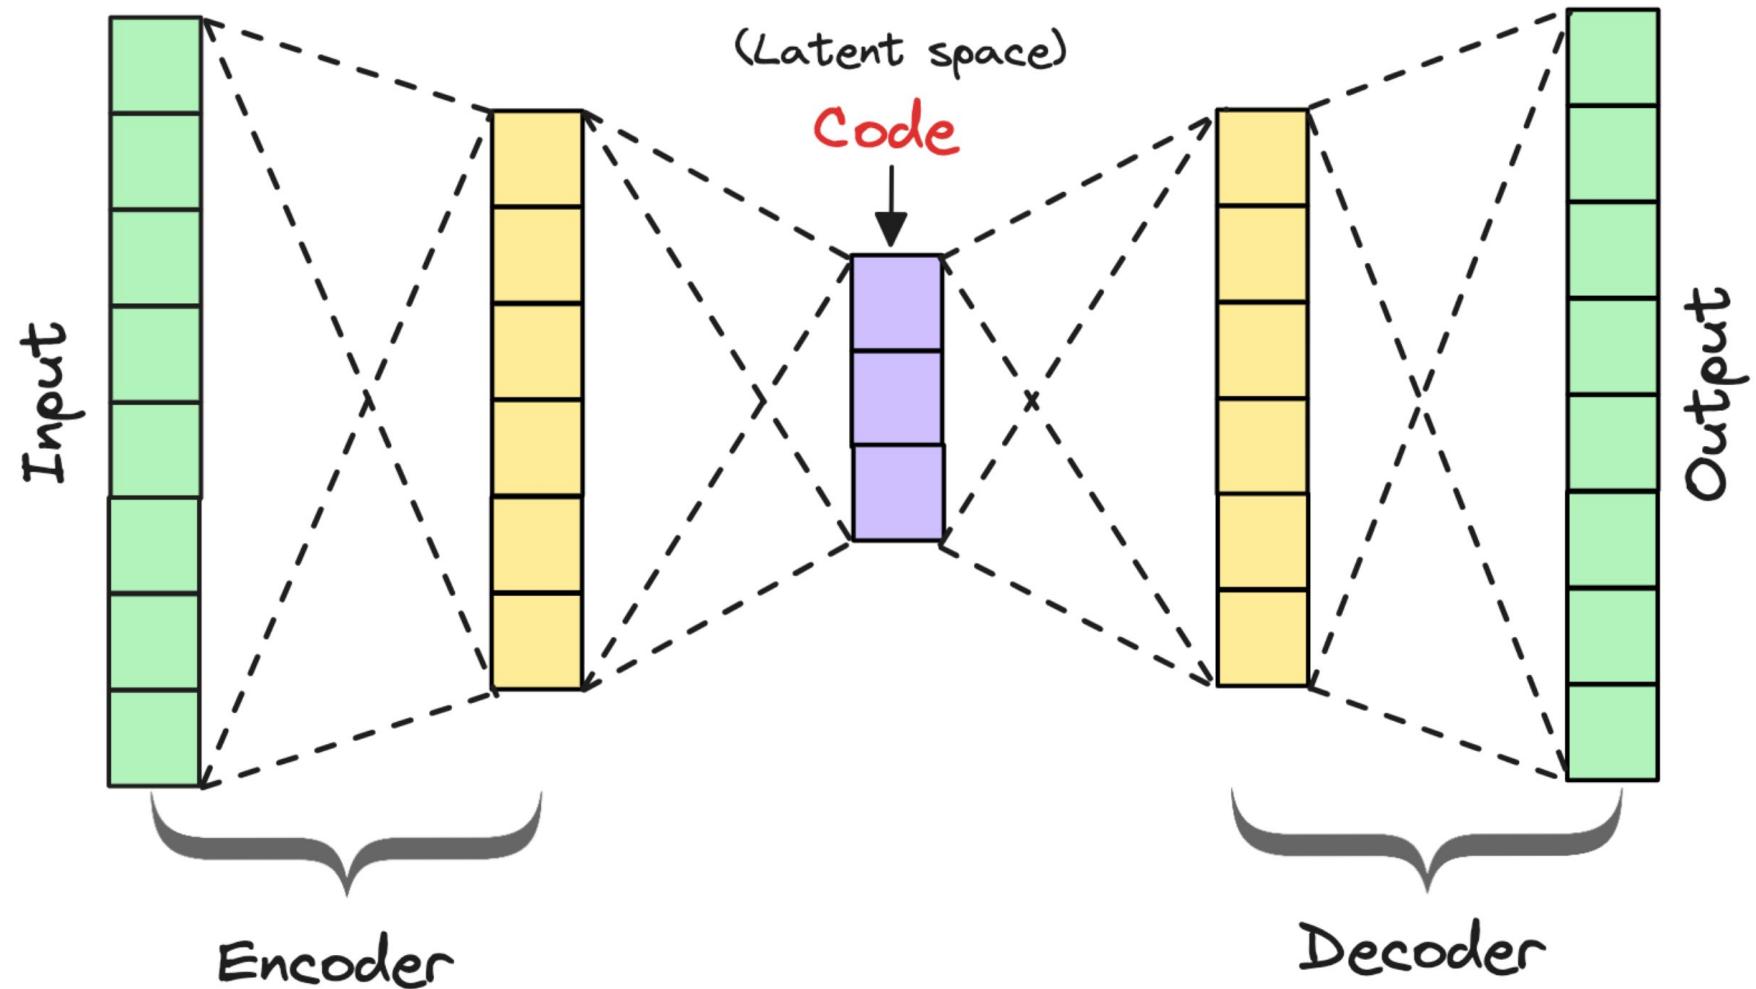


Understanding Autoencoders! 🚀



Implementation from scratch using PyTorch Lightning! ⚡

Autoencoders have two main parts:

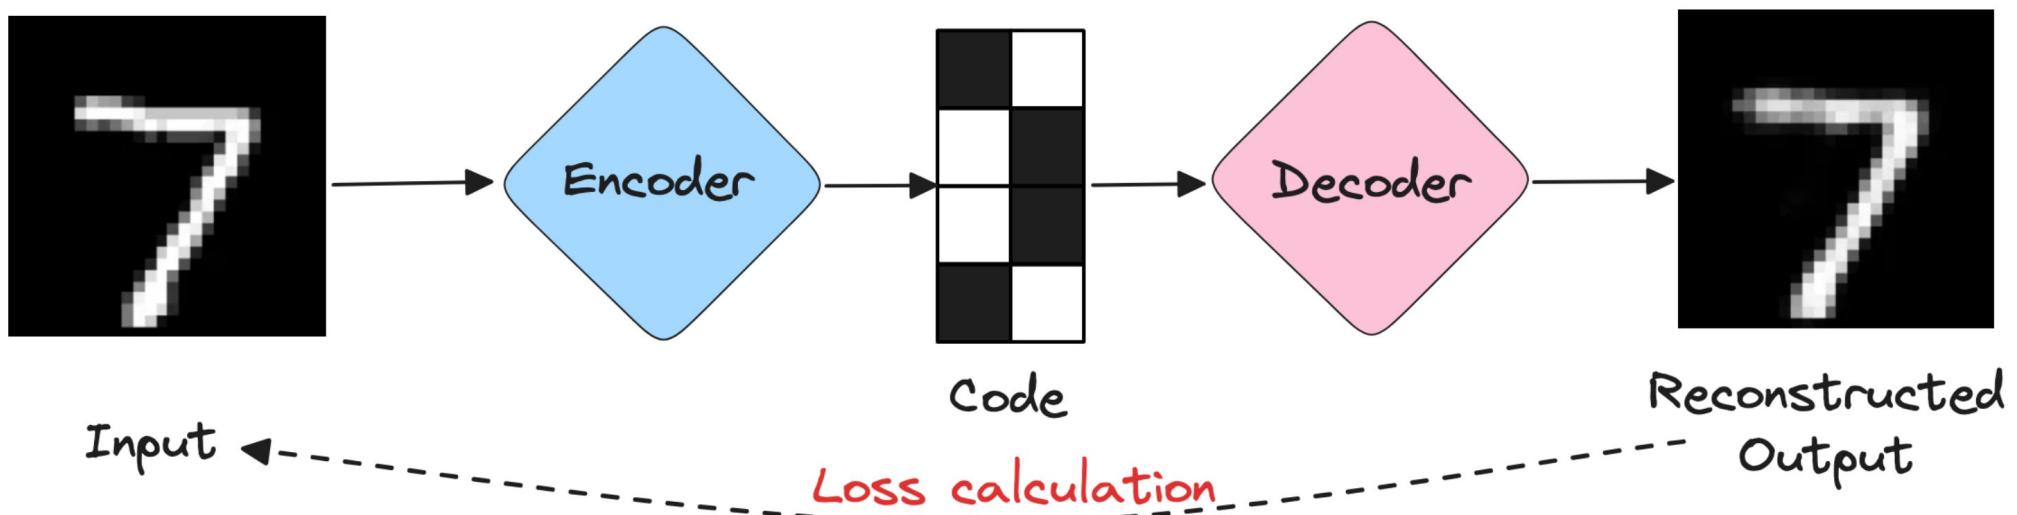
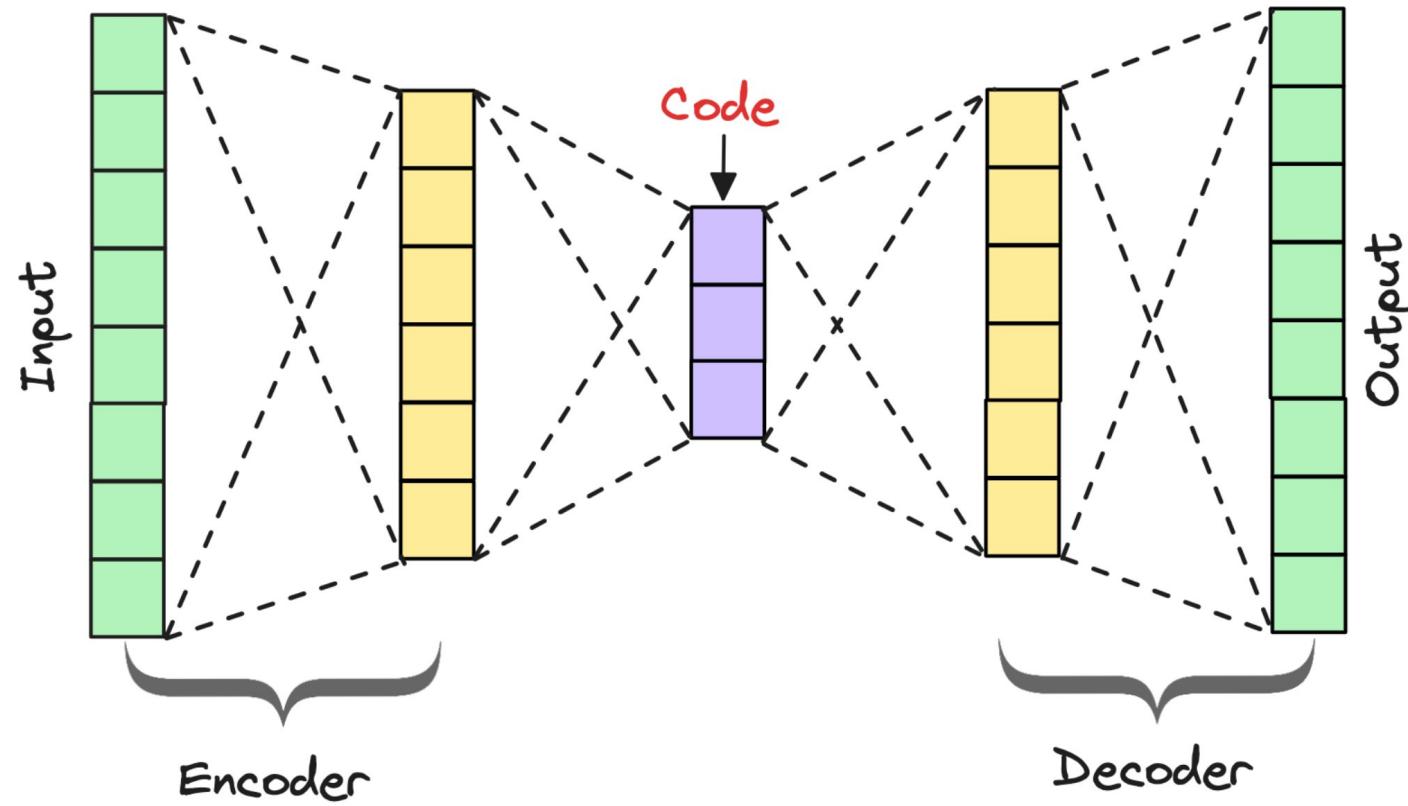
- 1 Encoder: Compresses the input into a dense representation (latent space)
- 2 Decoder: Reconstructs the input from this dense representation.

The idea is to make the reconstructed output as close to the original input as possible

Swipe ➡



@akshay_pachaar



Loss is the difference between input and reconstructed output.

We train them using backpropagation, adjusting weights based on the reconstruction loss!



@akshay_pachaar

Applications of Autoencoders:

- Dimensionality Reduction: Like PCA but cooler. 😎
- Anomaly Detection: If reconstruction error is high, something's fishy!
- Data Denoising: Clean noisy data by training on noise.

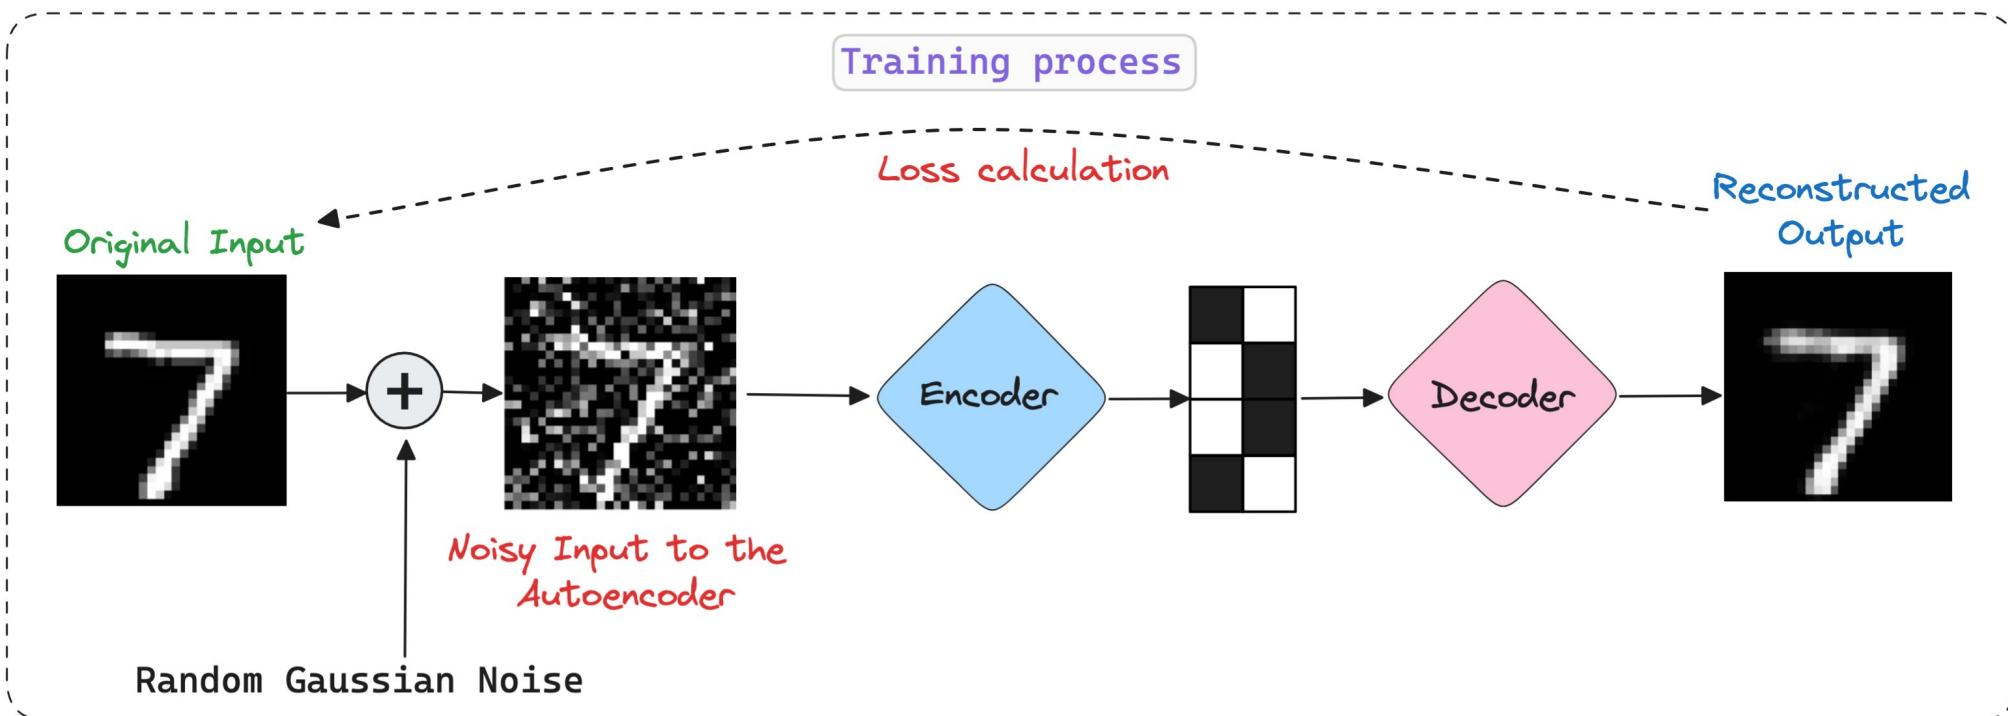
A glimpse on how a denoising autoencoder is trained

Swipe ➤

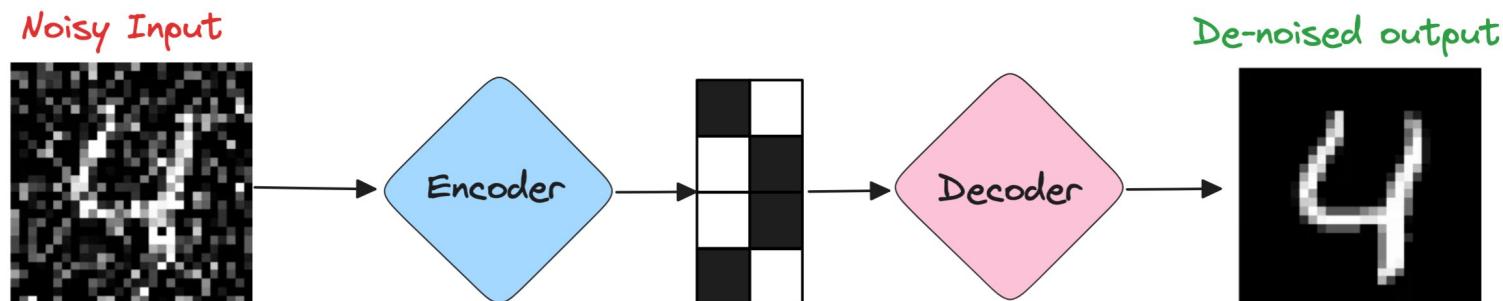


@akshay_pachaar

A denoising Autoencoder!



Inference



The encoder and decoder learns to remove noise!

Enough talk! Let's Code!

We'll use PyTorchLightning  for this.

- ◆ First, let's define our autoencoder!

A `LightningModule` enables your PyTorch `nn.Module` to play together in complex ways inside the `training_step`.

- ◆ I'll provide link to all the code at the end.

Swipe 



@akshay_pachaar

```
import os
from torch import optim, nn, utils, Tensor
from torchvision.datasets import MNIST
from torchvision.transforms import ToTensor
import lightning.pytorch as pl

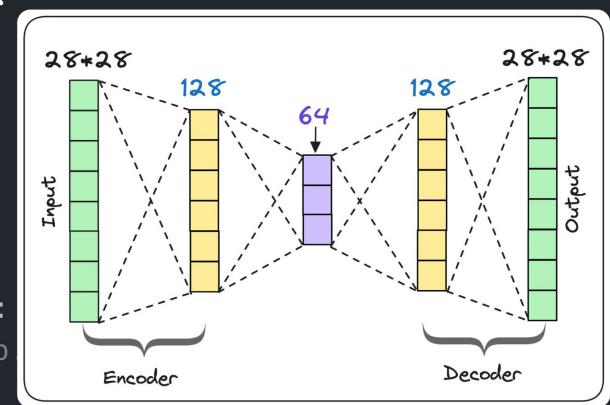
# define any number of nn.Modules (or use your current ones)
encoder = nn.Sequential(nn.Linear(28 * 28, 128), nn.ReLU(), nn.Linear(128, 64))
decoder = nn.Sequential(nn.Linear(64, 128), nn.ReLU(), nn.Linear(128, 28 * 28))
```

```
# define the LightningModule
class LitAutoEncoder(pl.LightningModule):
    def __init__(self, encoder, decoder):
        super().__init__()
        self.encoder = encoder
        self.decoder = decoder

    def training_step(self, batch, batch_idx):
        # training_step defines the train loop
        # it is independent of forward
        x, y = batch
        x = x.view(x.size(0), -1)
        z = self.encoder(x)
        x_hat = self.decoder(z)
        loss = nn.functional.mse_loss(x_hat, x)
        # Logging to TensorBoard (if installed) by default
        self.log("train_loss", loss)
        return loss

    def configure_optimizers(self):
        optimizer = optim.Adam(self.parameters(), lr=1e-3)
        return optimizer

# init the autoencoder
autoencoder = LitAutoEncoder(encoder, decoder)
```



- ◆ Define a dataset

Lightning supports ANY iterable (DataLoader, numpy, etc...) for the train/val/test/predict splits.

Swipe ➡



@akshay_pachaar



```
# setup data
dataset = MNIST(os.getcwd(), download=True, transform=ToTensor())
train_loader = utils.data.DataLoader(dataset)
```

- ◆ Train the model

The Lightning Trainer automates 40+ tricks including:

- Epoch and batch iteration
- `optimizer.step()`, `loss.backward()` etc.
- Calling of `model.eval()`
- enabling/disabling grads during evaluation
- Ckpt Saving and Loading
- Multi-GPU
- 16-bit precision AMP

Swipe ➡



@akshay_pachaar



```
# train the model
# (hint: here are some helpful Trainer arguments for rapid idea iteration)
trainer = pl.Trainer(limit_train_batches=100, max_epochs=1)
trainer.fit(model=autoencoder, train_dataloaders=train_loader)
```

Use the model

Once you've trained the model you can export to onnx, torchscript and put it into production or simply load the weights and run predictions.

Swipe ➡



@akshay_pachaar



```
# load checkpoint
checkpoint = "./lightning_logs/version_0/checkpoints/epoch=0-step=100.ckpt"
autoencoder = LitAutoEncoder.load_from_checkpoint(checkpoint, encoder=encoder, decoder=decoder)

# choose your trained nn.Module
encoder = autoencoder.encoder
encoder.eval()

# embed 4 fake images!
fake_image_batch = torch.rand(1, 28 * 28, device=autoencoder.device)
embeddings = encoder(fake_image_batch)
print("⚡" * 5, "\nPredictions (Embedding for an image):\n", embedding, "\n", "⚡" * 20)
```

⚡⚡⚡⚡⚡ *Predictions (Embedding for an image):*

tensor([[-0.0851, 0.1755, 0.3341, -0.3455, 0.1206, 0.3917,
-0.2160, 0.1877, 0.4324, -0.0497, -0.1032, 0.0408, -0.1427, -0.0320, 0.6609, - ...]])

Visualize results

Let's encode/decode an image using our trained model!

Visualization shows an original image which we encode & a reconstructed image!

Swipe ➡



@akshay_pachaar



```
import matplotlib.pyplot as plt
from torchvision.utils import make_grid

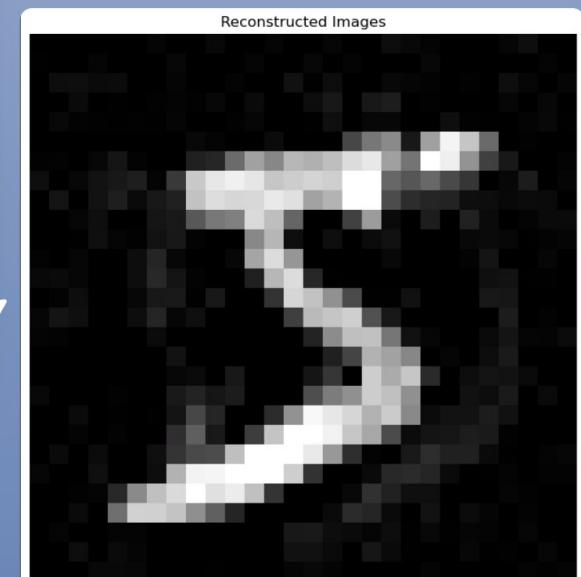
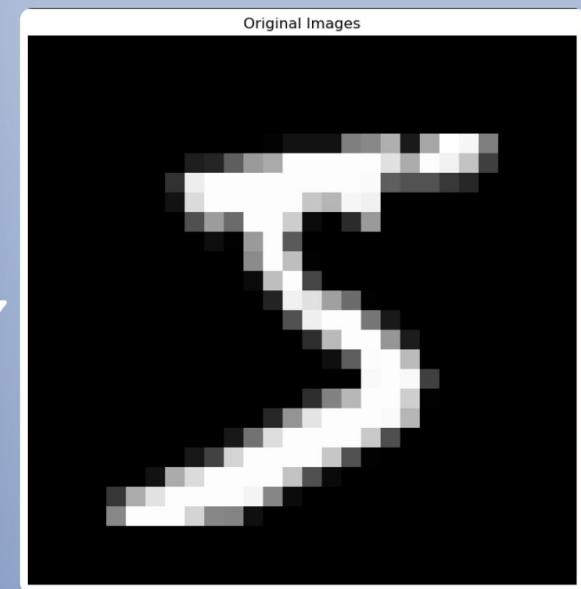
def show_images(images, title="Images"):
    """Utility function to display a batch of images."""
    grid_img = make_grid(images, nrow=4)
    plt.figure(figsize=(8, 8))
    plt.imshow(grid_img.permute(1, 2, 0))
    plt.title(title)
    plt.axis('off')
    plt.show()

# Load a batch of images from the dataset
images, _ = next(iter(train_loader))
show_images(images, title="Original Images")
```



```
# Preprocess the images
images = images.view(images.size(0), -1)

# Generate image from embeddings
embeddings = autoencoder.encoder(images)
reconstructed_images = autoencoder.decoder(embeddings).view(-1, 1, 28, 28)
show_images(reconstructed_images, title="Reconstructed Images")
```



That's a wrap!

If you interested in:

- Python 
- Data Science 
- Machine Learning 
- MLOps 
- NLP 
- Computer Vision 
- LLMs 

Follow me on LinkedIn 

Everyday, I share tutorials on above topics!

Cheers !! 



@akshay_pachaar