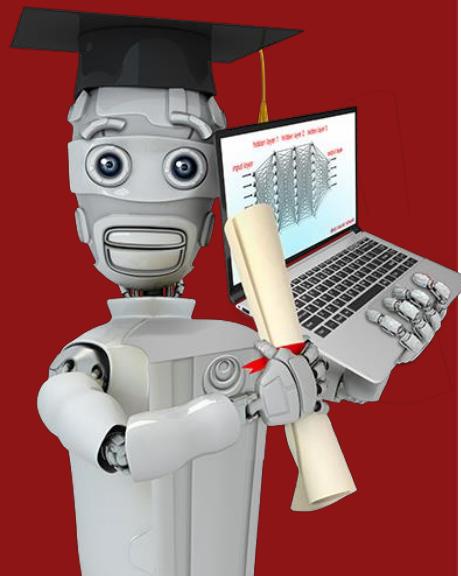


 DeepLearning.AI

Stanford
ONLINE



Advanced Learning Algorithms

Welcome!

Advanced learning algorithms

Neural Networks 
inference (prediction)
training
Practical advice for building machine learning systems 
Decision Trees 



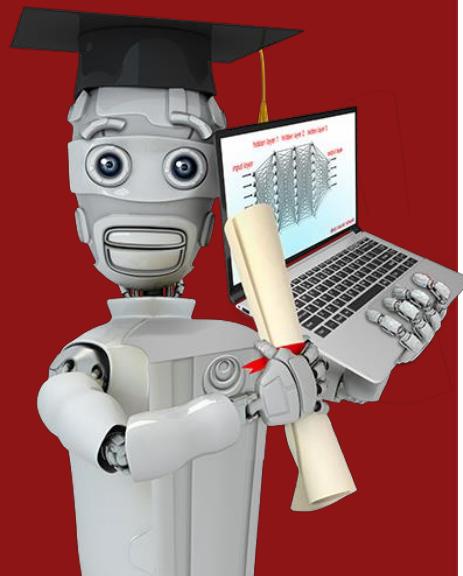
DeepLearning.AI

Stanford ONLINE

Andrew Ng

DeepLearning.AI

Stanford
ONLINE

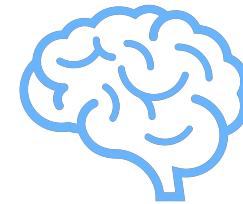


Neural Networks Intuition

Neurons and the brain

Neural networks

Origins: Algorithms that try to mimic the brain.



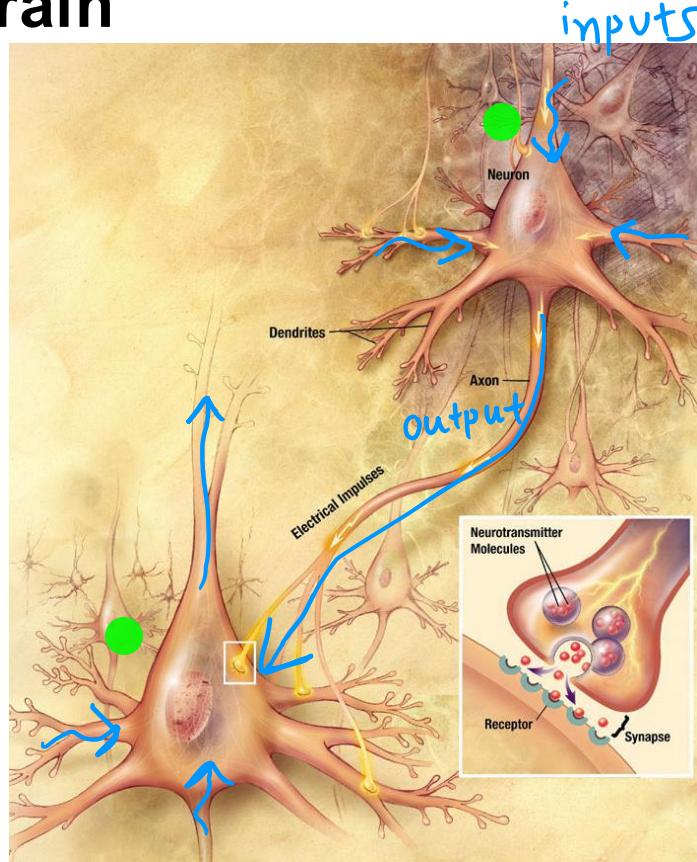
Used in the 1980's and early 1990's.
Fell out of favor in the late 1990's.

Resurgence from around 2005.

speech → images → text (NLP) → ...



Neurons in the brain

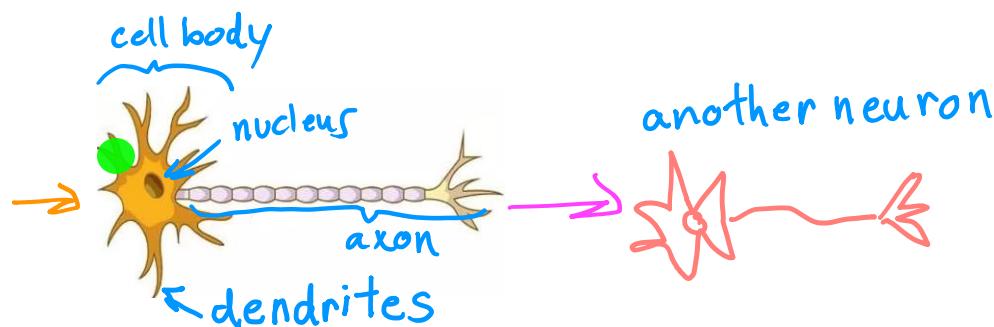


[Credit: US National Institutes of Health, National Institute on Aging]

Biological neuron

inputs

outputs



Simplified mathematical model of a neuron

inputs

outputs

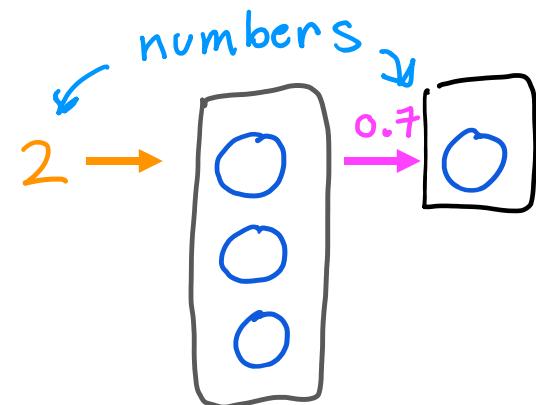
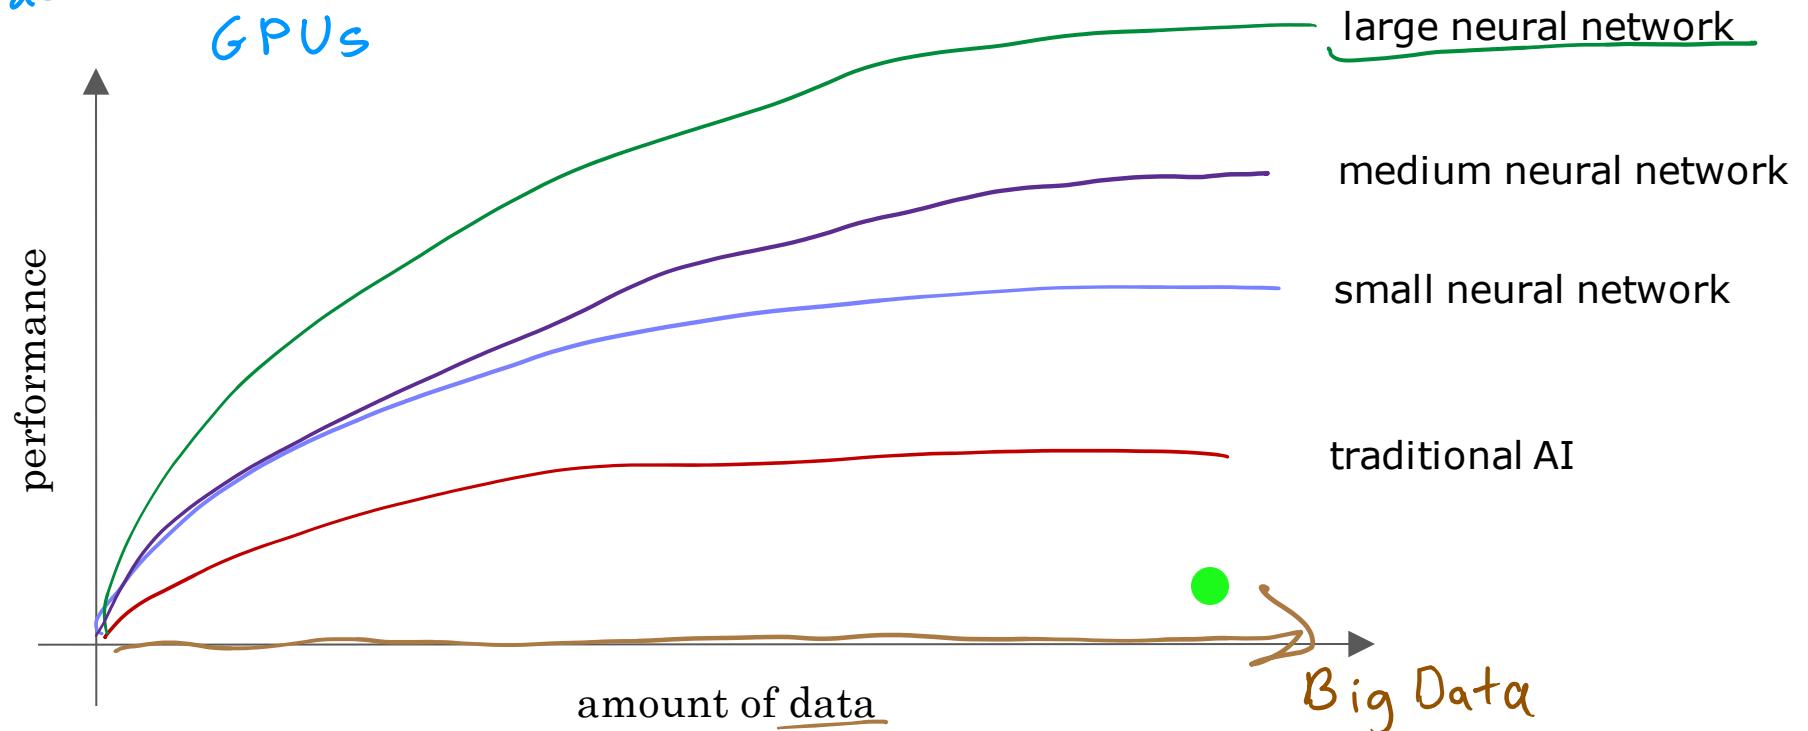


image source: <https://biologydictionary.net/sensory-neuron/>

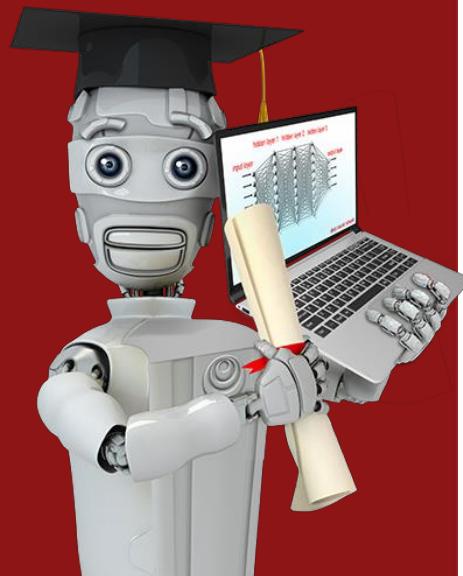
Faster computer processors
GPUs

Why Now?



DeepLearning.AI

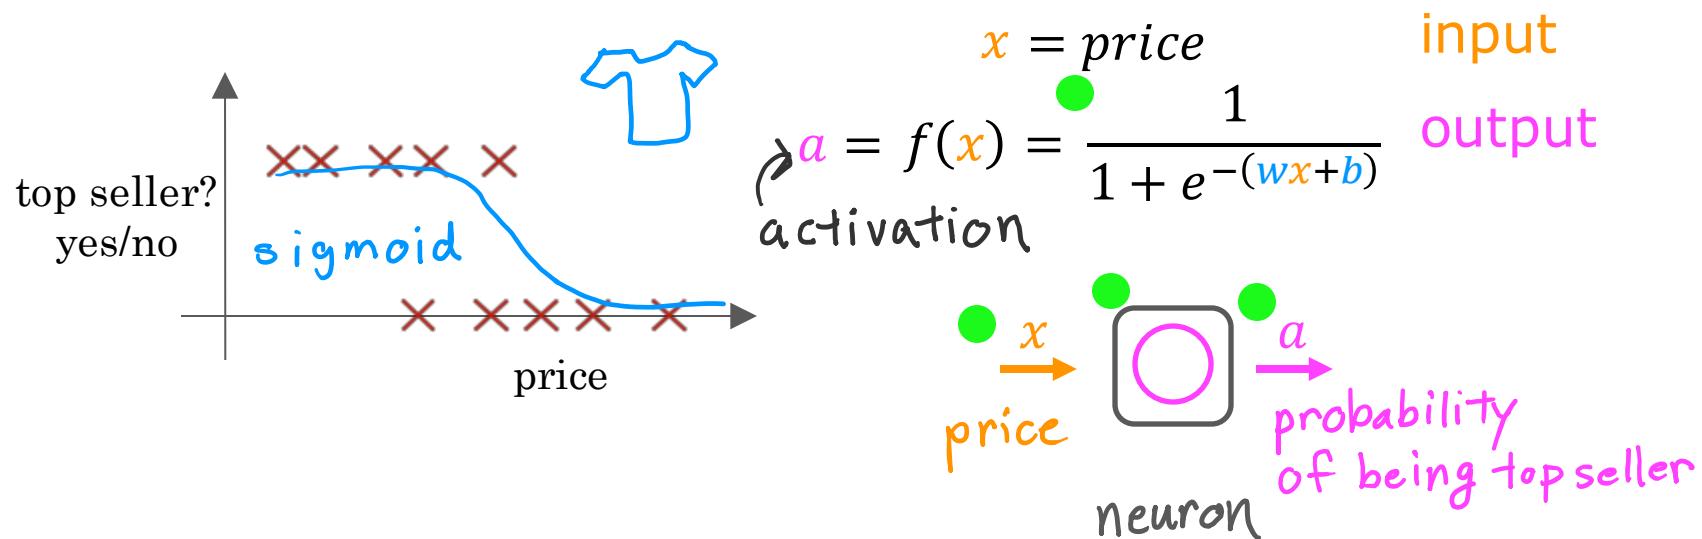
Stanford
ONLINE



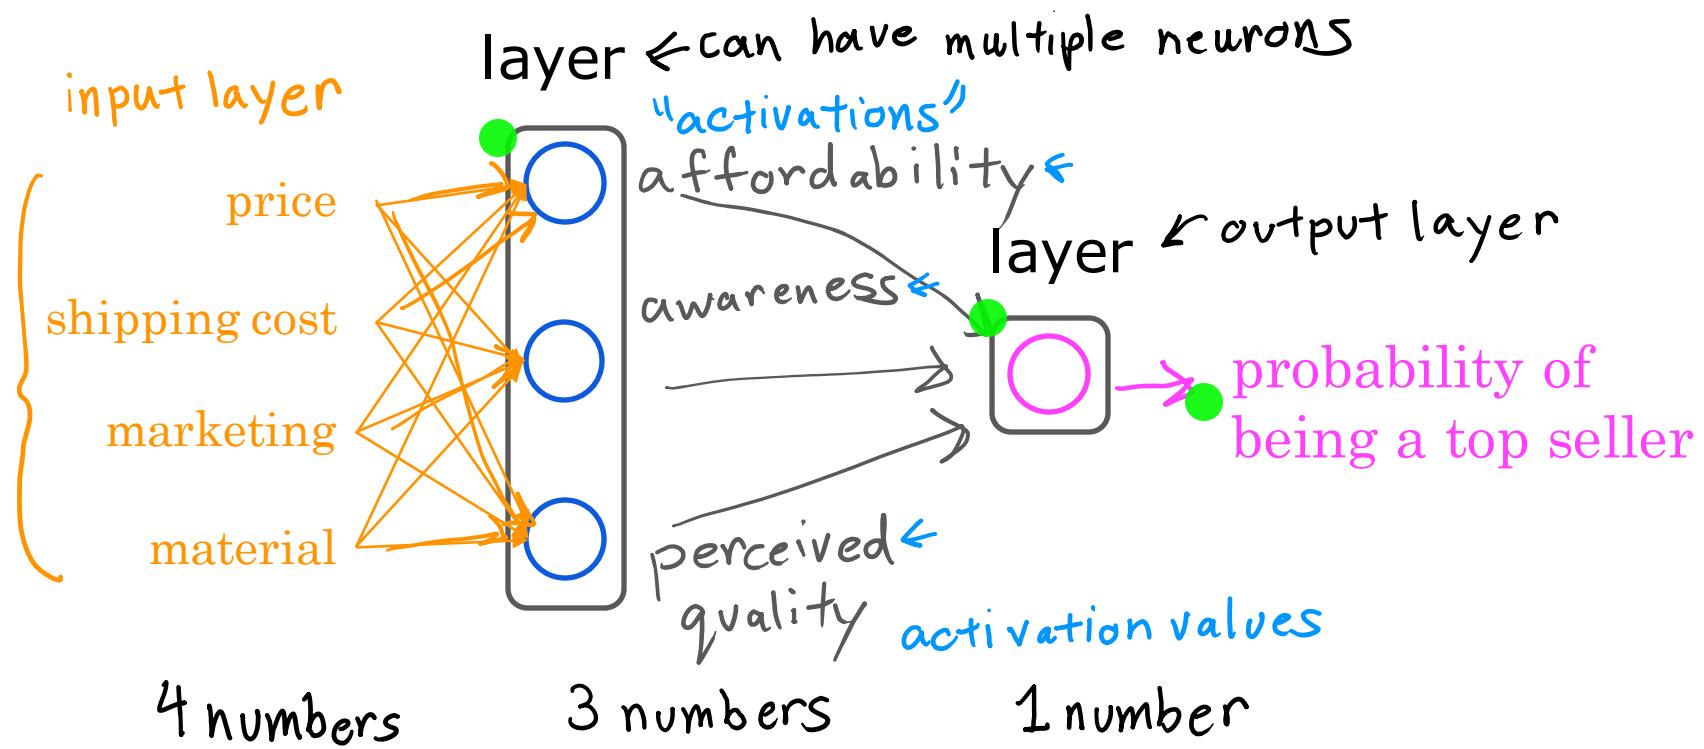
Neural Network Intuition

Demand Prediction

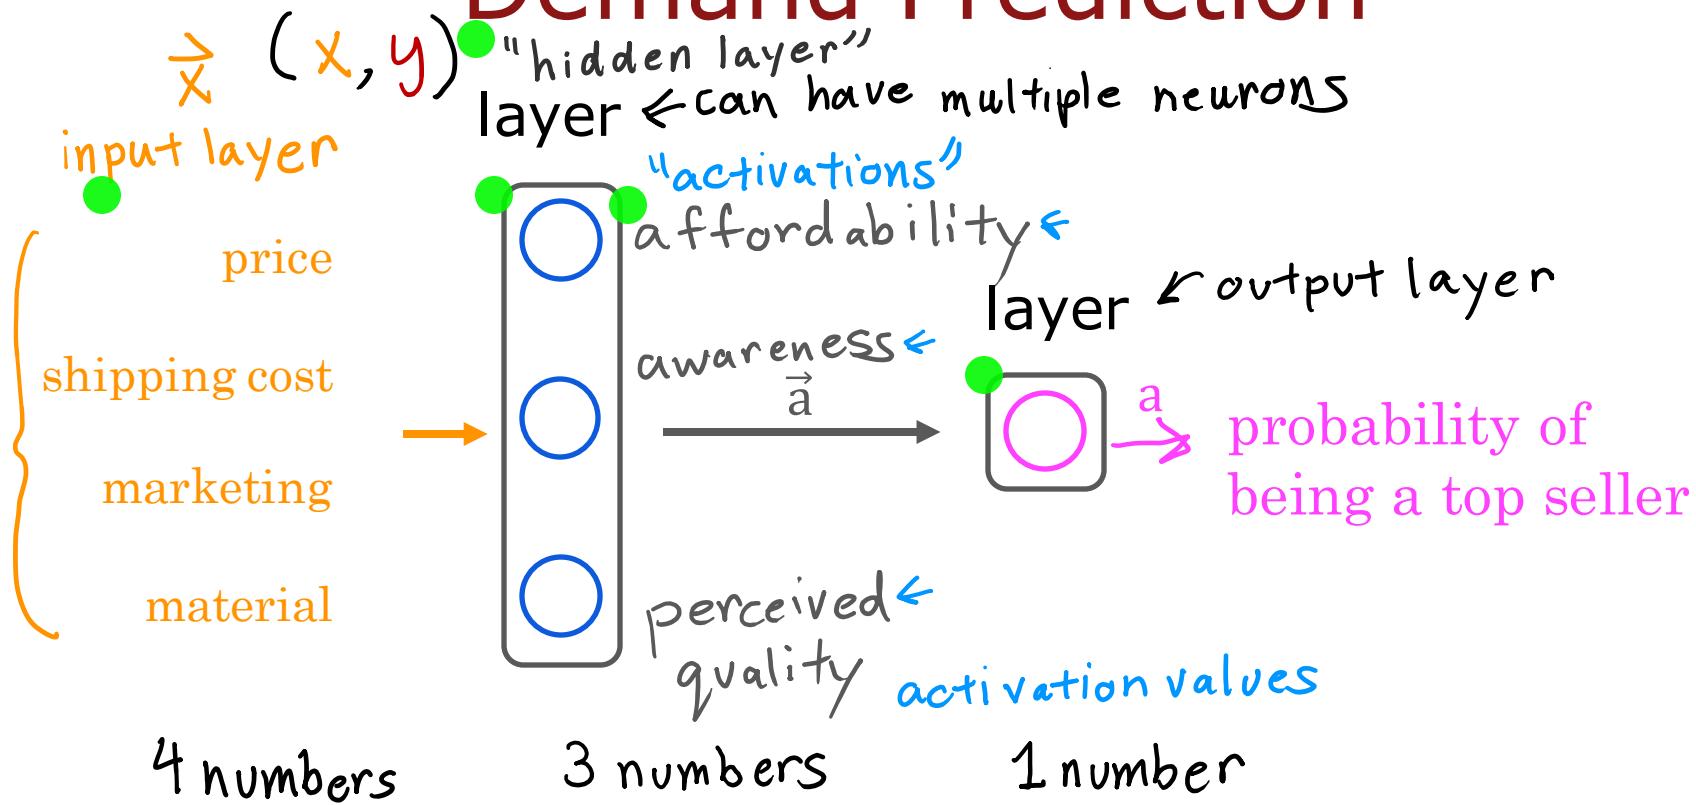
Demand Prediction



Demand Prediction

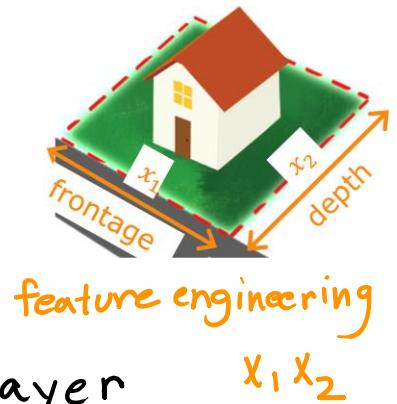
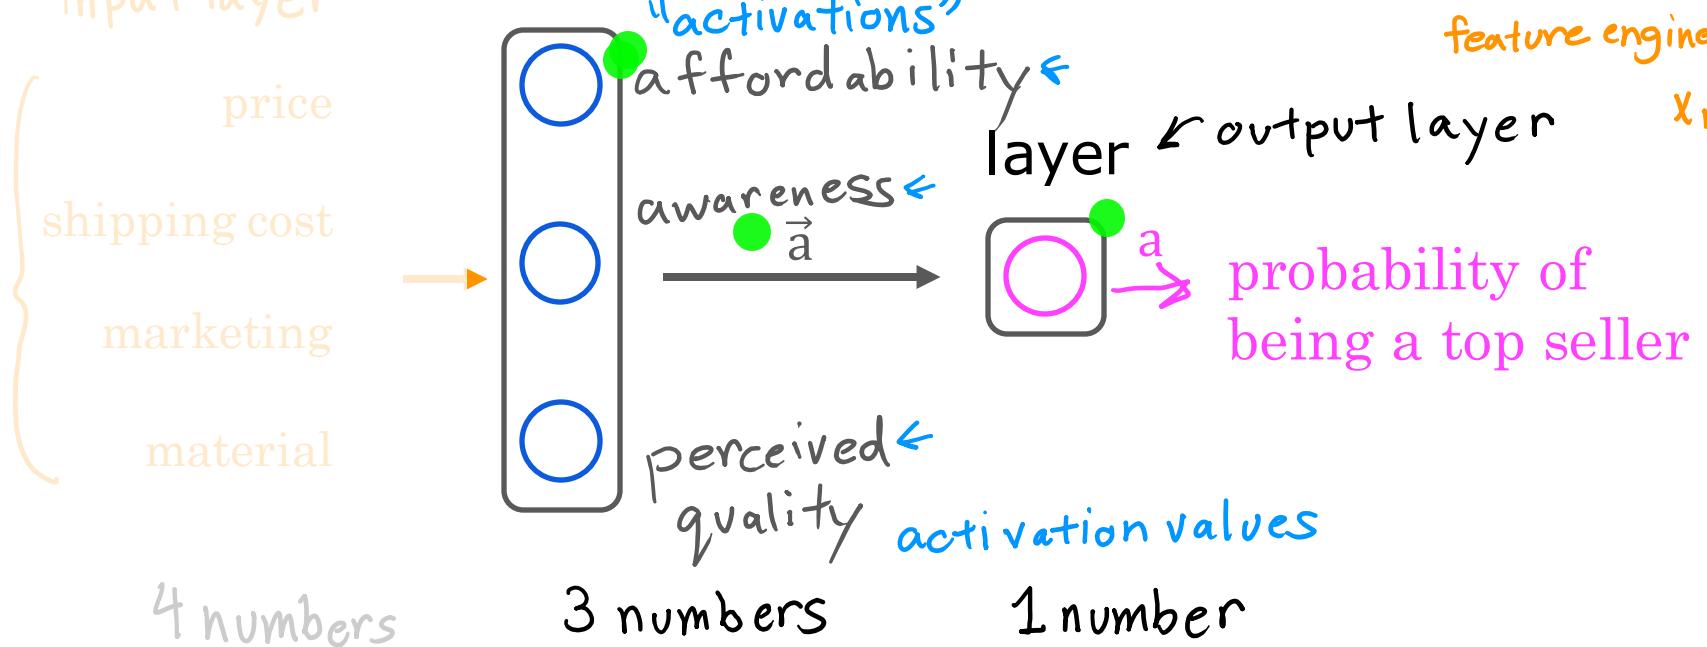


Demand Prediction

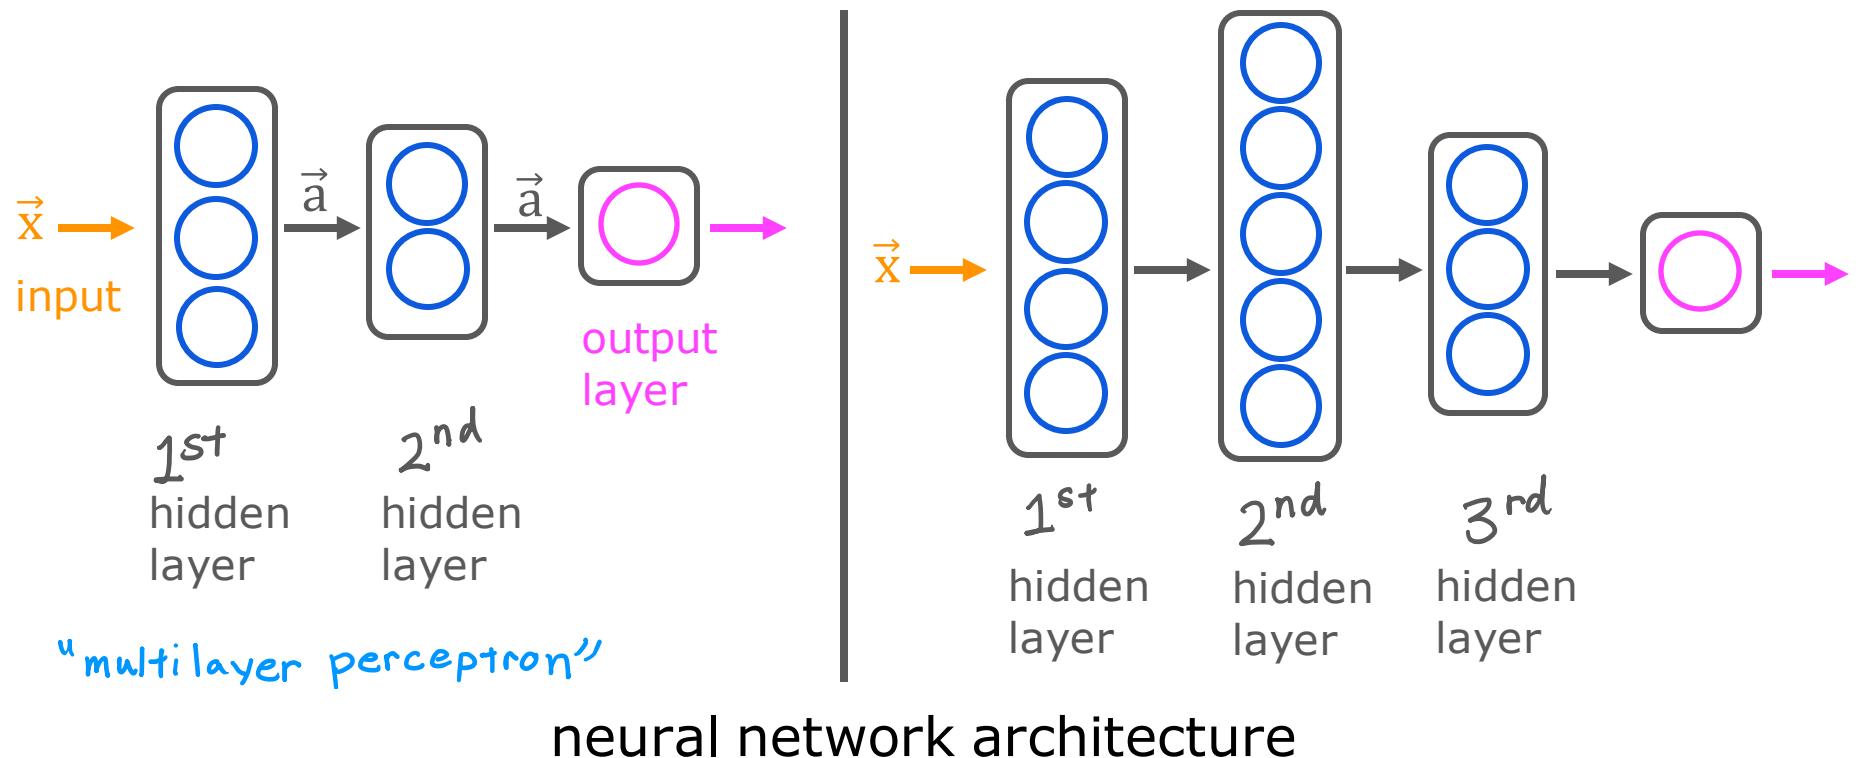


Demand Prediction

• (x, y) "hidden layer" layer can have multiple neurons

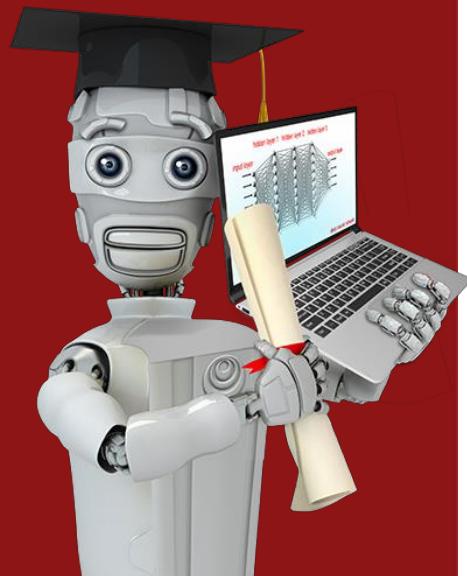


Multiple hidden layers



 DeepLearning.AI

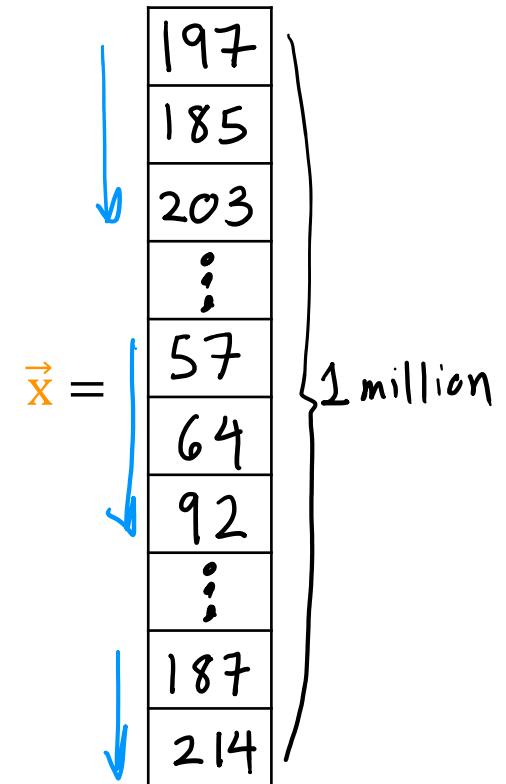
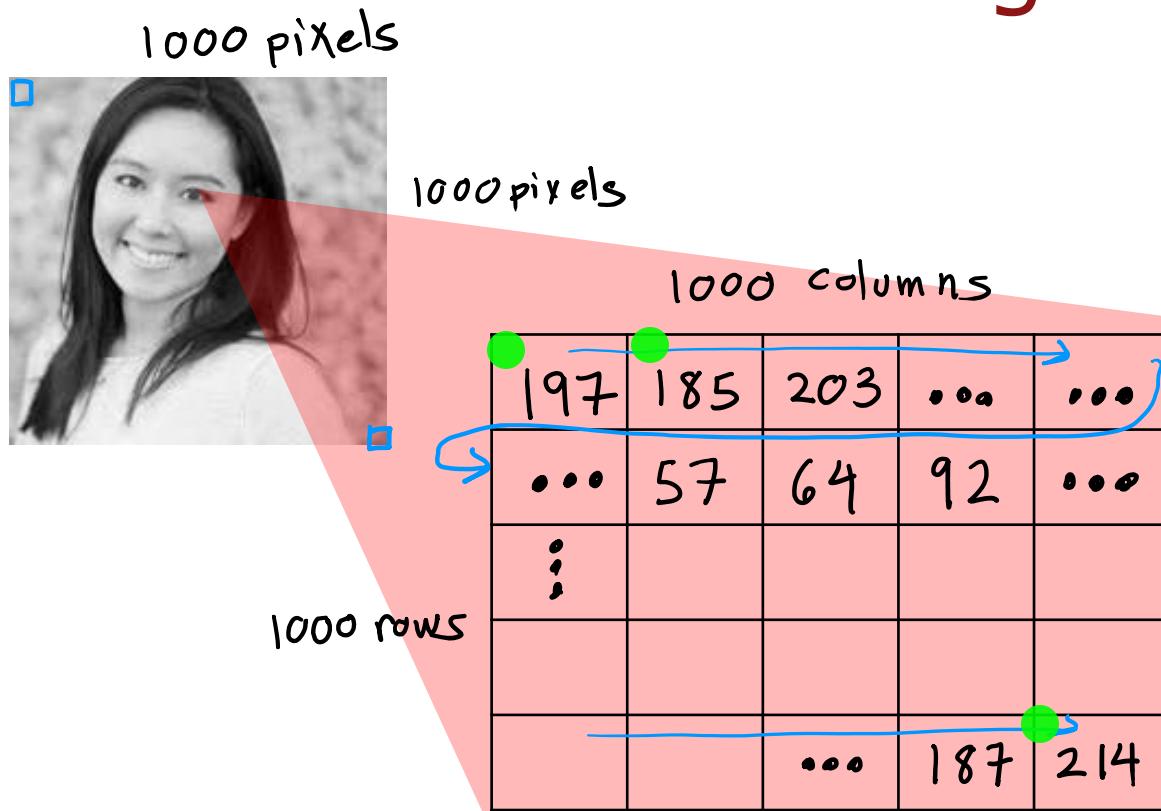
Stanford
ONLINE



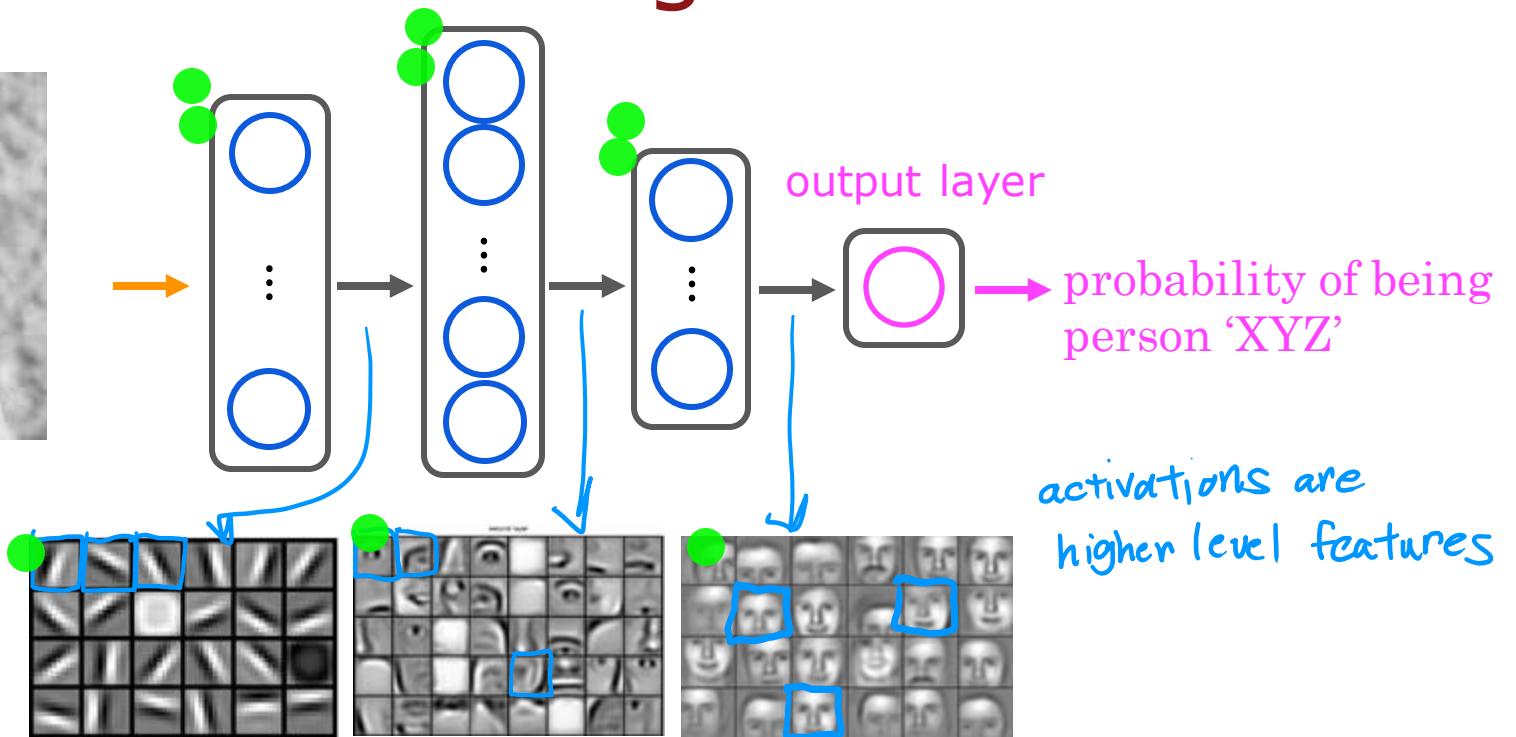
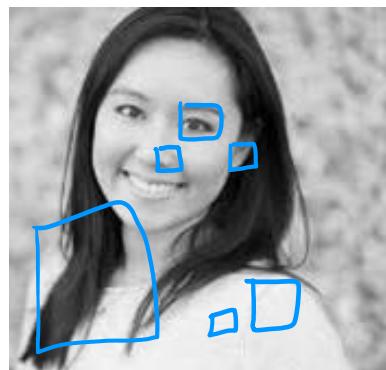
Neural Networks Intuition

Example:
Recognizing Images

Face recognition

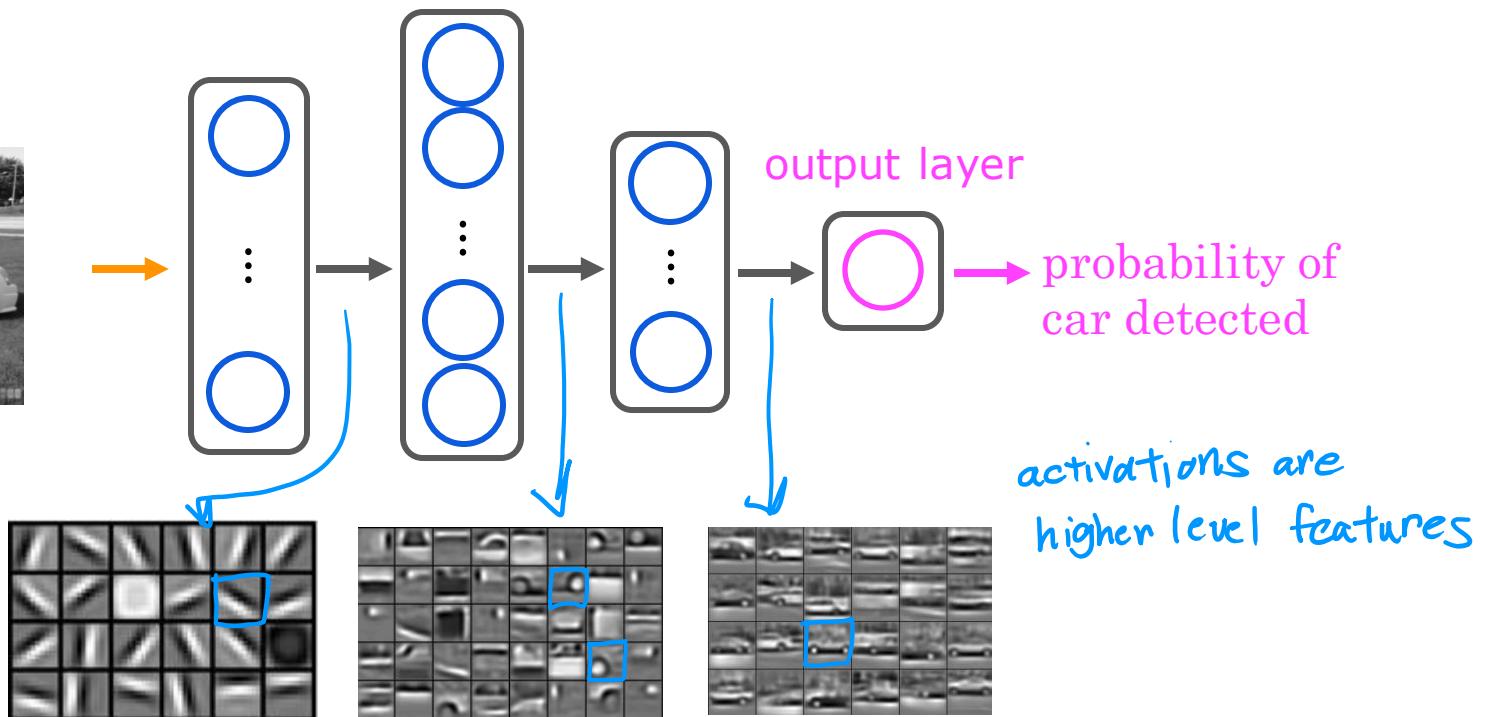


Face recognition



source: Convolutional Deep Belief Networks for Scalable Unsupervised Learning of Hierarchical Representations
by Honglak Lee, Roger Grosse, Ranganath Andrew Y. Ng

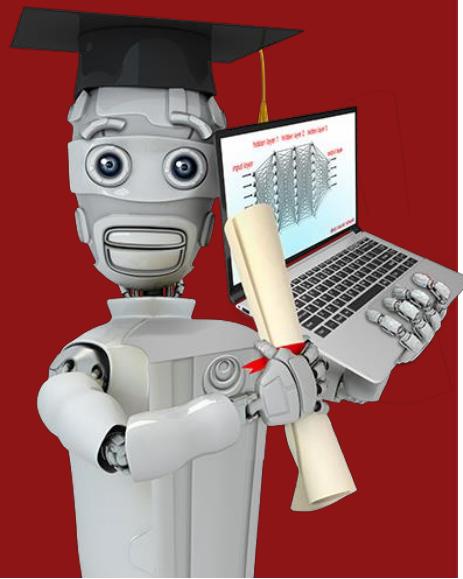
Car classification



source: Convolutional Deep Belief Networks for Scalable Unsupervised Learning of Hierarchical Representations
by Honglak Lee, Roger Grosse, Ranganath Andrew Y. Ng

DeepLearning.AI

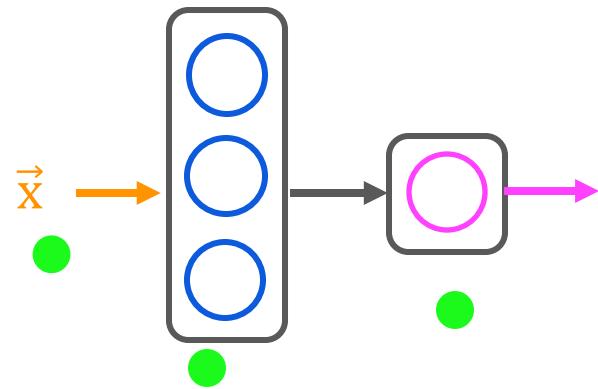
Stanford
ONLINE



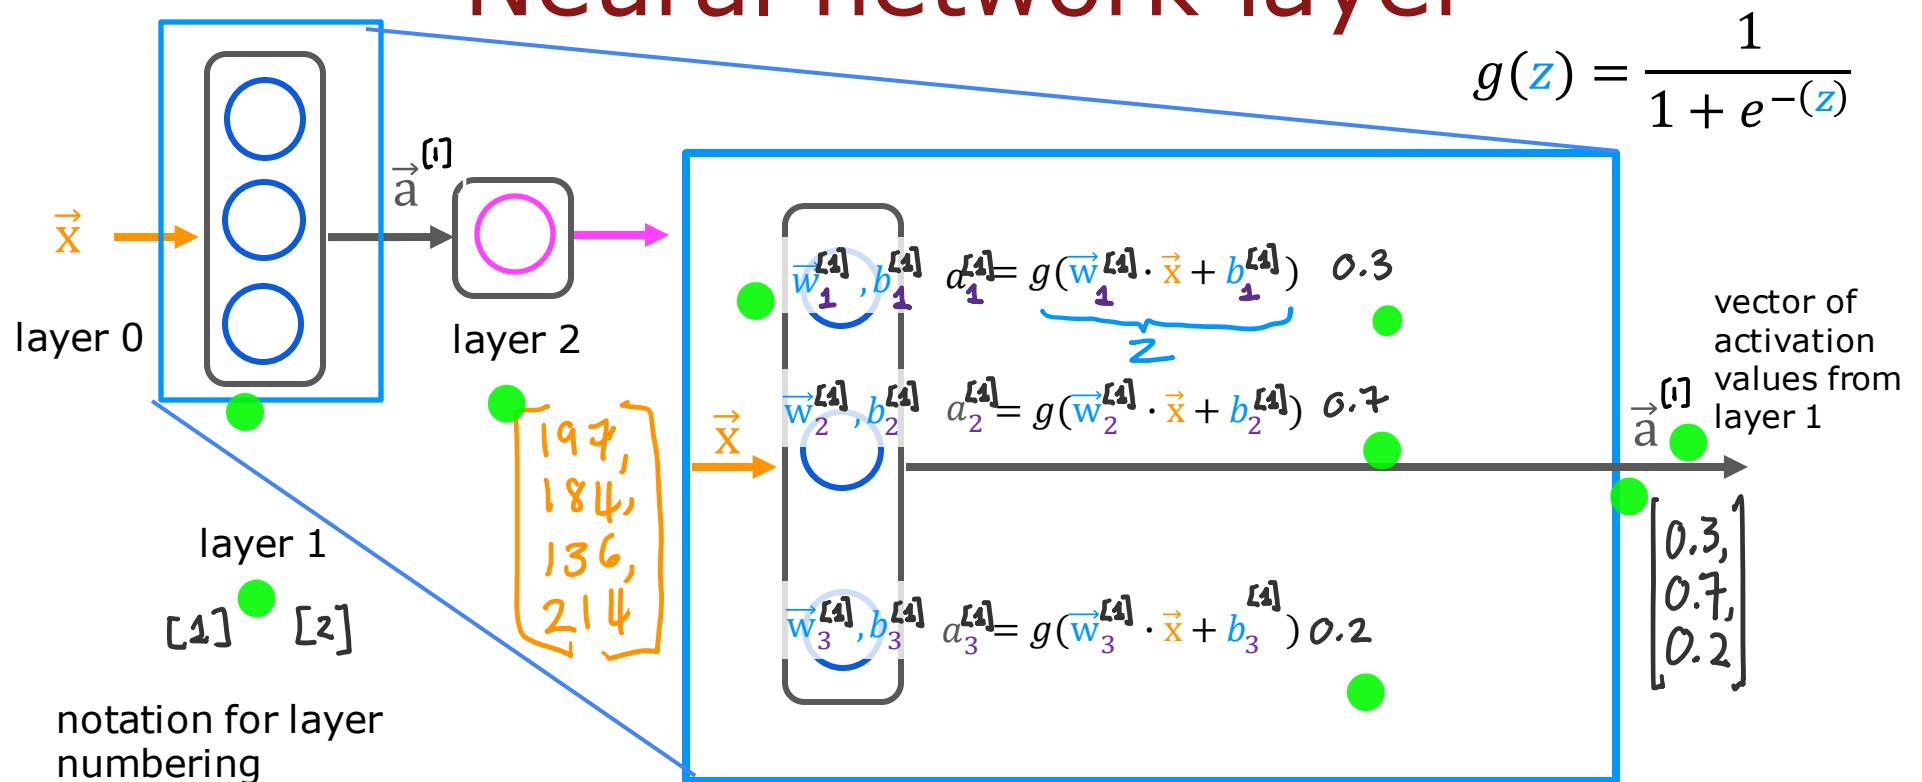
Neural network model

Neural network layer

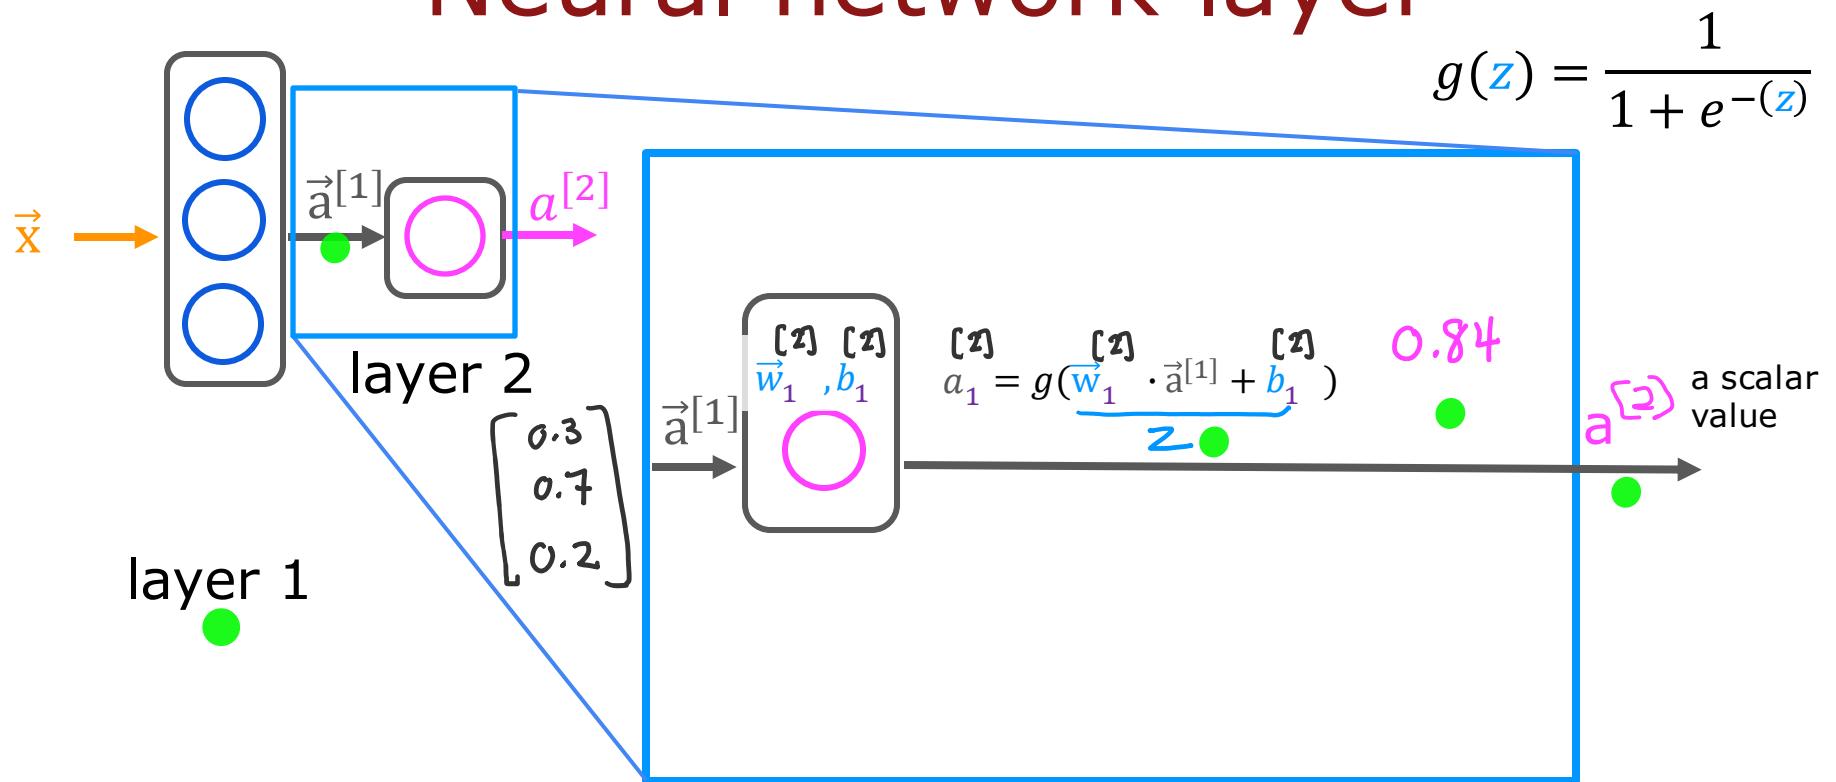
Neural network layer



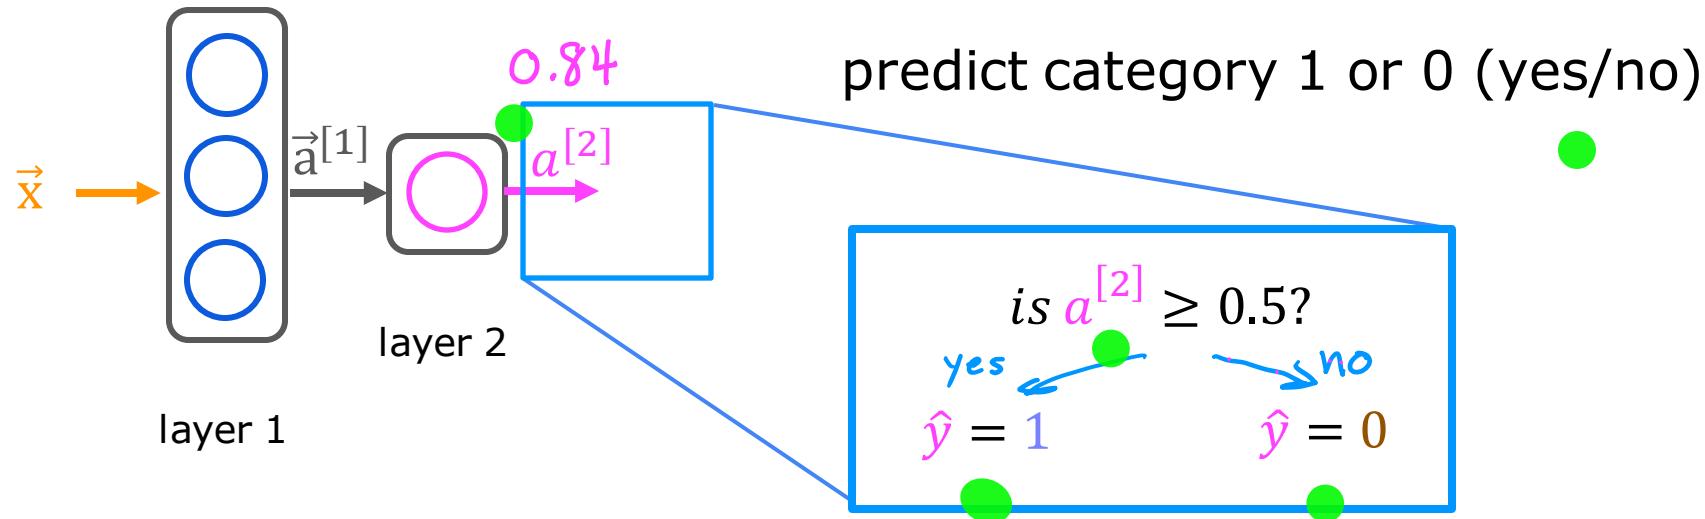
Neural network layer



Neural network layer

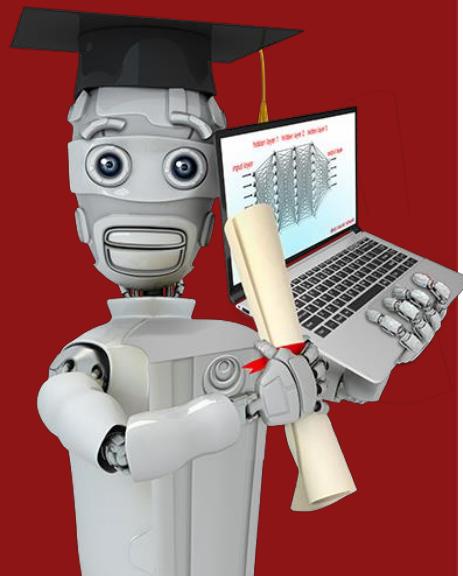


Neural network layer



 DeepLearning.AI

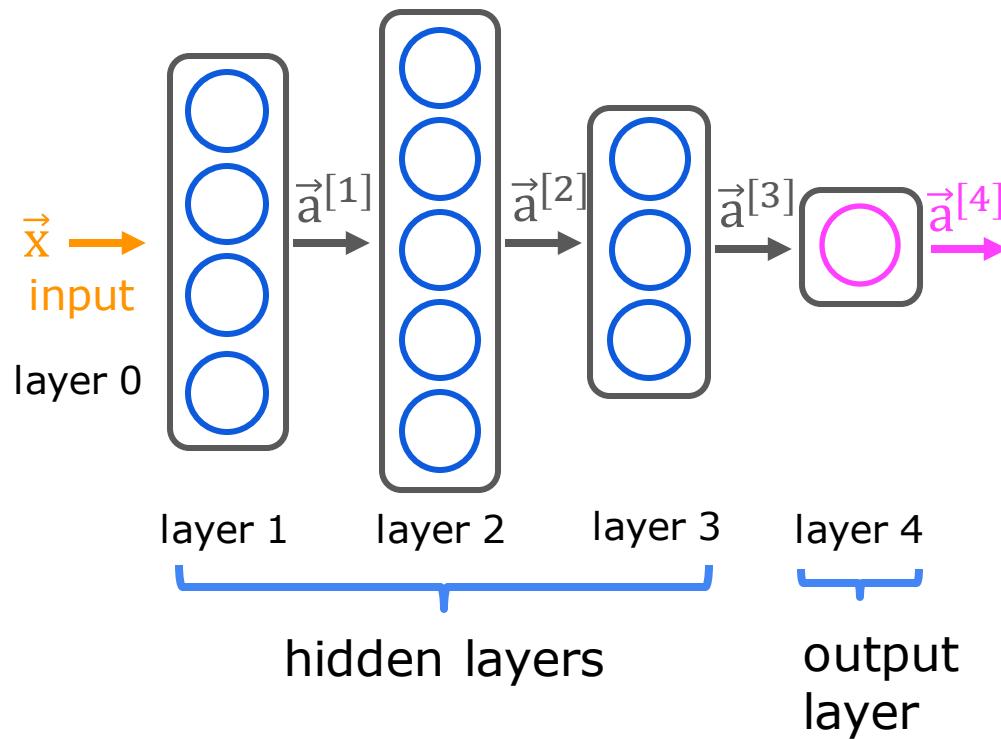
Stanford
ONLINE



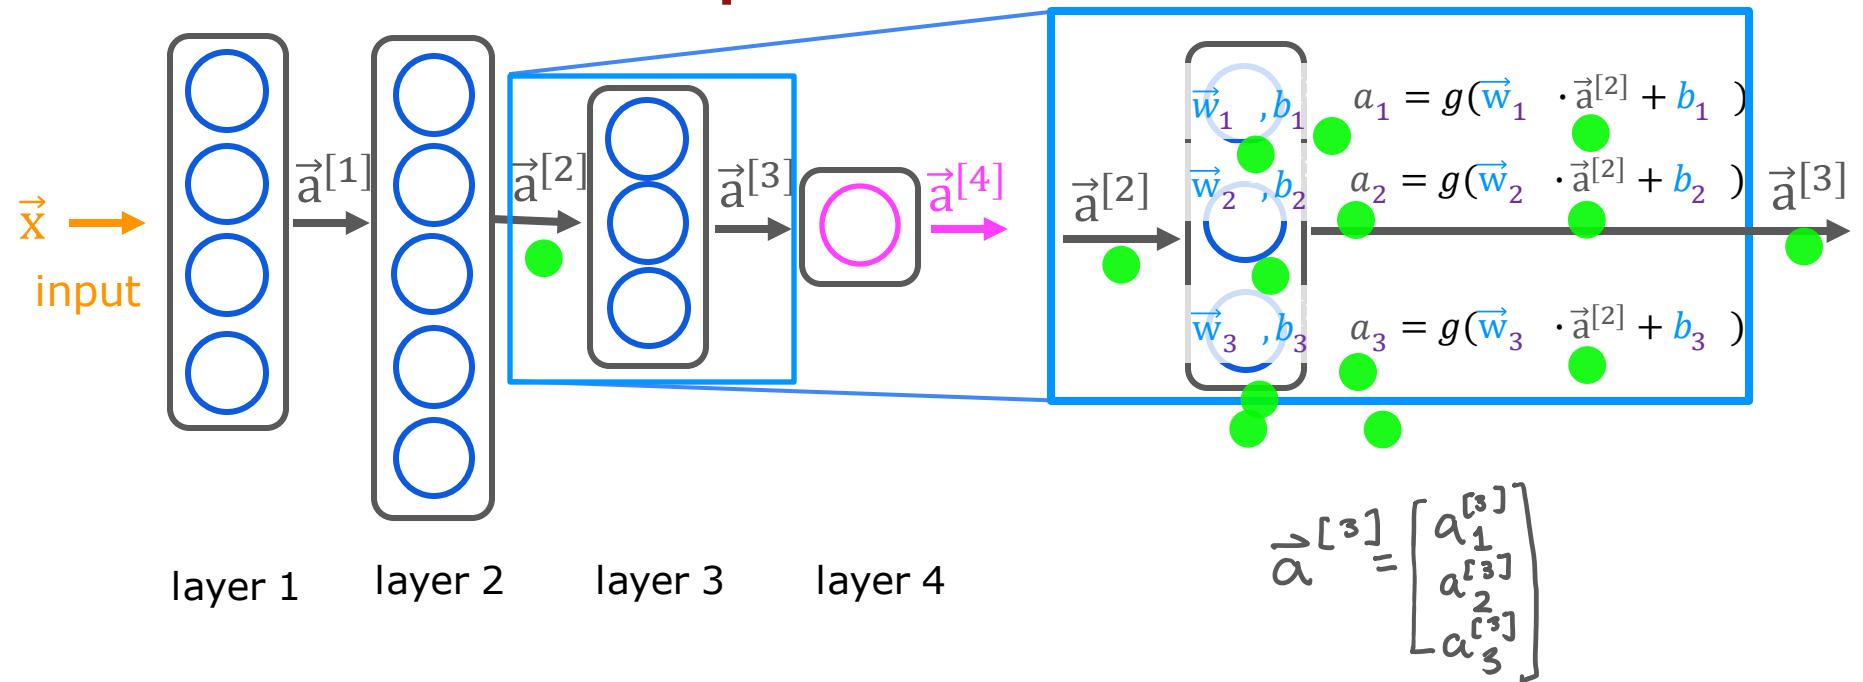
Neural Network Model

More complex neural networks

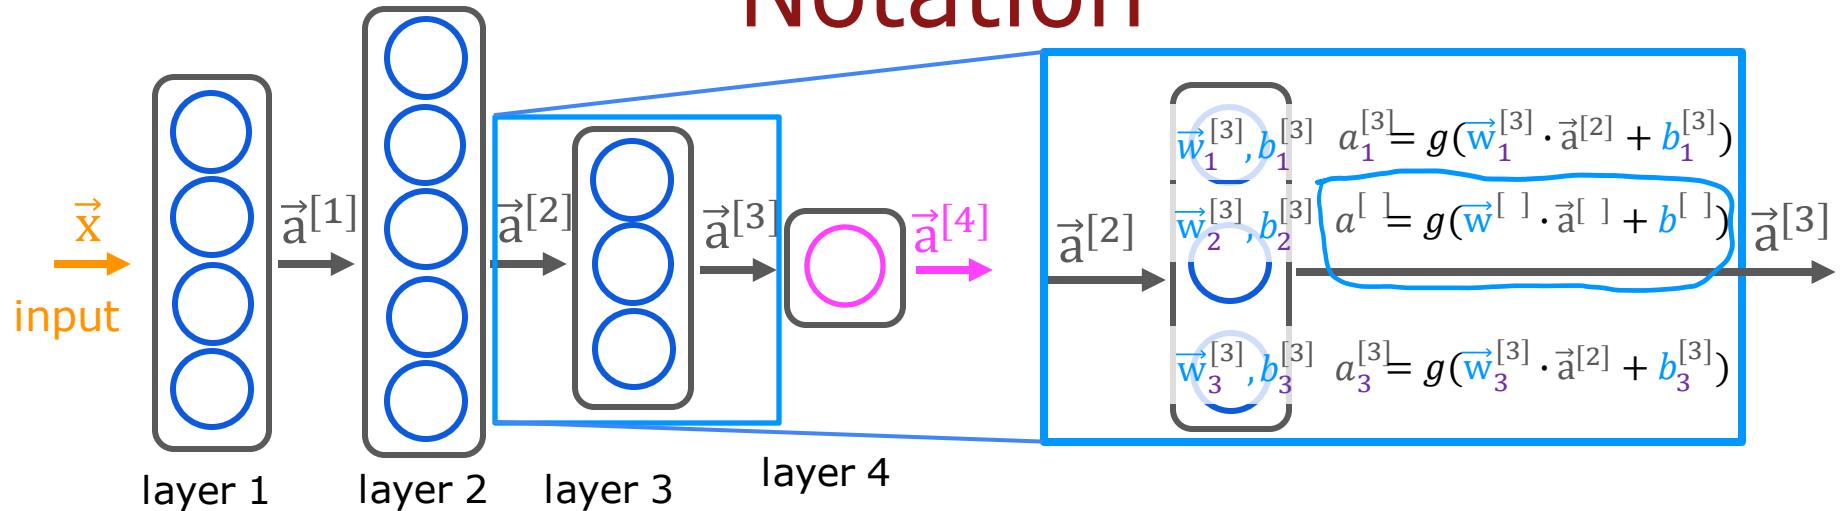
More complex neural network



More complex neural network

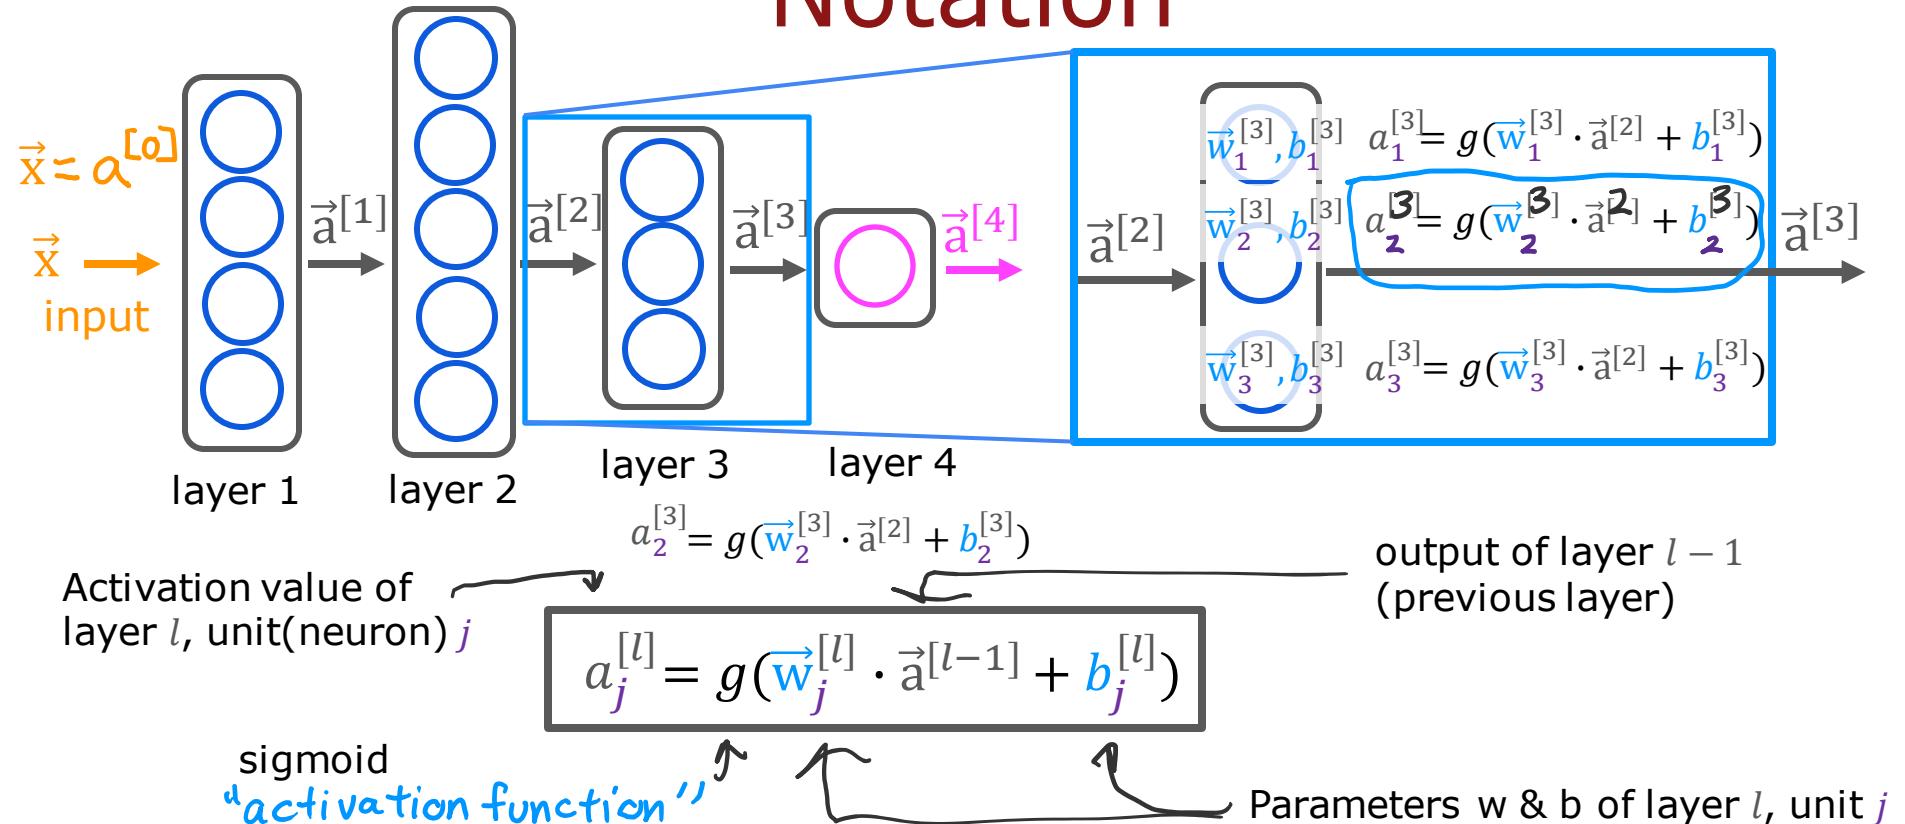


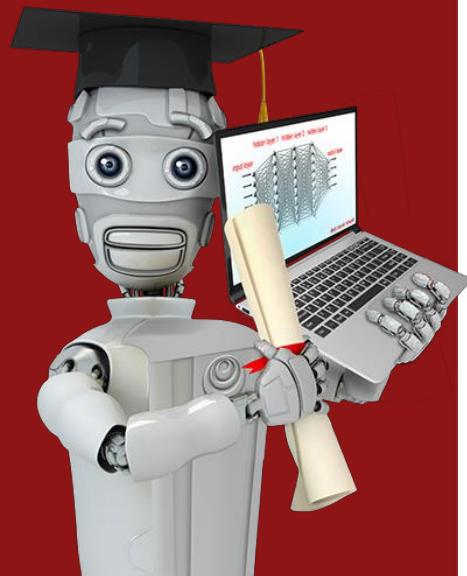
Notation



Question:
Can you fill in the superscripts and
subscripts for the second neuron?

Notation

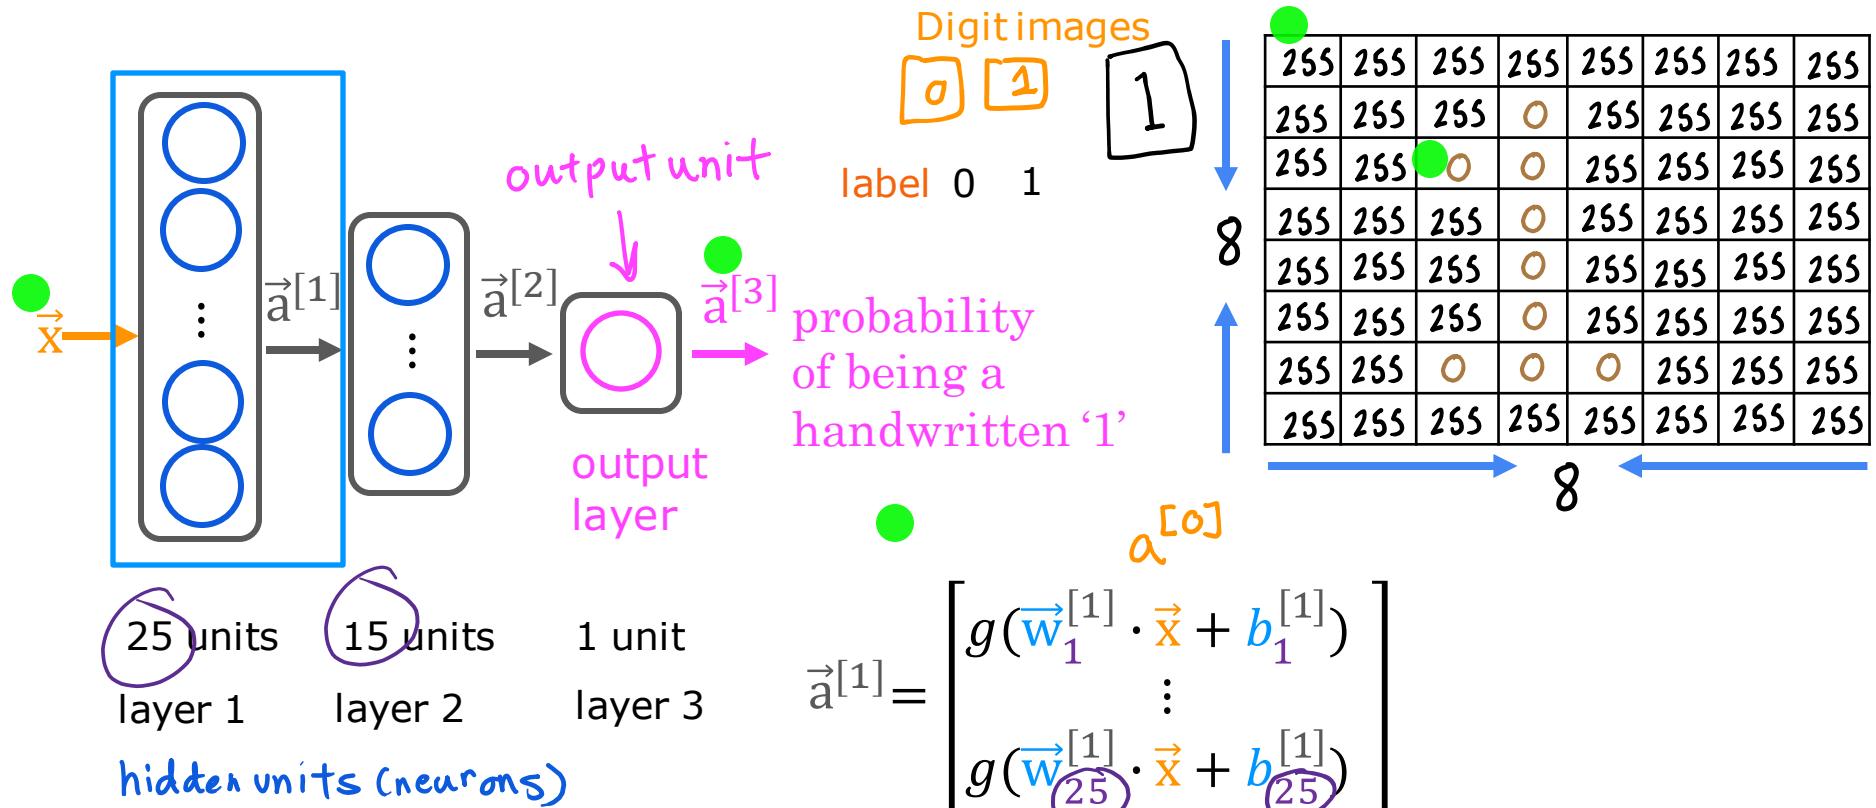




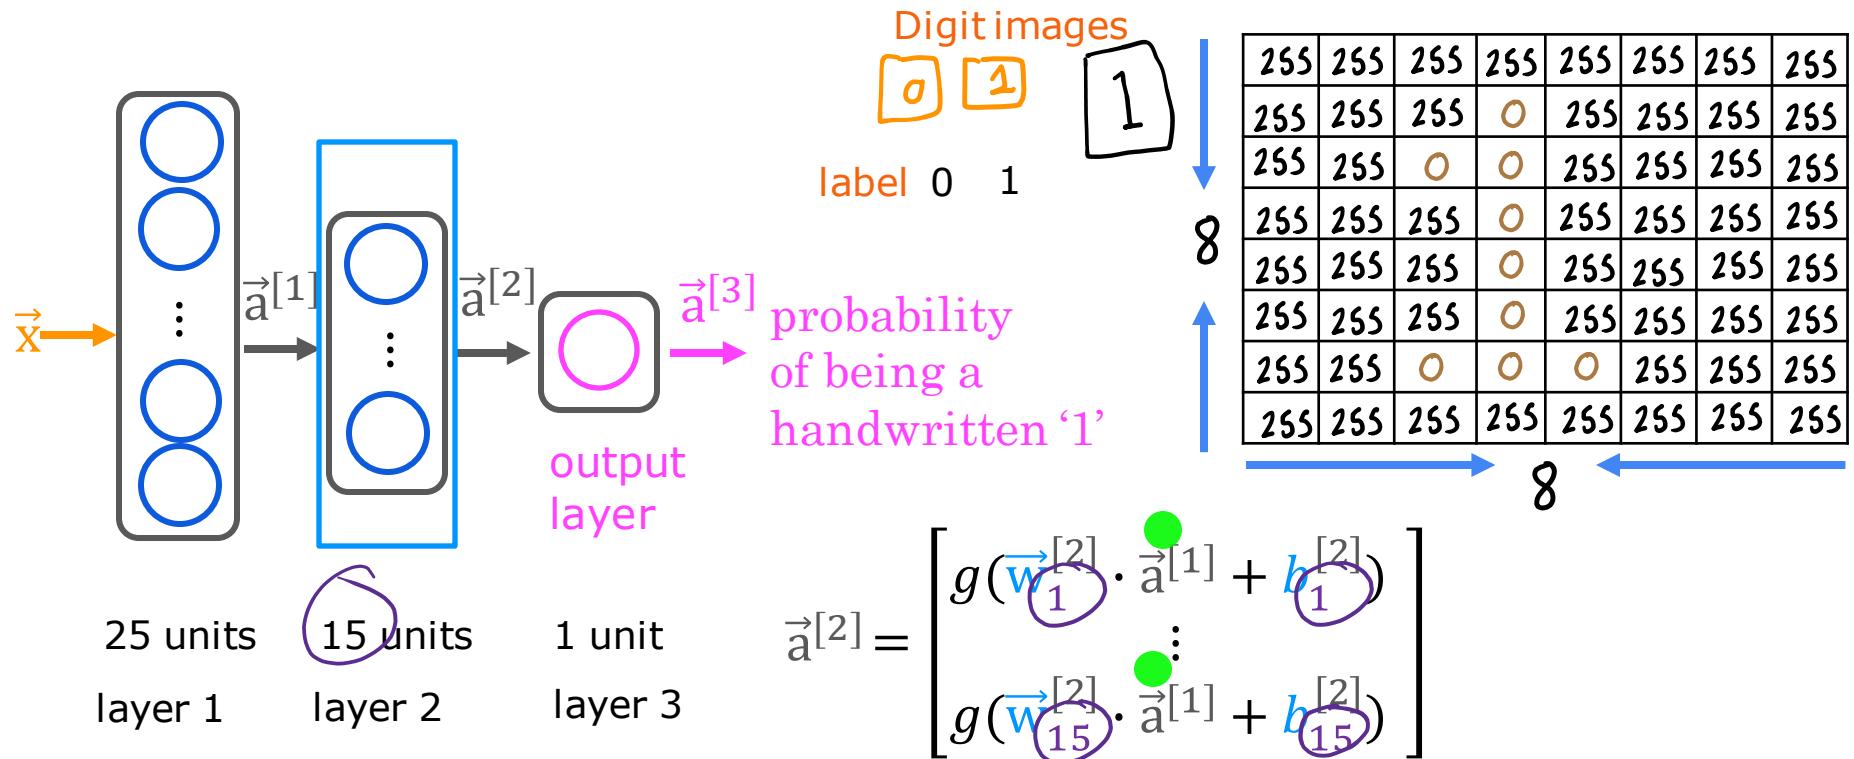
Neural Network Model

**Inference: making predictions
(forward propagation)**

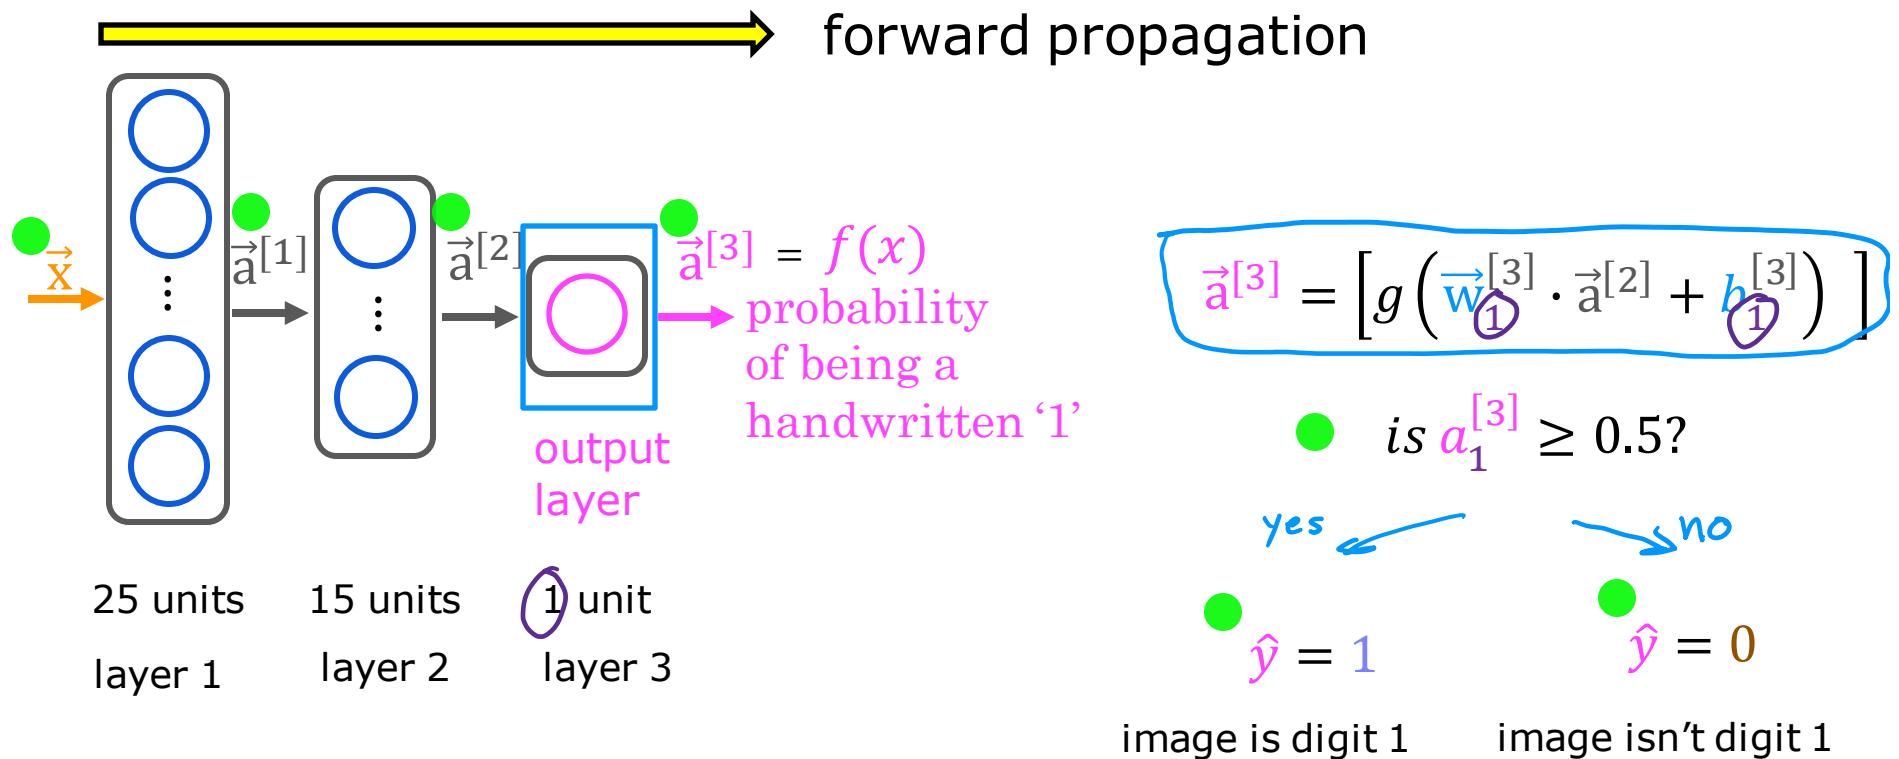
Handwritten digit recognition



Handwritten digit recognition

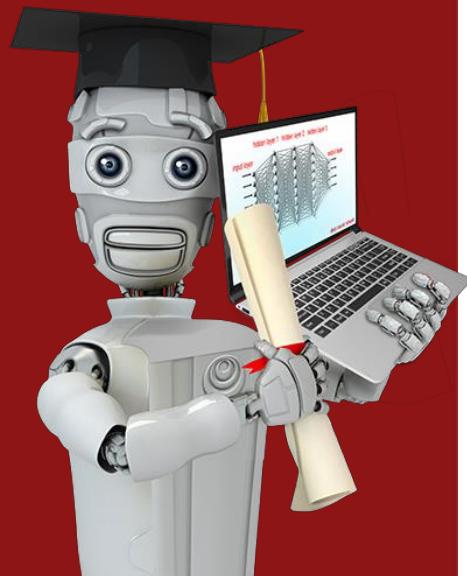


Handwritten digit recognition



 DeepLearning.AI

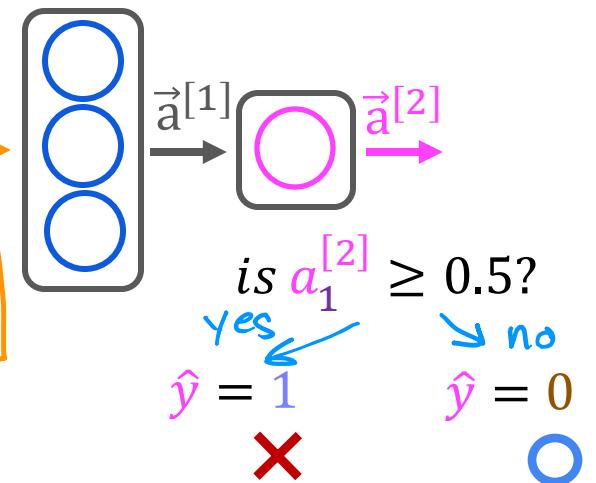
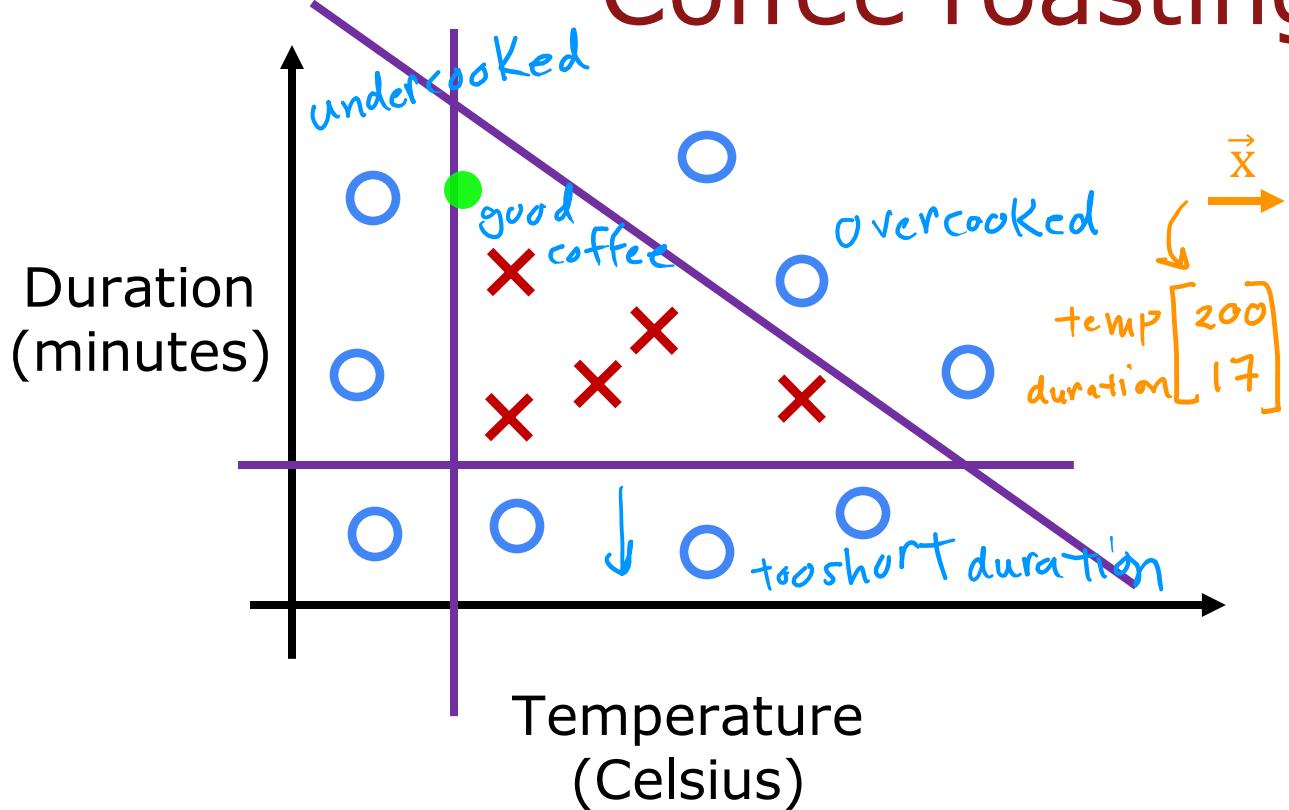
Stanford
ONLINE

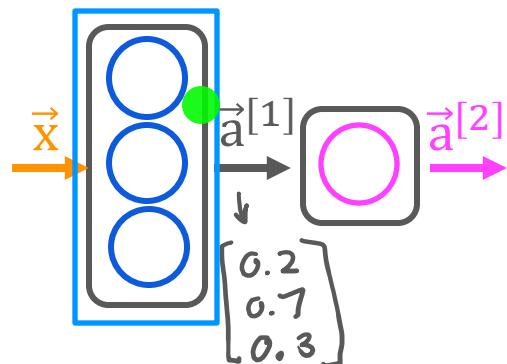


TensorFlow implementation

Inference in Code

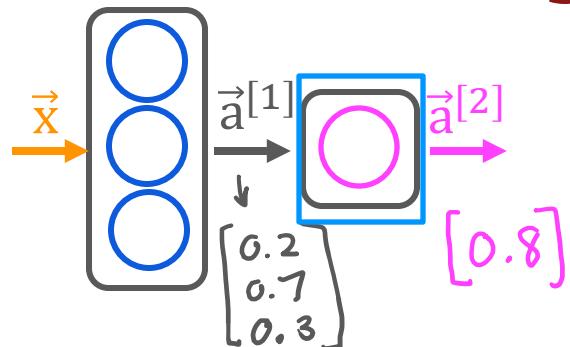
Coffee roasting





```
x = np.array([[200.0, 17.0]])
layer_1 = Dense(units=3, activation='sigmoid')
a1 = layer_1(x)
```

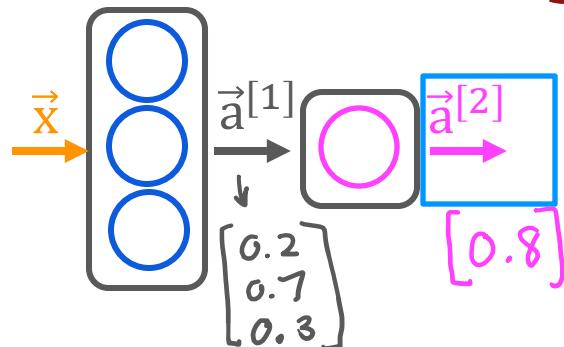
Build the model using TensorFlow



```
x = np.array([[200.0, 17.0]])
layer_1 = Dense(units=3, activation='sigmoid')
a1 = layer_1(x)
```

```
layer_2 = Dense(units=1, activation='sigmoid')
a2 = layer_2(a1)
```

Build the model using TensorFlow



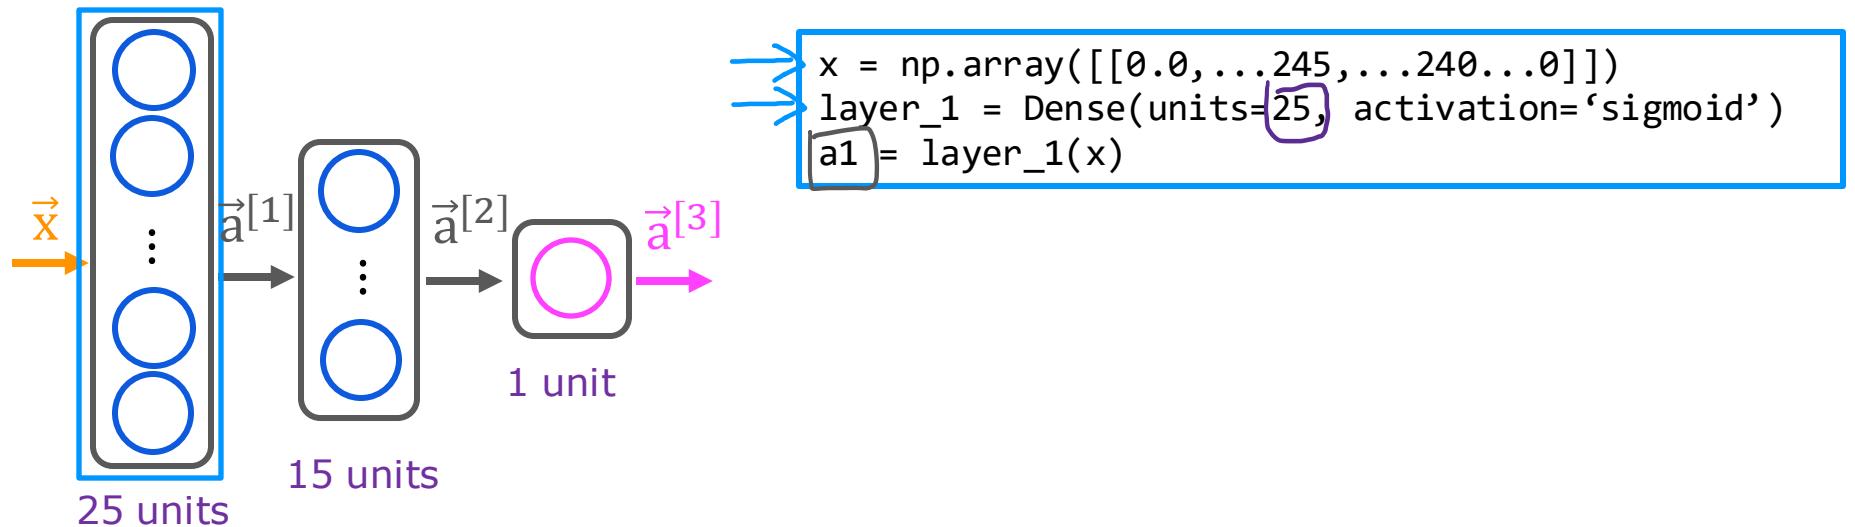
```
x = np.array([[200.0, 17.0]])
layer_1 = Dense(units=3, activation='sigmoid')
a1 = layer_1(x)
```

```
layer_2 = Dense(units=1, activation='sigmoid')
a2 = layer_2(a1)
```

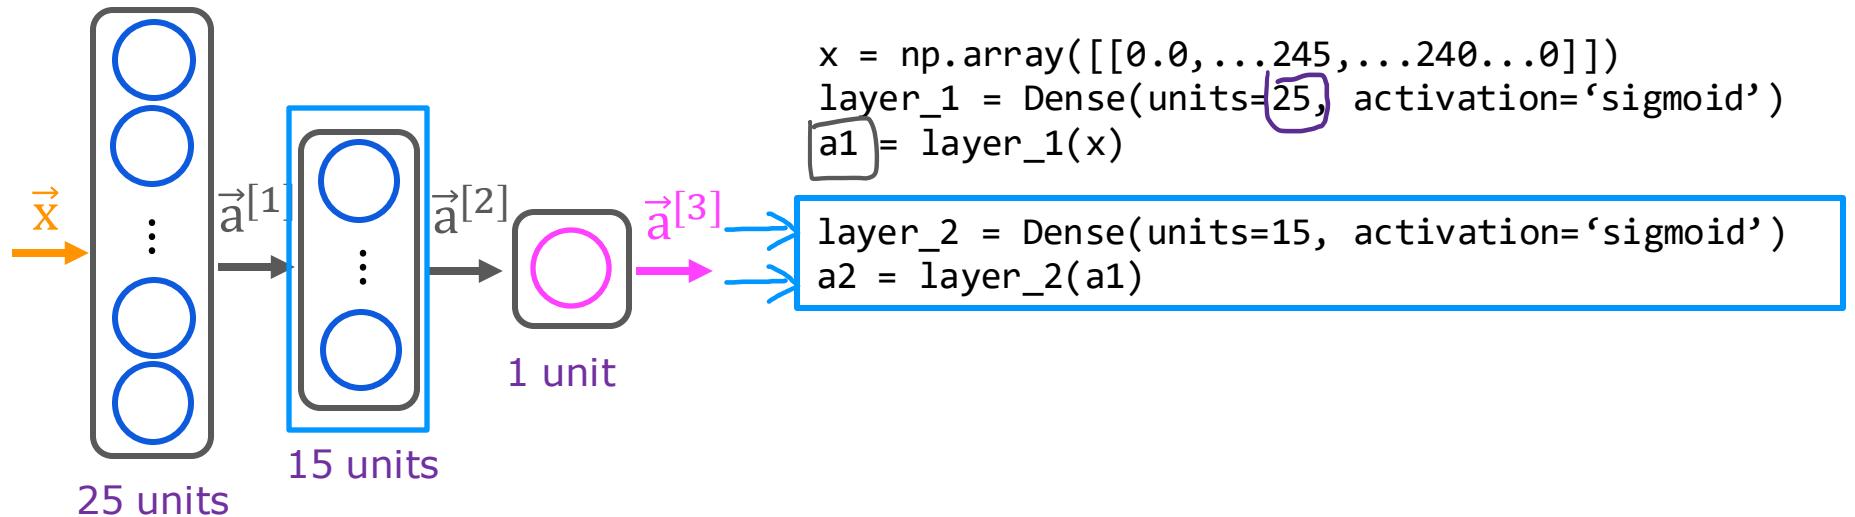
is $a_1^{[2]} \geq 0.5$?
yes $\hat{y} = 1$ no $\hat{y} = 0$

```
if a2 >= 0.5:
    yhat = 1
else:
    yhat = 0
```

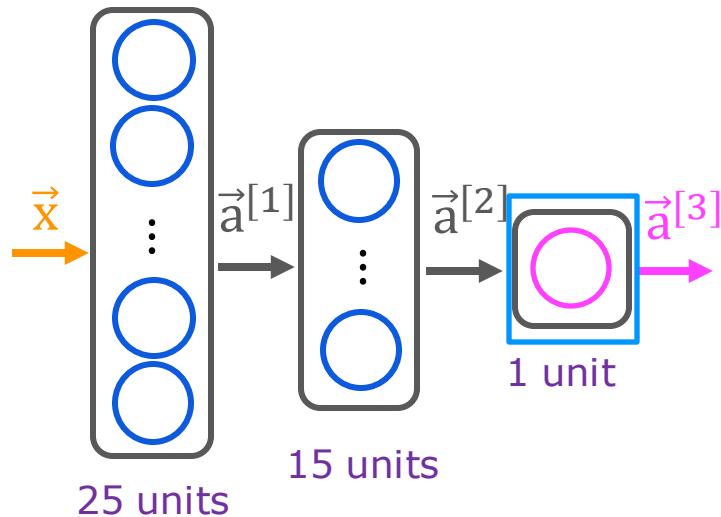
Model for digit classification



Model for digit classification



Model for digit classification



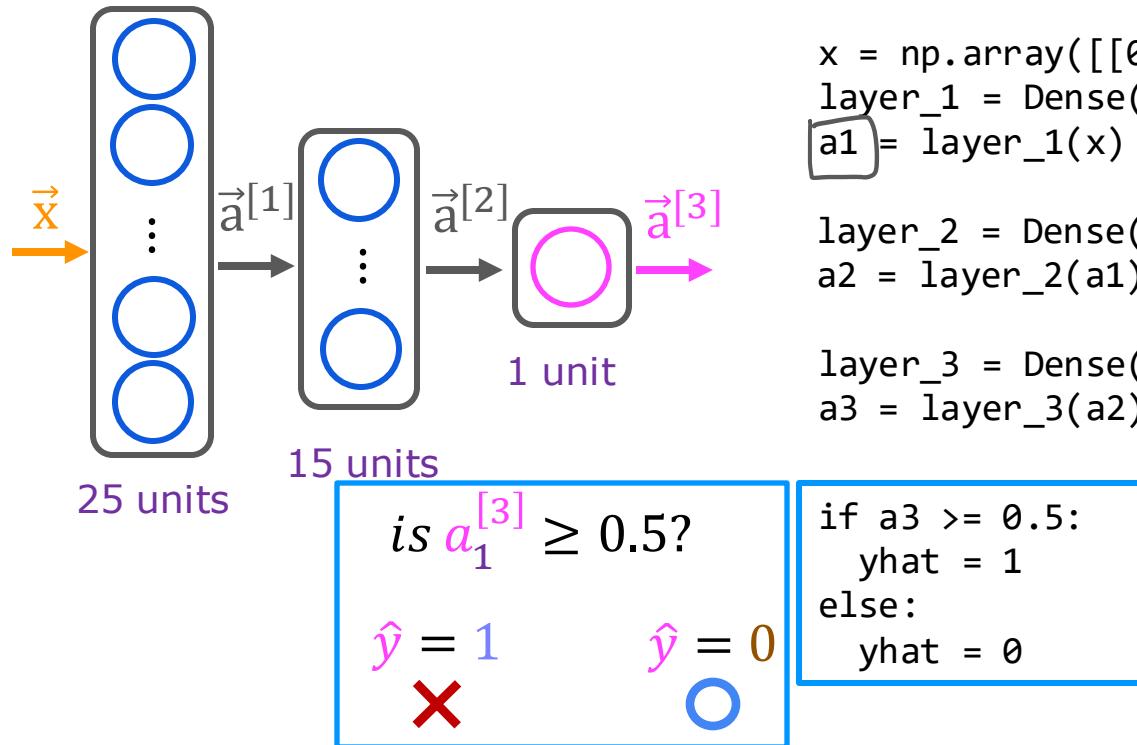
```
x = np.array([[0.0, ..., 245, ..., 240, ..., 0]])  
layer_1 = Dense(units=25, activation='sigmoid')  
a1 = layer_1(x)
```



```
layer_2 = Dense(units=15, activation='sigmoid')  
a2 = layer_2(a1)
```

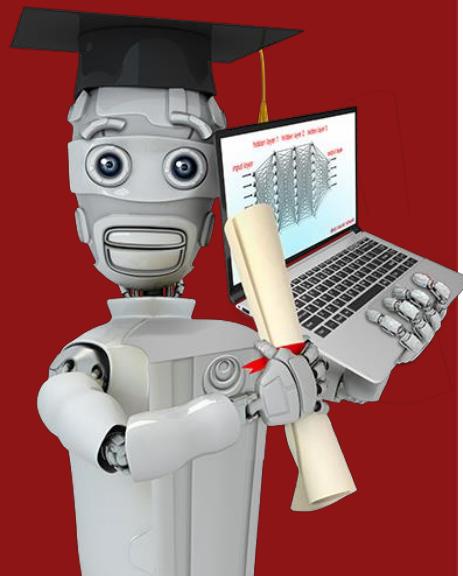
```
layer_3 = Dense(units=1, activation='sigmoid')  
a3 = layer_3(a2)
```

Model for digit classification



DeepLearning.AI

Stanford
ONLINE



TensorFlow implementation

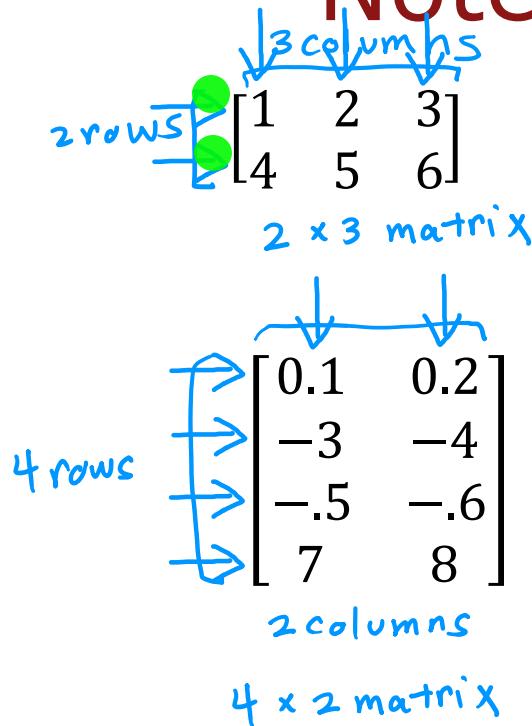
Data in TensorFlow

Feature vectors

temperature (Celsius)	duration (minutes)	Good coffee? (1/0)
200.0	17.0	1
425.0	18.5	0
...

$x = \text{np.array}([[200.0, 17.0]])$ ←
[[200.0, 17.0]]
why?

Note about numpy arrays



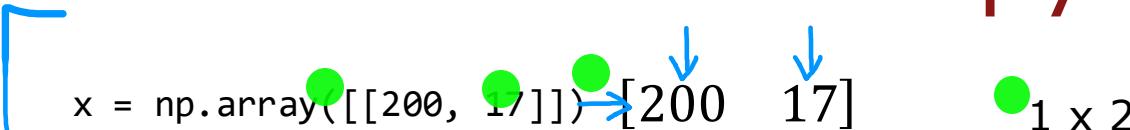
`x = np.array([[1, 2, 3],
[4, 5, 6]])` 2D array 2 x 3

`x = np.array([[0.1, 0.2],
[-3.0, -4.0],
[-0.5, -0.6],
[7.0, 8.0]])` 4 x 2 1 x 2 2 x 1

`[[0.1, 0.2],
[-3.0, -4.0],
[-0.5, -0.6],
[7.0, 8.0]]`



Note about numpy arrays

`x = np.array([[200, 17]])` 

`x = np.array([200, 17])` 

`x = np.array([200, 17])`

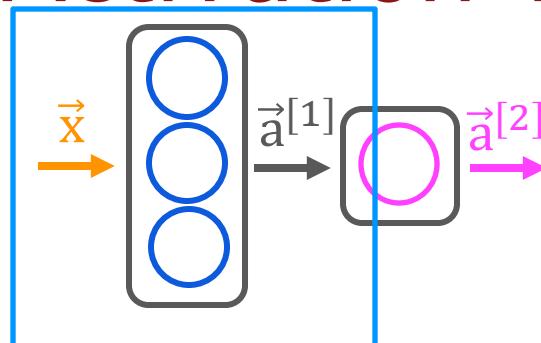
*1D
"Vector"*

Feature vectors

temperature (Celsius)	duration (minutes)	Good coffee? (1/0)
200.0	17.0	1
425.0	18.5	0
...

```
x = np.array([[200.0, 17.0]]) ←  
[[200.0, 17.0]]  
1 x 2  
→ [200.0 17.0]
```

Activation vector

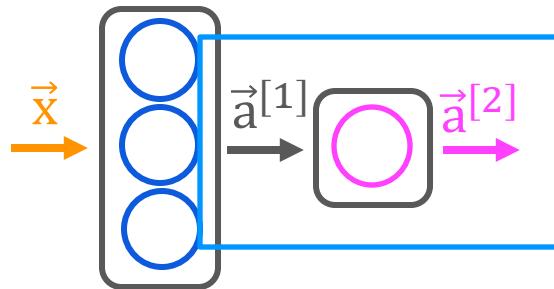


```
x = np.array([[200.0, 17.0]])
layer_1 = Dense(units=3, activation='sigmoid')
a1 = layer_1(x)
```

→ $\rightarrow [0.2, 0.7, 0.3]$ 1 x 3 matrix
→ $\rightarrow \text{tf.Tensor}([[0.2 0.7 0.3]], \text{shape}=(1, 3), \text{dtype}=\text{float32})$
→ $\rightarrow \text{a1.numpy()}$

```
array([[1.4661001, 1.125196 , 3.2159438]], dtype=float32)
```

Activation vector

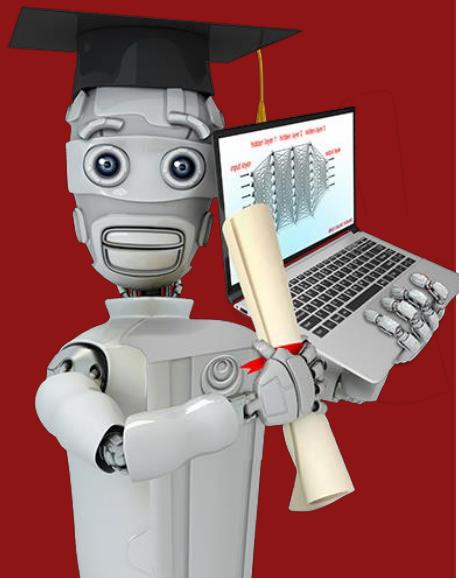


```
→ layer_2 = Dense(units=1, activation='sigmoid')
→ a2 = layer_2(a1)
    ↘ [[0.8]] ↙
→ tf.Tensor([[0.8]], shape=(1, 1), dtype=float32)
→ a2.numpy()
→ array([[0.8]], dtype=float32)
```

1×1

DeepLearning.AI

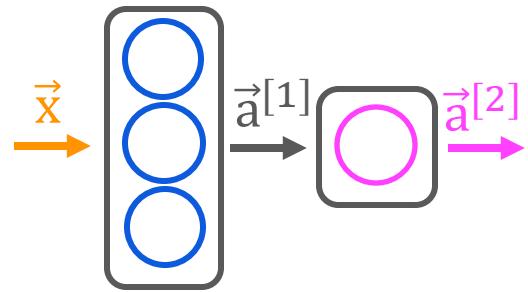
Stanford
ONLINE



TensorFlow implementation

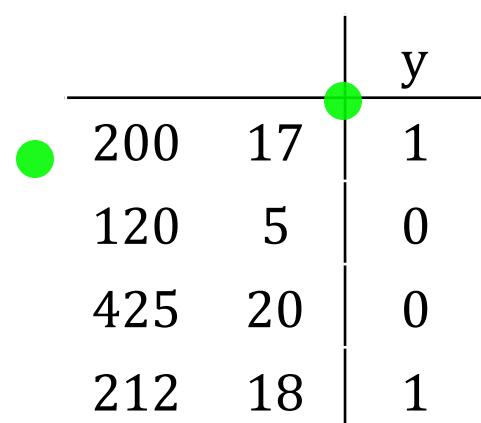
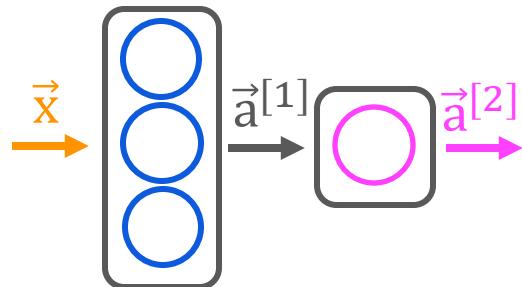
Building a neural network

What you saw earlier



```
→ x = np.array([[200.0, 17.0]])  
→ layer_1 = Dense(units=3, activation="sigmoid")  
→ a1 = layer_1(x)  
  
→ layer_2 = Dense(units=1, activation="sigmoid")  
→ a2 = layer_2(a1)
```

Building a neural network architecture



```
→ layer_1 = Dense(units=3, activation="sigmoid")  
→ layer_2 = Dense(units=1, activation="sigmoid")  
→ model = Sequential([layer_1, layer_2])
```

```
x = np.array([[200.0, 17.0],  
              [120.0, 5.0],  
              [425.0, 20.0],  
              [212.0, 18.0]])  
4 x 2
```

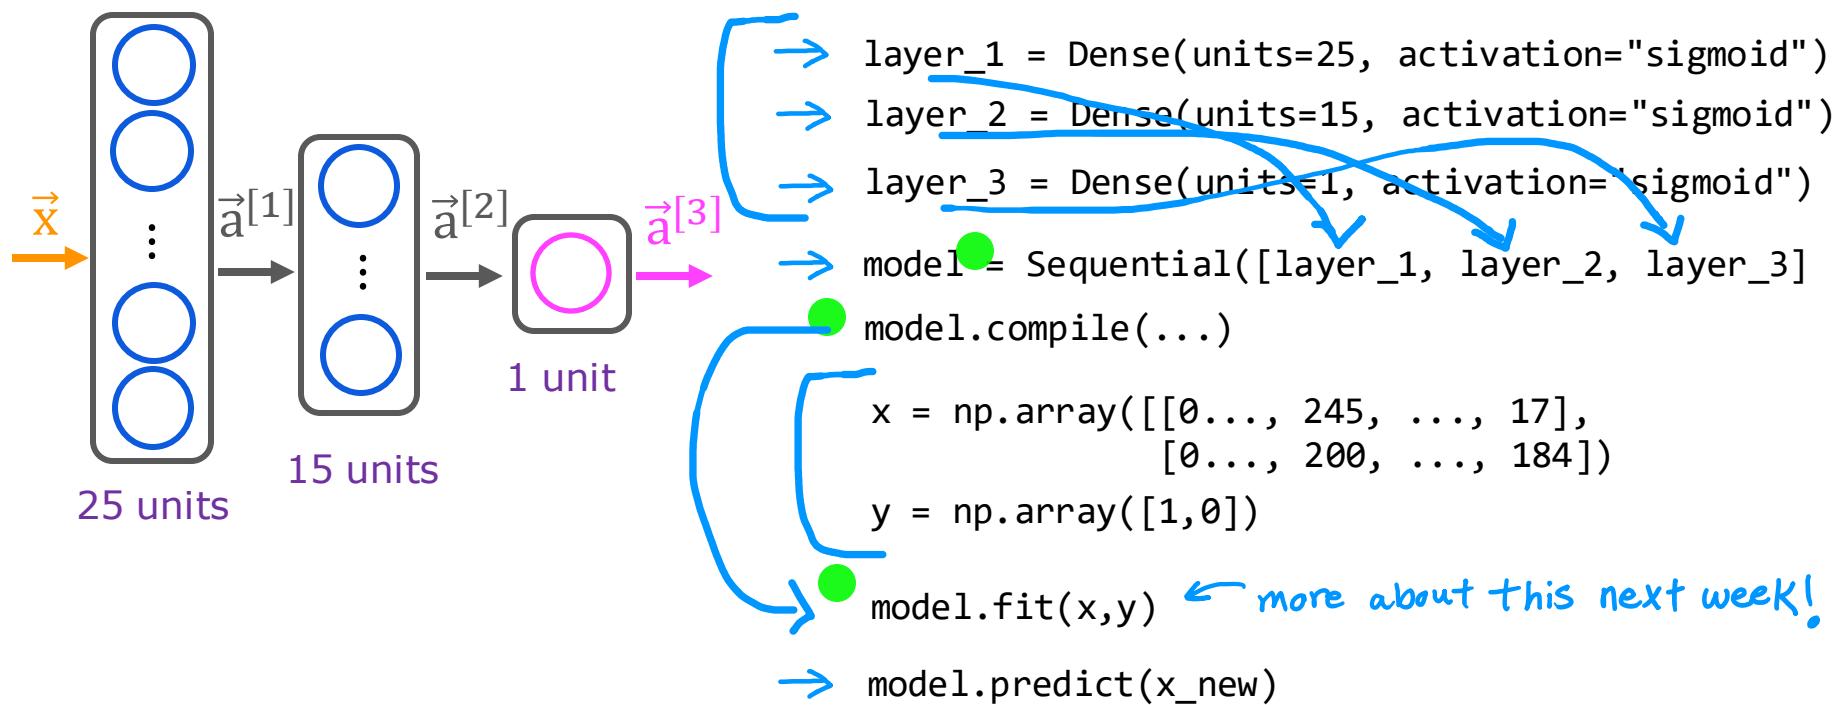
targets $y = \text{np.array}([1,0,0,1])$

```
model.compile(...)  
model.fit(x,y)
```

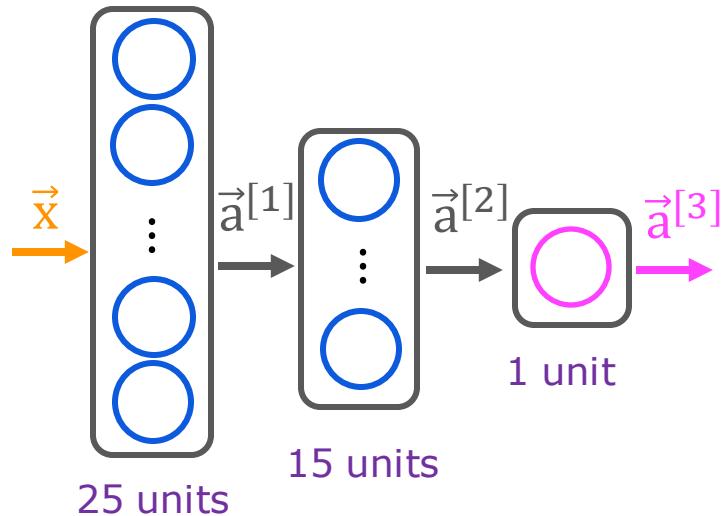
```
→ model.predict(x_new)
```

more about this next week!

Digit classification model



Digit classification model



```
model = Sequential([
    Dense(units=25, activation="sigmoid"),
    Dense(units=15, activation="sigmoid"),
    Dense(units=1, activation="sigmoid")])

model.compile(...)

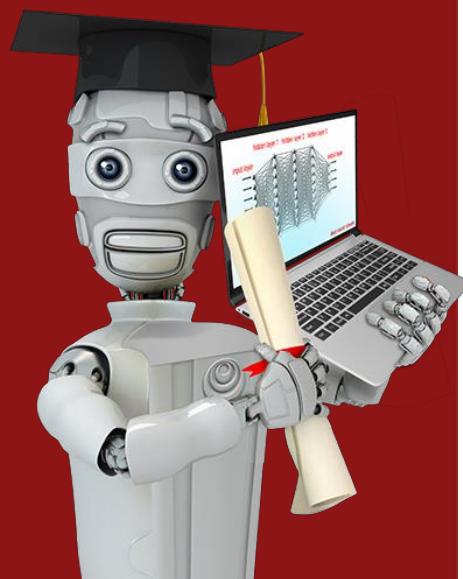
x = np.array([[0..., 245, ..., 17],
              [0..., 200, ..., 184]])
y = np.array([1,0])

model.fit(x,y)

model.predict(x_new)
```

DeepLearning.AI

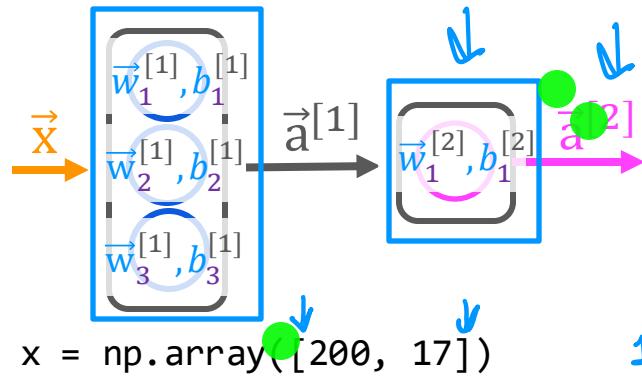
Stanford
ONLINE



Neural network implementation in Python

Forward prop in a single layer

forward prop (coffee roasting model)



```
x = np.array([200, 17])
```

$$a_1^{[1]} = g(\vec{w}_1^{[1]} \cdot \vec{x} + b_1^{[1]})$$

1D arrays

$$a_1^{[1]} = g(\vec{w}_2^{[1]} \cdot \vec{x} + b_2^{[1]})$$

$$a_1^{[2]} = g(\vec{w}_3^{[1]} \cdot \vec{x} + b_3^{[1]})$$

$$a_1^{[2]} = g(\vec{w}_1^{[2]} \cdot \vec{a}_1^{[1]} + b_1^{[2]})$$

$$w2_1 = np.array([-7, 8])$$

$$b2_1 = np.array([3])$$

$$z2_1 = np.dot(w2_1, a1) + b2_1$$

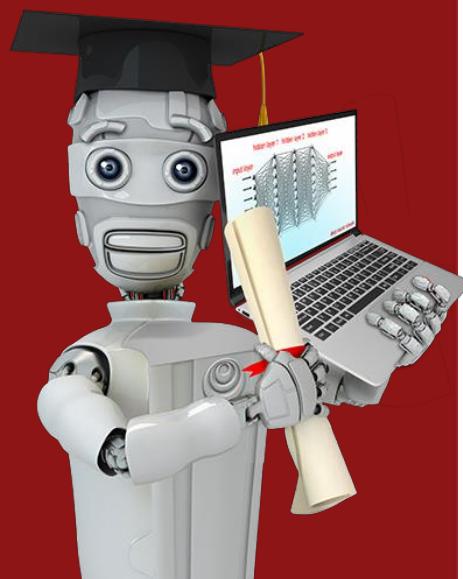
$$a2_1 = \text{sigmoid}(z1_1)$$

$$w_1^{[2]} \quad w_2_1$$

$$\begin{aligned} w1_1 &= np.array([1, 2]) & w1_2 &= np.array([-3, 4]) & w1_3 &= np.array([5, -6]) \\ b1_1 &= np.array([-1]) & b1_2 &= np.array([1]) & b1_3 &= np.array([2]) \\ z1_1 &= np.dot(w1_1, x) + b & z1_2 &= np.dot(w1_2, x) + b & z1_3 &= np.dot(w1_3, x) + b \\ a1_1 &= \text{sigmoid}(z1_1) & a1_2 &= \text{sigmoid}(z1_2) & a1_3 &= \text{sigmoid}(z1_3) \\ a1 &= np.array([a1_1, a1_2, a1_3]) \end{aligned}$$

DeepLearning.AI

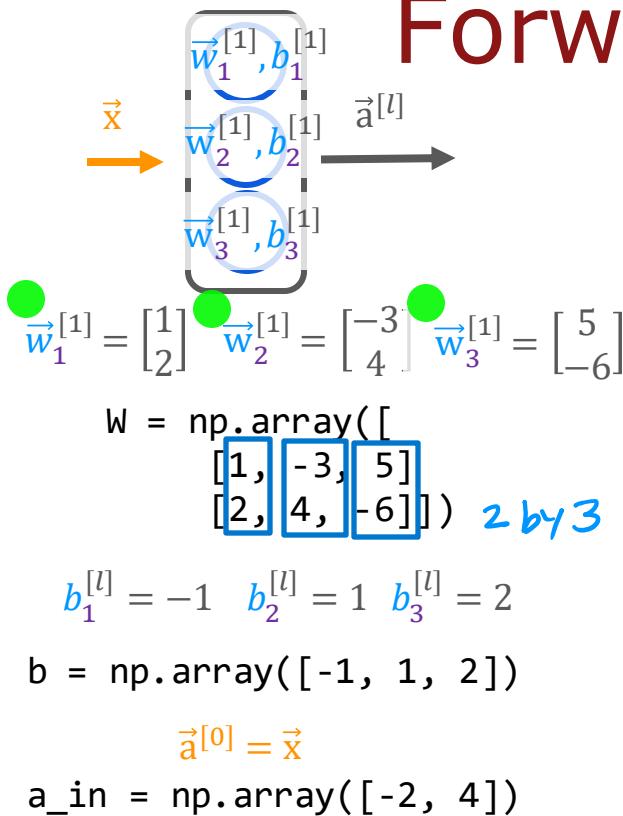
Stanford
ONLINE



Neural network implementation in Python

**General implementation of
forward propagation**

Forward prop in NumPy



```

def dense(a_in, W, b, g):
    units = W.shape[1] [0,0,0]
    a_out = np.zeros(units)
    for j in range(units): 0,1,2
        w = W[:,j]
        z = np.dot(w, a_in) + b[j]
        a_out[j] = g(z)
    return a_out
  
```

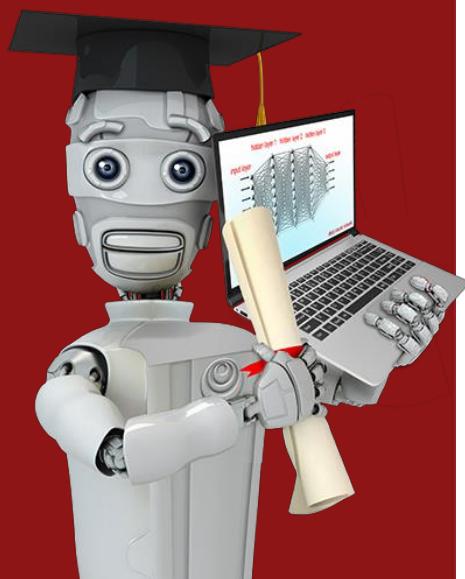
```

def sequential(x):
    a1 = dense(x, W1, b1, g)
    a2 = dense(a1, W2, b2, g)
    a3 = dense(a2, W3, b3, g)
    a4 = dense(a3, W4, b4, g)
    f_x = a4
    return f_x
  
```

capital W refers to a matrix

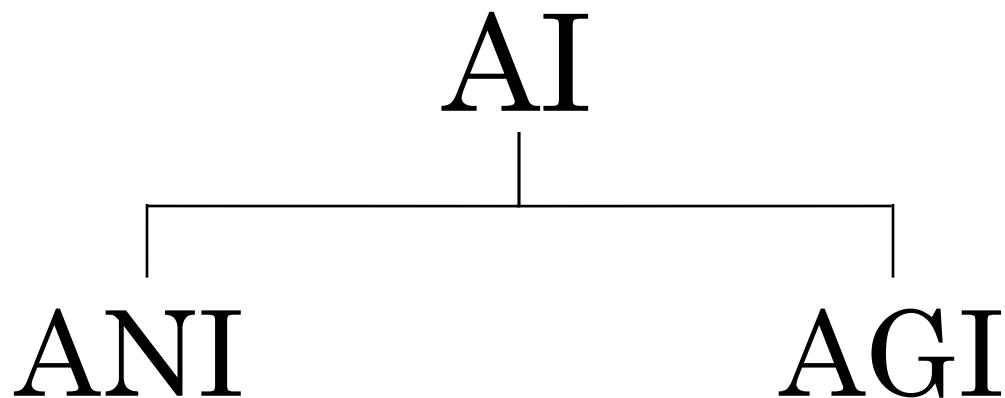
DeepLearning.AI

Stanford
ONLINE



Speculations on artificial general intelligence (AGI)

Is there a path to AGI?



(artificial narrow intelligence)

E.g., smart speaker,
self-driving car, web search,
AI in farming and factories

(artificial general intelligence)

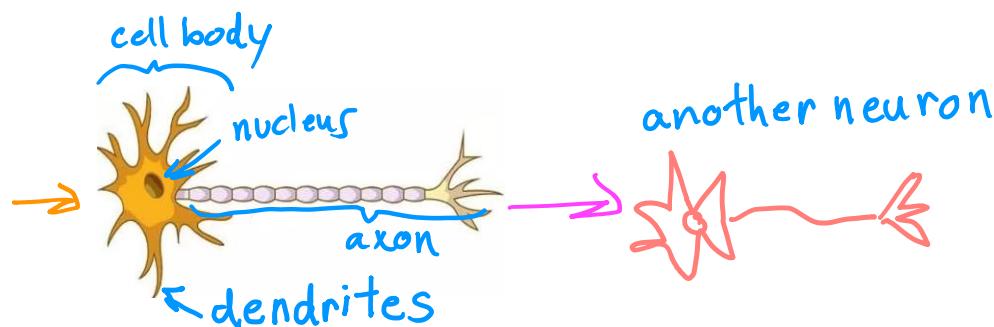
Do anything a human can do



Biological neuron

inputs

outputs



Simplified mathematical model of a neuron

inputs

outputs

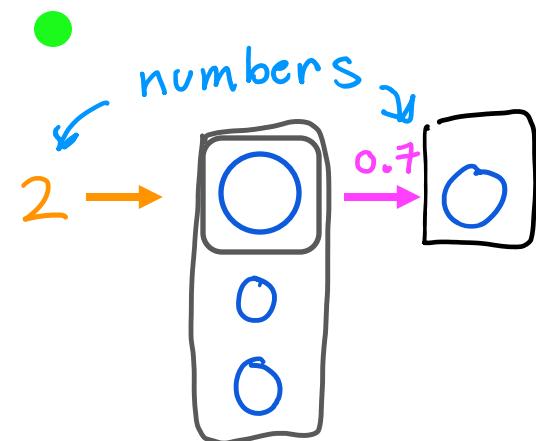
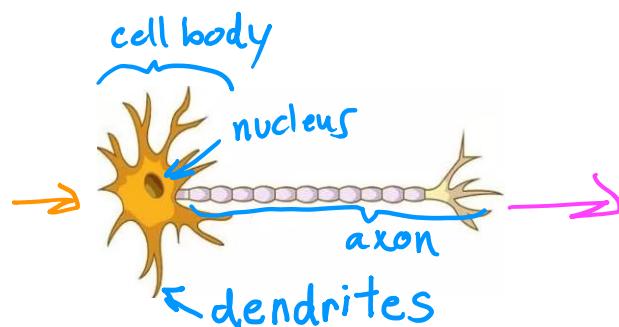


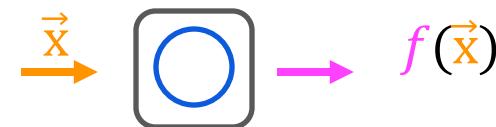
image source: <https://biologydictionary.net/sensory-neuron/>

Neural network and the brain

Can we mimic the human brain?

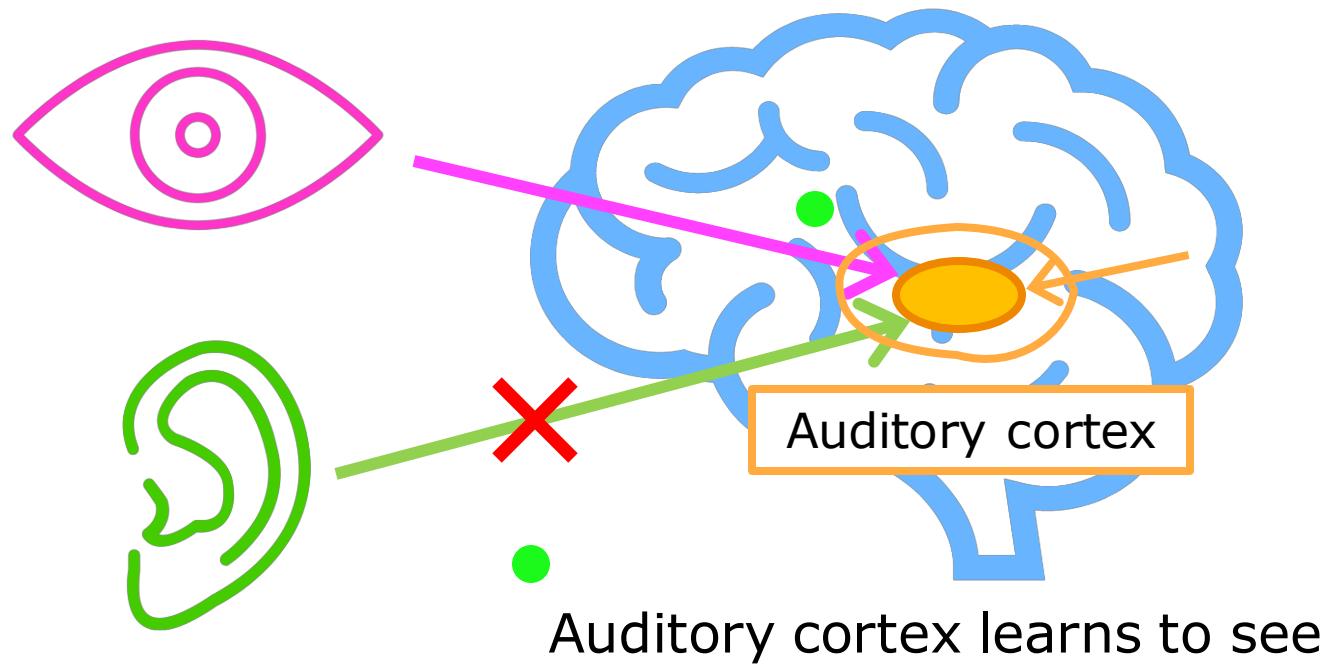


vs



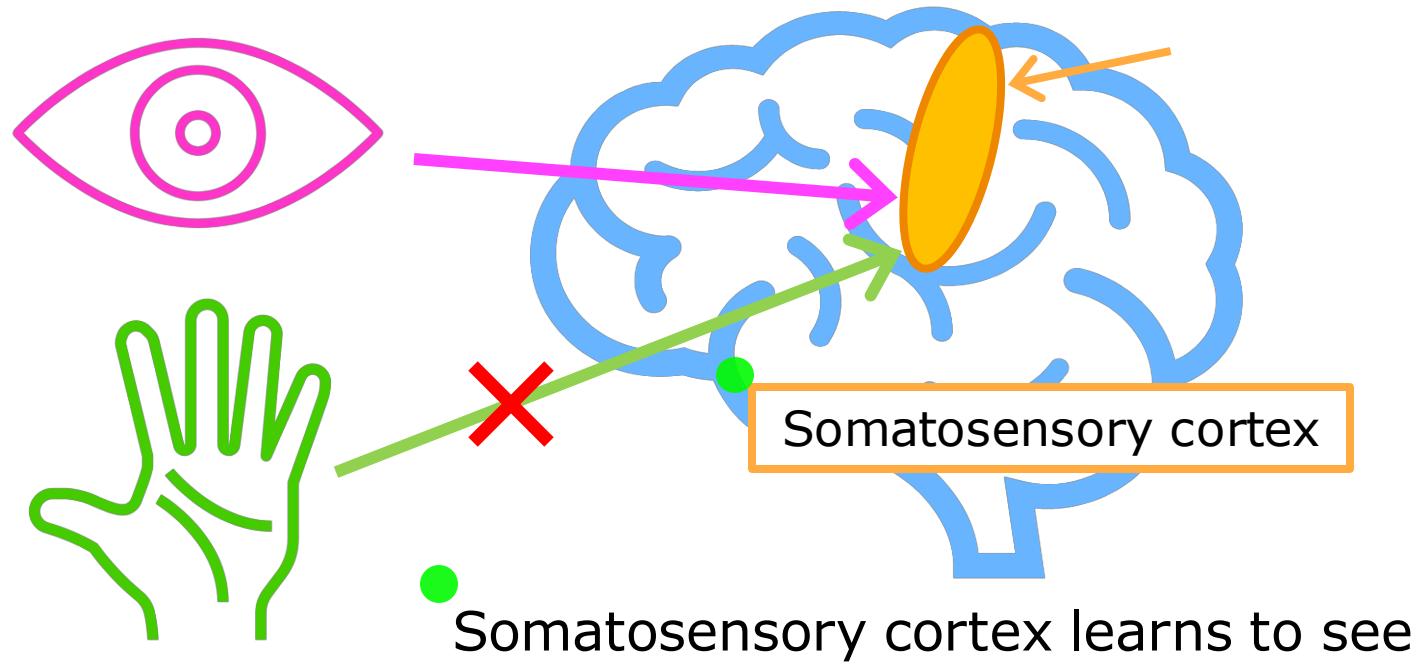
We have (almost) no idea how the brain works

The “one learning algorithm” hypothesis



[Roe et al., 1992]

The “one learning algorithm” hypothesis



[Metin & Frost, 1989]

Sensor representations in the brain



Seeing with your tongue

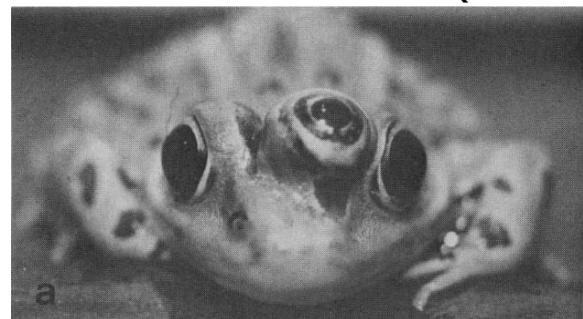


Haptic belt: Direction sense

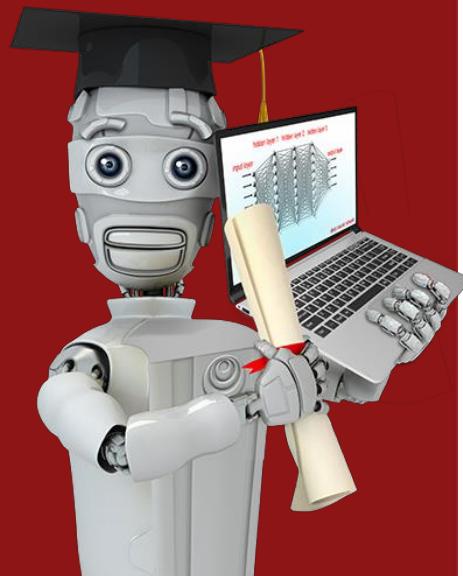
[BrainPort; Welsh & Blasch, 1997; Nagel et al., 2005; Constantine-Paton & Law, 2009]



Human echolocation (sonar)



Implanting a 3rd eye



Vectorization (optional)

**How neural networks are
implemented efficiently**

For loops vs. vectorization

```
x = np.array([200, 17])  
W = np.array([[1, -3, 5],  
             [-2, 4, -6]])  
b = np.array([-1, 1, 2])  
  
def dense(a_in,W,b):  
    a_out = np.zeros(units)  
    for j in range(units):  
        w = W[:,j]  
        z = np.dot(w,x) + b[j]  
        a[j] = g(z)  
    return a
```

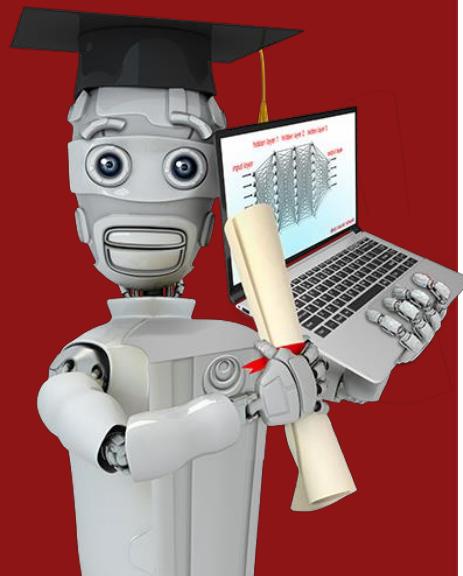
vectorized

```
x = np.array([[200, 17]]) 2D array  
W = np.array([[1, -3, 5],  
             [-2, 4, -6]]) same  
B = np.array([[-1, 1, 2]]) 1x3 2D array  
  
def dense(A_in,W,B): all 2D arrays  
    Z = np.matmul(A_in,W) + B  
    A_out = g(Z) matrix multiplication  
    return A_out  
  
[[1,0,1]]
```

[1,0,1]

DeepLearning.AI

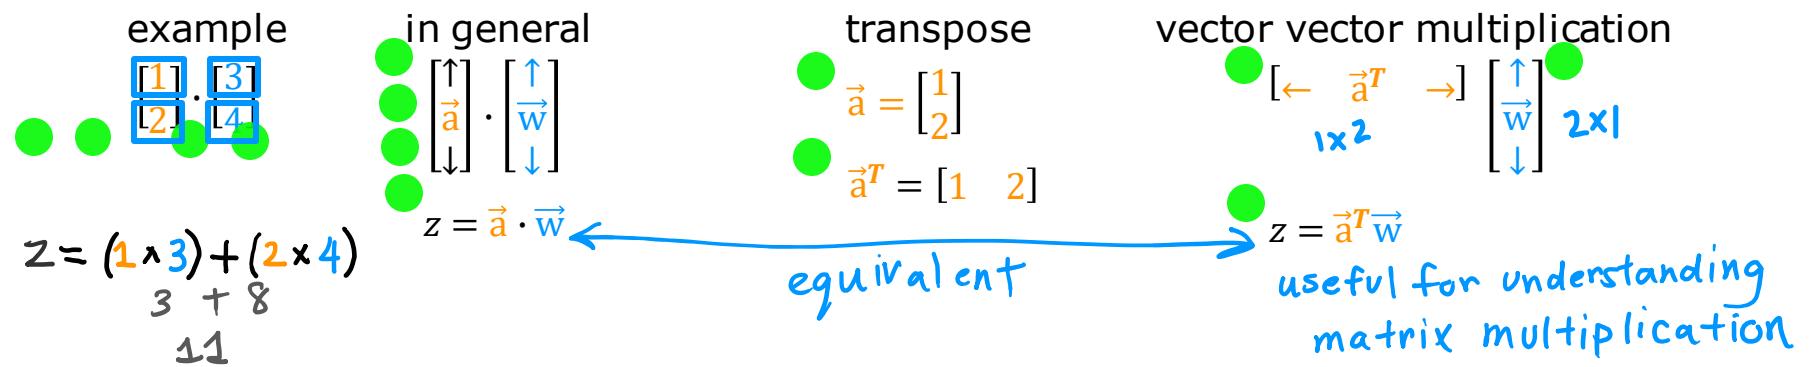
Stanford
ONLINE



Vectorization (optional)

Matrix multiplication

Dot products



Vector matrix multiplication

$$\mathbf{a} = \begin{bmatrix} 1 \\ 2 \end{bmatrix}$$
$$\mathbf{a} = \begin{bmatrix} 1 & 2 \end{bmatrix} = \begin{bmatrix} 3 & 5 \\ 4 & 6 \end{bmatrix} = \mathbf{a} \begin{bmatrix} \uparrow & \uparrow \\ \mathbf{w}_1 & \mathbf{w}_2 \\ \downarrow & \downarrow \end{bmatrix}$$

1 by 2

$$\begin{bmatrix} 1 & 3 \\ 2 & 4 \end{bmatrix} + \begin{bmatrix} 1 & 5 \\ 2 & 6 \end{bmatrix} = \begin{bmatrix} \mathbf{a} & \mathbf{w}_1 & \mathbf{w}_2 \end{bmatrix}$$
$$\begin{bmatrix} 3+8 \\ 11 \end{bmatrix} + \begin{bmatrix} 5+12 \\ 17 \end{bmatrix} = \begin{bmatrix} 11 & 17 \end{bmatrix}$$

matrix matrix multiplication

$$\begin{aligned} &= \begin{bmatrix} 1 & 1 \\ 2 & 2 \end{bmatrix} \\ &= \begin{bmatrix} 1 & 2 \\ 1 & 2 \end{bmatrix} \quad = \quad \begin{bmatrix} 3 & 5 \\ 4 & 6 \end{bmatrix} \\ &\text{rows} \quad \text{Columns} \end{aligned}$$
$$\begin{bmatrix} a_1 & a_2 \\ & \end{bmatrix} \quad \begin{bmatrix} \uparrow & \uparrow \\ w_1 & w_2 \\ \downarrow & \downarrow \end{bmatrix}$$

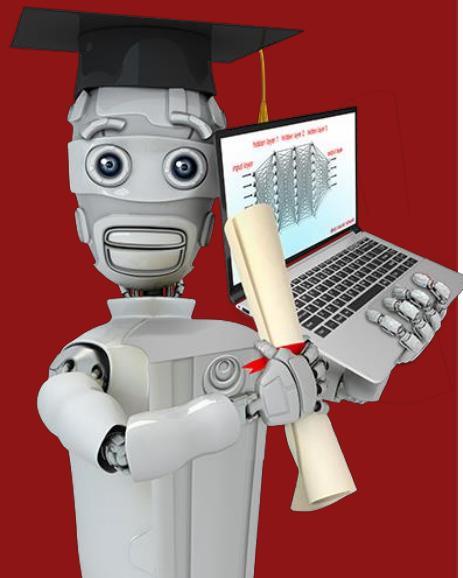
$$\begin{aligned} &\text{row1 col1} \quad \text{row1 col2} \\ &\text{row2 col1} \quad \text{row2 col2} \\ &(-1 \times 3) + (-2 \times 4) \quad (-1 \times 5) + (-2 \times 6) \\ &\quad -3 + -8 \quad -5 + -12 \\ &\quad -11 \quad -17 \\ &= \begin{bmatrix} 11 & 17 \\ 11 & 17 \end{bmatrix} \end{aligned}$$

general rules for
matrix multiplication
↳ next video!



DeepLearning.AI

Stanford
ONLINE



Vectorization (optional)

Matrix multiplication rules

Matrix multiplication rules

$$A = \begin{bmatrix} 1 & -1 & 0.1 \\ 2 & -2 & 0.2 \end{bmatrix} \quad A^T = \begin{bmatrix} 1 & 2 \\ -1 & -2 \\ 0.1 & 0.2 \end{bmatrix} \quad \vec{a}_1^T \quad \vec{a}_2^T \quad \vec{a}_3^T$$
$$W = \begin{bmatrix} 3 & 5 & 7 & 9 \\ 4 & 6 & 8 & 0 \end{bmatrix} \quad \vec{w}_1 \quad \vec{w}_2 \quad \vec{w}_3 \quad \vec{w}_4$$
$$Z = A^T W = \begin{bmatrix} & & & \\ & & & \\ & & & \\ & & & \end{bmatrix}$$



Matrix multiplication rules

$$A = \begin{bmatrix} 1 & -1 & 0.1 \\ 2 & -2 & 0.2 \end{bmatrix} \quad A^T = \begin{bmatrix} 1 & -1 & 0.1 \\ -1 & -2 & 0.2 \\ 0.1 & 0.2 & 1 \end{bmatrix} \quad W = \begin{bmatrix} 3 & 5 & 7 & 9 \\ 4 & 6 & 8 & 0 \end{bmatrix} \quad Z = A^T W = \begin{bmatrix} 1 & -1 & 0.1 & 1 \\ -1 & -2 & 0.2 & 1 \\ 0.1 & 0.2 & 1 & 1 \end{bmatrix}$$

$$\vec{a}_1^T \vec{w}_1 = (1 \times 3) + (2 \times 4) = 11$$

3 by 4 matrix

row 3 column 2

$$\vec{a}_3^T \vec{w}_2 = (0.1 \times 5) + (0.2 \times 6) = 1.7$$

0.5 + 1.2

row 2 column 3?

$$\vec{a}_2^T \vec{w}_3 = (-1 \times 7) + (-2 \times 8) = -23$$

-7 + -16



Matrix multiplication rules

$$A = \begin{bmatrix} 1 & -1 & 0.1 \\ 2 & -2 & 0.2 \end{bmatrix} \quad A^T = \begin{bmatrix} 1 & -1 & 0.1 \\ -1 & -2 & 0.2 \\ 0.1 & 0.2 & \end{bmatrix} \quad W = \begin{bmatrix} 3 & 5 & 7 & 9 \\ 4 & 6 & 8 & 0 \end{bmatrix} \quad Z = A^T W = \begin{bmatrix} 11 & 17 & 23 & 9 \\ -11 & -17 & -23 & -9 \\ 1.1 & 1.7 & 2.3 & 0.9 \end{bmatrix}$$

$$\vec{a}_1^T \vec{w}_1 = (1 \times 3) + (2 \times 4) = 11$$

3 by 4 matrix

row 3 column 2

$$\vec{a}_3^T \vec{w}_2 = (0.1 \times 5) + (0.2 \times 6) = 1.7$$

0.5 + 1.2

row 2 column 3?

$$\vec{a}_2^T \vec{w}_3 = (-1 \times 7) + (-2 \times 8) = -23$$

-7 + -16



Matrix multiplication rules

$$A = \begin{bmatrix} 1 & -1 & 0.1 \\ 2 & -2 & 0.2 \end{bmatrix} \quad A^T = \begin{bmatrix} 1 & -1 & 0.1 \\ -1 & -2 & 0.2 \\ 0.1 & 0.2 & 1 \end{bmatrix} \quad W = \begin{bmatrix} 3 & 5 & 7 & 9 \\ 4 & 6 & 8 & 0 \end{bmatrix}$$

3×2 2×4
can only take dot products
of vectors that are same length

$$\begin{bmatrix} 0.1 & 0.2 \end{bmatrix} \quad \begin{bmatrix} 5 \\ 6 \end{bmatrix}$$

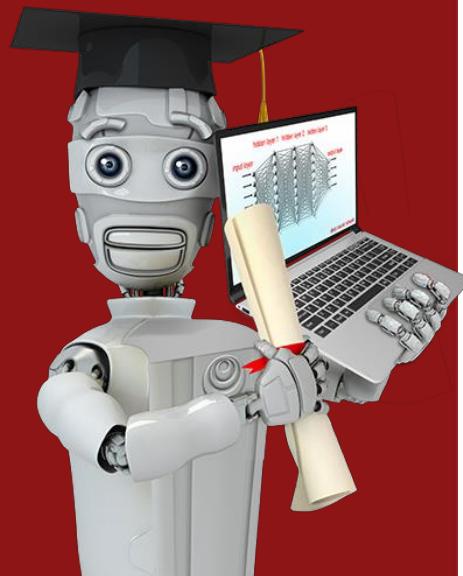
length 2 length 2

$$Z = A^T W = \begin{bmatrix} 11 & 17 & 23 & 9 \\ -11 & -17 & -23 & -9 \\ 1.1 & 1.7 & 2.3 & 0.9 \end{bmatrix}$$

3 by 4 matrix
↳ same # rows as A^T
↳ same # columns as W

DeepLearning.AI

Stanford
ONLINE



Vectorization (optional)

Matrix multiplication code

Matrix multiplication in NumPy

$$A = \begin{bmatrix} 1 & -1 & 0.1 \\ 2 & -2 & 0.2 \end{bmatrix} \quad A^T = \begin{bmatrix} 1 & -1 & 0.1 \\ -1 & -2 & 0.2 \\ 0.1 & 0.2 & 1 \end{bmatrix} \quad W = \begin{bmatrix} 3 & 5 & 7 & 9 \\ 4 & 6 & 8 & 0 \end{bmatrix} \quad Z = A^T W = \begin{bmatrix} 11 & 17 & 23 & 9 \\ -11 & -17 & -23 & -9 \\ 1.1 & 1.7 & 2.3 & 0.9 \end{bmatrix}$$

```
A=np.array([1,-1,0.1],  
           [2,-2,0.2]))
```

```
W=np.array([3,5,7,9],  
           [4,6,8,0])
```

or
Z = AT @ W

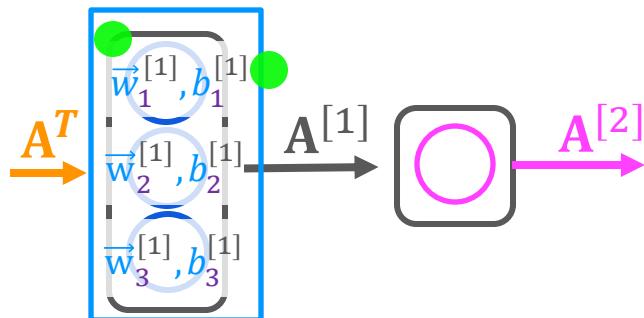
```
AT=np.array([1,2],  
           [-1,-2],  
           [0.1,0.2])
```

AT=A.T
transpose

result
[[11,17,23,9],
 [-11,-17,-23,-9],
 [1.1,1.7,2.3,0.9]]



Dense layer vectorized



$$A^T = [200 \quad 17]$$

$$W = \begin{bmatrix} 1 & -3 & 5 \\ -2 & 4 & -6 \end{bmatrix}$$

$$\vec{b} = [-1 \quad 1 \quad 2]$$

$$Z = A^T W + b$$

$$\begin{bmatrix} 165 \\ -531 \\ 900 \end{bmatrix} \quad z_1^{[1]} \quad z_2^{[1]} \quad z_3^{[1]}$$

$$A = g(Z)$$

$$\begin{bmatrix} 1 & 0 & 1 \end{bmatrix}$$

A

$AT = np.array([[200, 17]])$

$W = np.array([[1, -3, 5], [-2, 4, -6]])$

$b = np.array([-1, 1, 2])$

a_in

```
def dense(AT,W,b,g):
    z = np.matmul(AT,W) + b
    a_out = g(z)
    return a_out
```

$[[1,0,1]]$