

[PyTorch Operations] (CheatSheet)

1. Tensor Basics

- **Creating a Tensor:** `torch.tensor([1, 2, 3])`
- **Zeros Tensor:** `torch.zeros(3, 3)`
- **Ones Tensor:** `torch.ones(3, 3)`
- **Random Tensor:** `torch.rand(3, 3)`
- **Arange:** `torch.arange(0, 10)`
- **Linspace:** `torch.linspace(0, 10, steps=100)`
- **Eye (Identity Matrix):** `torch.eye(3)`
- **Tensor Shape:** `tensor.size()`
- **Reshaping a Tensor:** `tensor.view(3, -1)`
- **Flattening a Tensor:** `tensor.view(-1)`
- **Squeeze a Tensor (Remove Single-Dimensional Entries):** `tensor.squeeze()`
- **Unsqueeze (Add a Dimension):** `tensor.unsqueeze(dim=0)`
- **Tensor to Numpy Array:** `tensor.numpy()`
- **Numpy Array to Tensor:** `torch.from_numpy(numpy_array)`
- **Tensor Data Type:** `tensor.dtype`
- **Tensor to a Specified Device:** `tensor.to(device)`
- **Copying a Tensor:** `tensor.clone()`
- **Concatenating Tensors:** `torch.cat([tensor1, tensor2], dim=0)`
- **Stacking Tensors:** `torch.stack([tensor1, tensor2])`
- **Chunking a Tensor:** `torch.chunk(tensor, chunks=3, dim=0)`
- **Splitting a Tensor:** `torch.split(tensor, split_size_or_sections=3)`
- **Tiling a Tensor:** `tensor.repeat(3, 2, 1)`
- **Permute Tensor Dimensions:** `tensor.permute(2, 0, 1)`
- **Transpose a Tensor:** `tensor.t()`
- **Specific Element Indexing:** `tensor[0, 1]`
- **Slicing a Tensor:** `tensor[:, 1]`
- **Conditional Selection in Tensor:** `tensor[tensor > 0]`
- **Applying a Function Elementwise:** `torch.apply_along_axis(func, dim, tensor)`
- **Filling Tensor with a Value:** `tensor.fill_(5)`
- **In-Place Operations (Postfix with an Underscore):** `tensor.add_(5)`

2. Mathematical Operations

- **Addition:** `torch.add(tensor1, tensor2)`

- **Subtraction:** `tensor1 - tensor2`
- **Elementwise Multiplication:** `torch.mul(tensor1, tensor2)`
- **Elementwise Division:** `torch.div(tensor1, tensor2)`
- **Matrix Multiplication:** `torch.matmul(tensor1, tensor2)`
- **Dot Product:** `torch.dot(tensor1, tensor2)`
- **Power:** `torch.pow(tensor, exponent)`
- **Square Root:** `torch.sqrt(tensor)`
- **Logarithm:** `torch.log(tensor)`
- **Exponential:** `torch.exp(tensor)`
- **Summation:** `torch.sum(tensor)`
- **Product:** `torch.prod(tensor)`
- **Mean:** `torch.mean(tensor)`
- **Median:** `torch.median(tensor)`
- **Standard Deviation:** `torch.std(tensor)`
- **Variance:** `torch.var(tensor)`
- **Max:** `torch.max(tensor)`
- **Min:** `torch.min(tensor)`
- **Absolute Value:** `torch.abs(tensor)`
- **Clamping Values:** `torch.clamp(tensor, min=-1, max=1)`

3. Linear Algebra

- **Eigenvalues and Eigenvectors:** `torch.eig(tensor, eigenvectors=True)`
- **Determinant:** `torch.det(tensor)`
- **Inverse of Matrix:** `torch.inverse(tensor)`
- **Singular Value Decomposition:** `torch.svd(tensor)`
- **QR Decomposition:** `torch.qr(tensor)`
- **Cholesky Decomposition:** `torch.cholesky(tensor)`
- **Cross Product:** `torch.cross(tensor1, tensor2)`

4. Advanced Tensor Operations

- **Diagonal Elements:** `torch.diag(tensor)`
- **Trace (Sum of Diagonal Elements):** `torch.trace(tensor)`
- **Rank of a Tensor:** `torch.matrix_rank(tensor)`
- **Kronecker Product:** `torch.kron(tensor1, tensor2)`
- **Flatten a Tensor:** `tensor.flatten()`
- **Tensor Unfolding (Matricization):** `tensor.unfold(dimension, size, step)`
- **Tensor Normalization:** `torch.nn.functional.normalize(tensor, p=2, dim=1)`

- **Finding Unique Elements:** `torch.unique(tensor)`
- **Sorting a Tensor:** `torch.sort(tensor)`
- **Top-k Elements:** `torch.topk(tensor, k=5)`
- **Comparing Two Tensors:** `torch.equal(tensor1, tensor2)`
- **Where Operation:** `torch.where(condition, tensor1, tensor2)`
- **Index Select:** `torch.index_select(tensor, dim, index)`

5. Random and Probability

- **Setting the Seed for Randomness:** `torch.manual_seed(42)`
- **Random Sampling from Uniform Distribution:** `torch.rand(size)`
- **Random Sampling from Normal Distribution:** `torch.randn(size)`
- **Random Integer Generation:** `torch.randint(low=0, high=10, size=size)`
- **Shuffling Elements:** `tensor[torch.randperm(tensor.size(0))]`
- **Bernoulli Sampling:** `torch.bernoulli(tensor)`

6. Gradients and Autograd

- **Enabling Gradient Tracking:** `tensor.requires_grad_(True)`
- **Computing Gradients:** `tensor.backward()`
- **Accessing the Gradient:** `tensor.grad`
- **Disabling Gradient Tracking:** `with torch.no_grad():`
- **Detaching Gradients:** `tensor.detach()`
- **Zeroing Gradients:** `optimizer.zero_grad()`

7. Neural Network Building Blocks

- **Defining a Linear Layer:** `torch.nn.Linear(in_features, out_features)`
- **Convolutional Layer:** `torch.nn.Conv2d(in_channels, out_channels, kernel_size)`
- **Pooling Layer:** `torch.nn.MaxPool2d(kernel_size)`
- **Non-linear Activations (ReLU, Sigmoid, etc.):** `torch.nn.ReLU()`
- **Batch Normalization:** `torch.nn.BatchNorm2d(num_features)`
- **Dropout:** `torch.nn.Dropout(p=0.5)`
- **Sequential Container:** `torch.nn.Sequential(torch.nn.Linear(), torch.nn.ReLU())`

8. Loss Functions

- **Mean Squared Error Loss:** `torch.nn.MSELoss()`

- **Cross Entropy Loss:** `torch.nn.CrossEntropyLoss()`
- **Binary Cross Entropy Loss:** `torch.nn.BCELoss()`
- **L1 Loss (Absolute Error):** `torch.nn.L1Loss()`
- **Negative Log Likelihood Loss:** `torch.nn.NLLLoss()`

9. Optimizers

- **Stochastic Gradient Descent:** `torch.optim.SGD(model.parameters(), lr=0.01)`
- **Adam Optimizer:** `torch.optim.Adam(model.parameters(), lr=0.001)`
- **Adagrad Optimizer:** `torch.optim.Adagrad(model.parameters(), lr=0.01)`
- **RMSprop Optimizer:** `torch.optim.RMSprop(model.parameters(), lr=0.01)`
- **Learning Rate Scheduling:** `torch.optim.lr_scheduler.StepLR(optimizer, step_size=30, gamma=0.1)`

10. Data Preprocessing and Loaders

- **Tensor Dataset:** `torch.utils.data.TensorDataset(features, targets)`
- **Data Loader:** `torch.utils.data.DataLoader(dataset, batch_size=32, shuffle=True)`
- **Transforms for Data Augmentation:**
`torchvision.transforms.RandomRotation(degrees=30)`
- **Normalizing Data:** `torchvision.transforms.Normalize(mean, std)`

11. Model Training and Evaluation

- **Training Loop:** Loop through `DataLoader` and update weights with `optimizer`
- **Evaluation Loop:** Loop through `DataLoader` for validation or testing
- **Saving a Model:** `torch.save(model.state_dict(), 'model.pth')`
- **Loading a Model:** `model.load_state_dict(torch.load('model.pth'))`
- **Setting Model to Train Mode:** `model.train()`
- **Setting Model to Evaluation Mode:** `model.eval()`

12. GPU Utilization

- **Check GPU Availability:** `torch.cuda.is_available()`
- **Send Model to GPU:** `model.to('cuda')`
- **Send Data to GPU:** `tensor.to('cuda')`
- **Copying Data Back to CPU:** `tensor.to('cpu')`

13. Advanced Operations

- **Custom Autograd Function:** Define forward and backward functions
- **Using Multiple GPUs:** `torch.nn.DataParallel(model)`
- **Gradient Clipping:** `torch.nn.utils.clip_grad_norm_(model.parameters(), max_norm)`
- **Weight Initialization:** `torch.nn.init.xavier_uniform_(tensor)`

14. PyTorch Extensions

- **Using PyTorch Lightning for Training Abstraction:** Utilize `pytorch_lightning.LightningModule`
- **Using Torchvision for Pretrained Models:** Load models with `torchvision.models`
- **TorchText for NLP:** Preprocessing and loading text data
- **TorchAudio for Audio Processing:** Working with audio data

15. Debugging and Profiling

- **PyTorch Profiler:** with `torch.profiler.profile()` as `prof`:
- **Inspecting Tensor Values:** Print or log tensor values during debugging
- **Checking Model Summary:** Use `torchsummary` for model architecture and parameters

16. Visualization and Interpretation

- **TensorBoard Integration:** `from torch.utils.tensorboard import SummaryWriter`
- **Visualizing Model Graphs:** Use `'torchviz'` for visualizing computational graphs
- **Feature Map Visualization:** Extract and visualize intermediate layers
- **Activation Visualization:** Plotting activations of specific layers
- **Weights and Gradients Visualization:** Monitor weights and gradients during training

17. Working with Sequences and Time Series

- **Recurrent Neural Networks (RNN):** `torch.nn.RNN(input_size, hidden_size)`
- **Long Short-Term Memory (LSTM):** `torch.nn.LSTM(input_size, hidden_size)`
- **Gated Recurrent Units (GRU):** `torch.nn.GRU(input_size, hidden_size)`
- **Sequence Padding:** `torch.nn.utils.rnn.pad_sequence(sequences)`

- **Packing Padded Sequences:** `torch.nn.utils.rnn.pack_padded_sequence(input, lengths)`

18. Custom Layers and Models

- **Defining Custom Layer:** Create a class inheriting from `torch.nn.Module`
- **Writing Forward Pass:** Define `forward` method for custom layers
- **Composite Model Construction:** Combining multiple layers into a single model

19. Advanced Techniques

- **Attention Mechanisms:** Implement attention in sequence models
- **Generative Adversarial Networks (GANs):** Building generator and discriminator models
- **Transfer Learning:** Fine-tuning pretrained models
- **Multi-Task Learning:** Shared representations for multiple tasks

20. Best Practices and Tips

- **Weight Decay for Regularization:** Using `weight_decay` in optimizers
- **Batch Normalization Tuning:** Properly placing `BatchNorm` layers
- **Avoiding Overfitting:** Techniques like dropout, data augmentation
- **Hyperparameter Tuning:** Experimenting with learning rates, batch sizes, etc.

21. Miscellaneous Operations

- **Image to Tensor Conversion:** `transforms.ToTensor()`
- **Gradient Accumulation:** Summing gradients over multiple mini-batches
- **TorchScript for Deployment:** `torch.jit.script(model)`
- **Quantization for Model Compression:**
`torch.quantization.quantize_dynamic(model, {torch.nn.Linear}, dtype=torch.qint8)`