

hw6i7sngo

May 29, 2024

1 The goal of this project is to create a CNN that can classify images into one of 10 classes:

- airplane,
- automobile,
- bird,
- cat,
- deer,
- dog,
- frog,
- horse,
- ship, and
- truck.

2 The CIFAR-10 dataset is a popular benchmark in the field of computer vision.

3 Import Libraries

```
[1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import tensorflow as tf

from matplotlib.patches import Circle

from sklearn.metrics import precision_recall_curve
from sklearn.metrics import average_precision_score
from itertools import cycle

from sklearn.metrics import roc_curve, auc
from matplotlib import cm
import cv2
from sklearn.model_selection import train_test_split
```

```

from keras.datasets import cifar10
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from keras.utils import to_categorical
from keras.models import Sequential
from keras.layers import Dense, Conv2D, MaxPooling2D
from keras.layers import Dropout, Flatten, BatchNormalization
from keras.regularizers import l2
from keras.optimizers import Adam
from keras.callbacks import ReduceLROnPlateau, EarlyStopping
from keras.models import load_model
%matplotlib inline

```

```

2024-05-23 06:52:36.094594: E
external/local_xla/xla/stream_executor/cuda/cuda_dnn.cc:9261] Unable to register
cuDNN factory: Attempting to register factory for plugin cuDNN when one has
already been registered
2024-05-23 06:52:36.094712: E
external/local_xla/xla/stream_executor/cuda/cuda_fft.cc:607] Unable to register
cuFFT factory: Attempting to register factory for plugin cuFFT when one has
already been registered
2024-05-23 06:52:36.229763: E
external/local_xla/xla/stream_executor/cuda/cuda_blas.cc:1515] Unable to
register cuBLAS factory: Attempting to register factory for plugin cuBLAS when
one has already been registered

```

```
[2]: print("TensorFlow version:", tf.__version__)
```

```
TensorFlow version: 2.15.0
```

```
[3]: # Downloading the CIFAR-10 data
(x_train, y_train), (x_test, y_test) = cifar10.load_data()
```

```

Downloading data from https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz
170498071/170498071          4s
0us/step

```

```
[4]: # split the dataset in train and validation
x_train, x_val, y_train, y_val = train_test_split(x_train, y_train, test_size=0.
↪2, random_state=42)
```

```
[5]: # printing the shape of training, validation and test data
print("x_train shape:", x_train.shape)
print("y_train shape:", y_train.shape)
print("x_val shape:", x_val.shape)
print("y_val shape:", y_val.shape)
```

```

x_train shape: (40000, 32, 32, 3)
y_train shape: (40000, 1)

```


```
x_val shape: (10000, 32, 32, 3)
y_val shape: (10000, 1)
```

```
[6]: # CIFAR-10 classes as a list
label_name = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
```

```
[7]: unique, counts = np.unique(y_train, return_counts=True)
class_counts = dict(zip(label_name, counts))

class_counts_df = pd.DataFrame(list(class_counts.items()), columns=['Class', 'Count'])

plt.figure(figsize=(12, 6))
sns.barplot(x='Class', y='Count', data=class_counts_df, palette='hsv')
plt.title('Class Distribution in CIFAR-10 Training Dataset')
plt.xlabel('Class')
plt.ylabel('Count')
plt.xticks(rotation=45)
plt.show()
```



```
[8]: class_counts_df
```

```
[8]:
```


	Class	Count
0	airplane	4027
1	automobile	4021

2	bird	3970
3	cat	3977
4	deer	4067
5	dog	3985
6	frog	4004
7	horse	4006
8	ship	3983
9	truck	3960

```
[9]: unique, counts = np.unique(y_train, return_counts=True)
class_counts = dict(zip(label_name, counts))

class_counts_df = pd.DataFrame(list(class_counts.items()), columns=['Class',
↪ 'Count'])

plt.figure(figsize=(10, 8))
plt.pie(class_counts_df['Count'], labels=class_counts_df['Class'], autopct='%1.
↪ 1f%', startangle=140, colors=sns.color_palette('RdBu', len(label_name)))
plt.title('Class Distribution in CIFAR-10 Training Dataset')
plt.axis('equal')
plt.show()
```



```
[10]: unique, counts = np.unique(y_train, return_counts=True)
class_counts = dict(zip(label_name, counts))

class_counts_df = pd.DataFrame(list(class_counts.items()), columns=['Class',
    ↪ 'Count'])


plt.figure(figsize=(10, 8))

wedges, texts, autotexts = plt.pie(class_counts_df['Count'],
    ↪ labels=class_counts_df['Class'], autopct='%1.1f%%', startangle=140,
    ↪ colors=sns.color_palette('seismic', len(label_name)))

circle = Circle((0, 0), 0.7, color='white')
plt.gca().add_artist(circle)

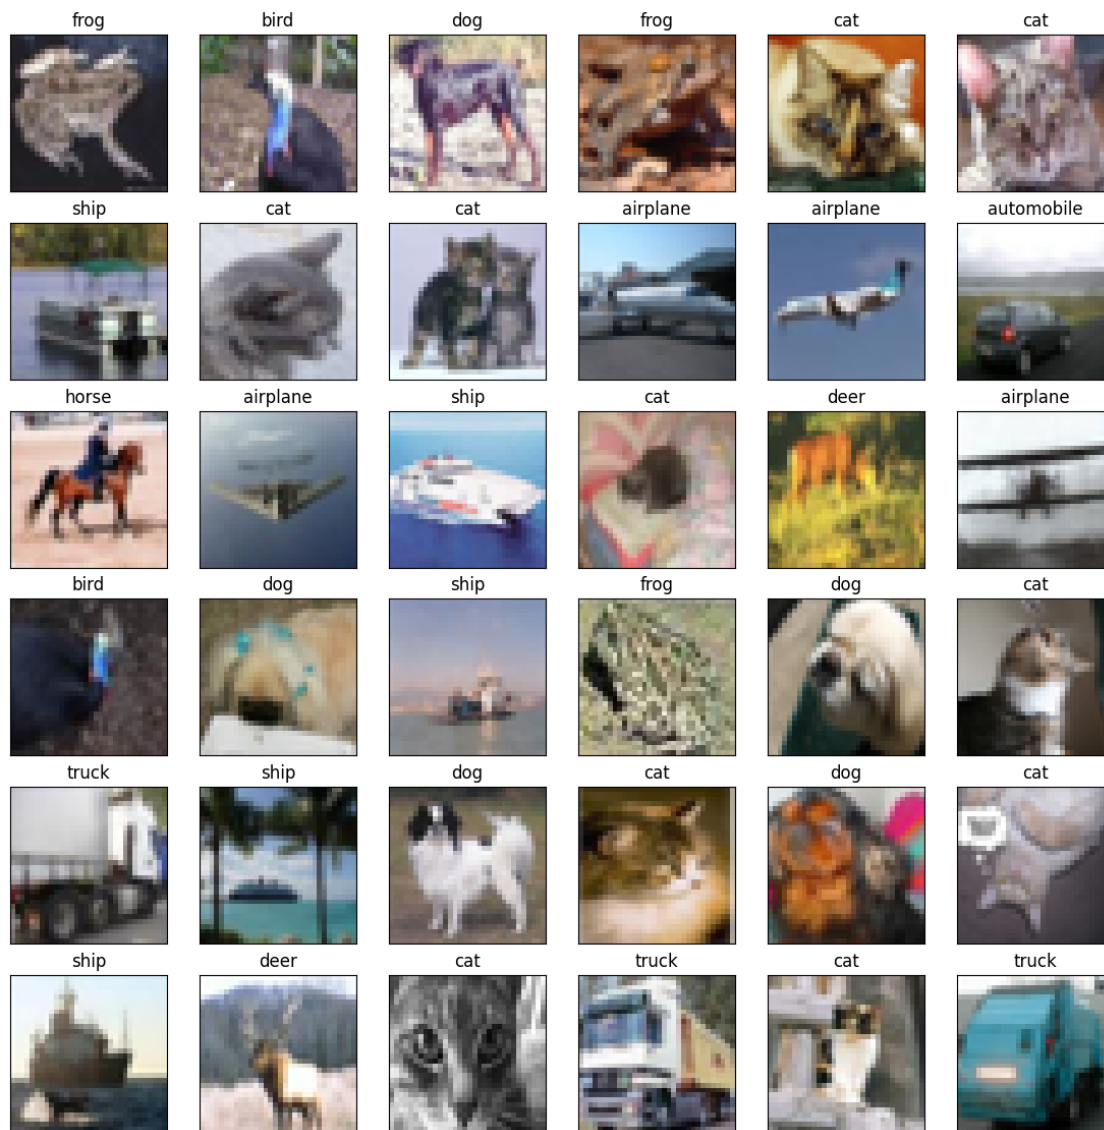
plt.axis('equal')
plt.title('Class Distribution in CIFAR-10 Training Dataset')
```

```
plt.show()
```



```
[11]: # let's see the first 25 images from our cifar-10 dataset
```

```
plt.figure(figsize=(14,14))
for i in range(36):
    plt.subplot(6,6,i+1)
    plt.xticks([])
    plt.yticks([])
    plt.grid(False)
    plt.imshow(x_train[i])
    plt.title(label_name[y_train[i][0]])
plt.show()
```



4 Data Preprocessing

- Normalization
- Encoding
- Data Augmentation

```
[12]: # pixel values to float
x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_val = x_val.astype('float32')
```

```
[13]: mean = np.mean(x_train)
      std = np.std(x_train)

      print(mean)
      print(std)
```

```
120.73008
64.074066
```

```
[14]: x_train = (x_train - mean) / (std+1e-7)
      x_test = (x_test - mean) / (std+1e-7)
      x_val = (x_val - mean) / (std+1e-7)
```

```
[15]: # converting the class labels into one-hot vectors
      y_train = to_categorical(y_train, 10)
      y_test = to_categorical(y_test, 10)
      y_val = to_categorical(y_val, 10)
```

```
[16]: datagen = ImageDataGenerator(
      rotation_range=20,
      width_shift_range=0.2,
      height_shift_range=0.2,
      horizontal_flip=True,
      zoom_range=0.3,
      brightness_range=[0.8,1.2],
      channel_shift_range=0.2
      )
```

```
[17]: x_train.shape[1:]
```

```
[17]: (32, 32, 3)
```

5 Create ZFNet Model

```
[18]: from tensorflow.keras import layers, models
      model = Sequential()

      # First layer (ZFNet used 7x7 filters)
      model.add(layers.Conv2D(96, (7, 7), strides=(2, 2), padding='same',
      ↪activation='relu', input_shape=x_train.shape[1:]))
      model.add(layers.MaxPooling2D(pool_size=(3, 3), strides=(2, 2)))

      # Second layer
      model.add(layers.Conv2D(256, (5, 5), padding='same', activation='relu'))
      model.add(layers.MaxPooling2D(pool_size=(3, 3), strides=(2, 2)))
```



```

# Third layer
model.add(layers.Conv2D(384, (3, 3), padding='same', activation='relu'))

# Fourth layer
model.add(layers.Conv2D(384, (3, 3), padding='same', activation='relu'))

# Fifth layer
model.add(layers.Conv2D(256, (3, 3), padding='same', activation='relu'))
model.add(layers.MaxPooling2D(pool_size=(3, 3), strides=(2, 2)))

# Flatten
model.add(layers.Flatten())

# Fully connected layers
model.add(layers.Dense(4096, activation='relu'))
model.add(layers.Dropout(0.5))
model.add(layers.Dense(4096, activation='relu'))
model.add(layers.Dropout(0.5))
model.add(layers.Dense(10, activation='softmax'))

optimizer = Adam(learning_rate=0.0001)
model.compile(optimizer=optimizer, loss='categorical_crossentropy',
              metrics=['accuracy'])
model.summary()

```

/opt/conda/lib/python3.10/site-packages/keras/src/layers/convolutional/base_conv.py:99: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 16, 16, 96)	14,208
max_pooling2d (MaxPooling2D)	(None, 7, 7, 96)	0
conv2d_1 (Conv2D)	(None, 7, 7, 256)	614,656
max_pooling2d_1 (MaxPooling2D)	(None, 3, 3, 256)	0
conv2d_2 (Conv2D)	(None, 3, 3, 384)	885,120

conv2d_3 (Conv2D)	(None, 3, 3, 384)	1,327,488
conv2d_4 (Conv2D)	(None, 3, 3, 256)	884,992
max_pooling2d_2 (MaxPooling2D)	(None, 1, 1, 256)	0
flatten (Flatten)	(None, 256)	0
dense (Dense)	(None, 4096)	1,052,672
dropout (Dropout)	(None, 4096)	0
dense_1 (Dense)	(None, 4096)	16,781,312
dropout_1 (Dropout)	(None, 4096)	0
dense_2 (Dense)	(None, 10)	40,970

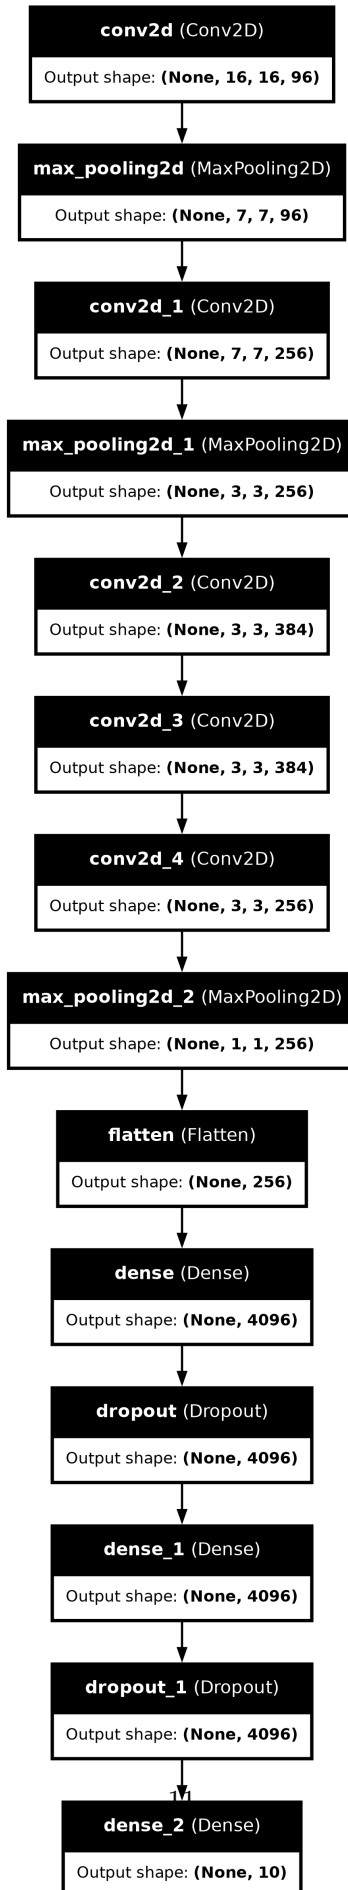
Total params: 21,601,418 (82.40 MB)

Trainable params: 21,601,418 (82.40 MB)

Non-trainable params: 0 (0.00 B)

```
[41]: from tensorflow.keras.utils import plot_model
plot_model(model, to_file='zfnet_cifar10.png', show_shapes=True,
          show_layer_names=True)
```

[41]:



```
[19]: batch_size = 64
epochs = 100

reduce_lr = ReduceLROnPlateau(monitor = 'val_loss', factor = 0.5, patience = 5,
    ↪ min_lr = 0.00001)
early_stop = EarlyStopping(monitor = 'val_loss', min_delta = 0, patience = 10,
    ↪ verbose = 1, mode = 'auto')

model.fit(x_train, y_train, batch_size = batch_size,
          epochs = epochs,
          validation_data = (x_val, y_val),
          callbacks = [reduce_lr, early_stop],
          verbose = 1)
```

Epoch 1/100

15/625 7s 11ms/step - accuracy:
0.1034 - loss: 2.3012

WARNING: All log messages before absl::InitializeLog() is called are written to STDERR

I0000 00:00:1716447187.423526 157 device_compiler.h:186] Compiled cluster
using XLA! This line is logged at most once for the lifetime of the process.
W0000 00:00:1716447187.441835 157 graph_launch.cc:671] Fallback to op-by-op
mode because memset node breaks graph update

625/625 0s 11ms/step -
accuracy: 0.2897 - loss: 1.8732

W0000 00:00:1716447195.136756 159 graph_launch.cc:671] Fallback to op-by-op
mode because memset node breaks graph update

625/625 19s 15ms/step -
accuracy: 0.2898 - loss: 1.8728 - val_accuracy: 0.5206 - val_loss: 1.3152 -
learning_rate: 1.0000e-04

Epoch 2/100

625/625 7s 11ms/step -
accuracy: 0.5393 - loss: 1.2634 - val_accuracy: 0.5874 - val_loss: 1.1365 -
learning_rate: 1.0000e-04

Epoch 3/100

625/625 7s 11ms/step -
accuracy: 0.6247 - loss: 1.0468 - val_accuracy: 0.6451 - val_loss: 1.0119 -
learning_rate: 1.0000e-04

Epoch 4/100

625/625 7s 11ms/step -
accuracy: 0.6863 - loss: 0.8828 - val_accuracy: 0.6648 - val_loss: 0.9211 -

```

learning_rate: 1.0000e-04
Epoch 5/100
625/625          7s 11ms/step -
accuracy: 0.7332 - loss: 0.7462 - val_accuracy: 0.6980 - val_loss: 0.8668 -
learning_rate: 1.0000e-04
Epoch 6/100
625/625          7s 11ms/step -
accuracy: 0.7854 - loss: 0.6195 - val_accuracy: 0.7153 - val_loss: 0.8450 -
learning_rate: 1.0000e-04
Epoch 7/100
625/625          7s 11ms/step -
accuracy: 0.8275 - loss: 0.4958 - val_accuracy: 0.7267 - val_loss: 0.8215 -
learning_rate: 1.0000e-04
Epoch 8/100
625/625          7s 11ms/step -
accuracy: 0.8670 - loss: 0.3843 - val_accuracy: 0.7240 - val_loss: 0.8707 -
learning_rate: 1.0000e-04
Epoch 9/100
625/625          7s 11ms/step -
accuracy: 0.8966 - loss: 0.2976 - val_accuracy: 0.7325 - val_loss: 0.8983 -
learning_rate: 1.0000e-04
Epoch 10/100
625/625          7s 11ms/step -
accuracy: 0.9210 - loss: 0.2271 - val_accuracy: 0.7337 - val_loss: 0.9643 -
learning_rate: 1.0000e-04
Epoch 11/100
625/625          7s 11ms/step -
accuracy: 0.9424 - loss: 0.1645 - val_accuracy: 0.7209 - val_loss: 1.1078 -
learning_rate: 1.0000e-04
Epoch 12/100
625/625          7s 11ms/step -
accuracy: 0.9529 - loss: 0.1378 - val_accuracy: 0.7264 - val_loss: 1.1221 -
learning_rate: 1.0000e-04
Epoch 13/100
625/625          7s 11ms/step -
accuracy: 0.9878 - loss: 0.0437 - val_accuracy: 0.7472 - val_loss: 1.2859 -
learning_rate: 5.0000e-05
Epoch 14/100
625/625          7s 11ms/step -
accuracy: 0.9966 - loss: 0.0129 - val_accuracy: 0.7426 - val_loss: 1.4302 -
learning_rate: 5.0000e-05
Epoch 15/100
625/625          7s 11ms/step -
accuracy: 0.9912 - loss: 0.0283 - val_accuracy: 0.7360 - val_loss: 1.4540 -
learning_rate: 5.0000e-05
Epoch 16/100
625/625          10s 11ms/step -
accuracy: 0.9919 - loss: 0.0247 - val_accuracy: 0.7491 - val_loss: 1.4797 -

```

```

learning_rate: 5.0000e-05
Epoch 17/100
625/625          7s 11ms/step -
accuracy: 0.9897 - loss: 0.0327 - val_accuracy: 0.7427 - val_loss: 1.5428 -
learning_rate: 5.0000e-05
Epoch 17: early stopping

```

```
[19]: <keras.src.callbacks.history.History at 0x7d58003a63e0>
```

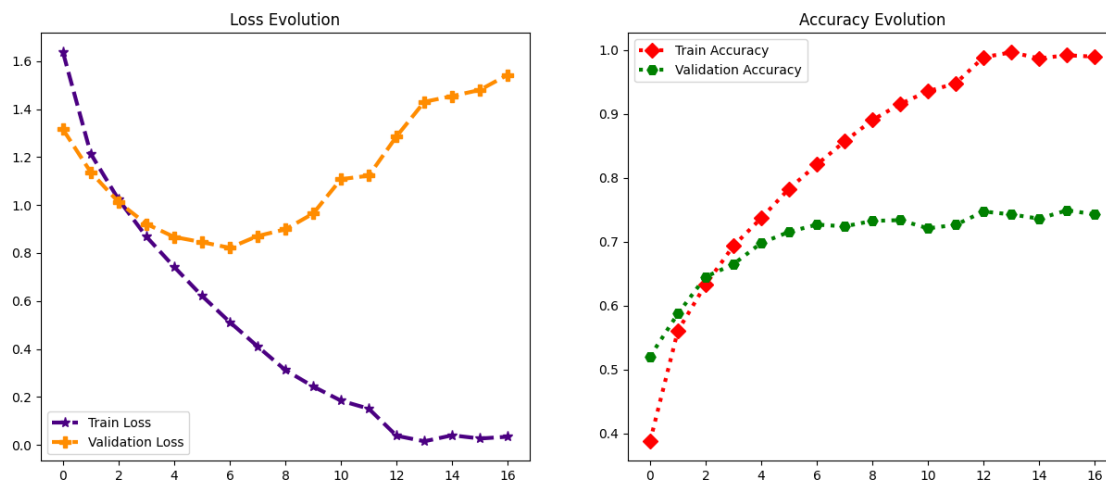
```
[20]: model.save('trained_model_10.h5')
```

```
[21]: plt.figure(figsize=(15,6))

# Plotting the training and validation loss
plt.subplot(1, 2, 1)
plt.plot(model.history.history['loss'], label='Train Loss',
         color='indigo',marker="*",lw=3,markersize=8,linestyle="--")
plt.plot(model.history.history['val_loss'], label='Validation Loss',
         color='darkorange',marker="P",lw=3,markersize=8,linestyle="--")
plt.legend()
plt.title('Loss Evolution')

# Plotting the training and validation accuracy
plt.subplot(1, 2, 2)
plt.plot(model.history.history['accuracy'], label='Train Accuracy',
         color="red",marker="D",lw=3,markersize=8,linestyle=":")
plt.plot(model.history.history['val_accuracy'], label='Validation Accuracy',
         color='green',marker="H",lw=3,markersize=8,linestyle=":")
plt.legend()
plt.title('Accuracy Evolution')

plt.show()
```



```
[22]: test_loss, test_acc = model.evaluate(x_test, y_test, verbose=1)
```

```
print('\nTest Accuracy:', test_acc)
print('Test Loss:      ', test_loss)
```

```
313/313          2s 3ms/step -
accuracy: 0.7451 - loss: 1.5669
```

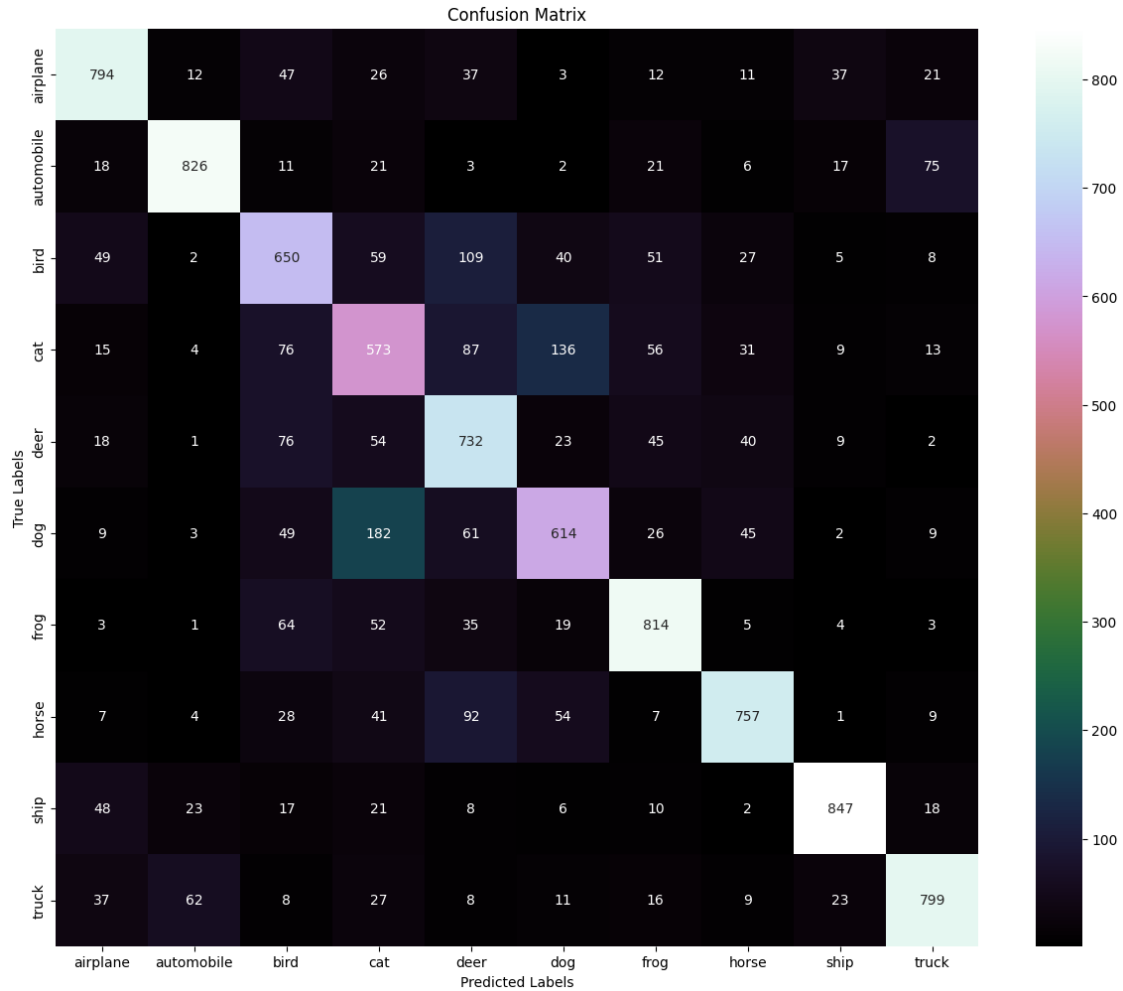
```
Test Accuracy: 0.7405999898910522
Test Loss:      1.611112356185913
```

```
[23]: pred=model.predict(x_test)
final_pred=np.argmax(pred,axis=1)
y_test_labels = np.argmax(y_test, axis=1)
```

```
313/313          2s 3ms/step
```

6 confusion matrix

```
[24]: from sklearn.metrics import confusion_matrix, roc_curve, roc_auc_score, \
      ↪classification_report, log_loss, accuracy_score
import seaborn as sns
cm = confusion_matrix(y_test_labels, final_pred)
plt.figure(figsize=(15,12))
sns.
      ↪heatmap(cm,annot=True,fmt="d",xticklabels=label_name,yticklabels=label_name,cmap="cubehelix
plt.title('Confusion Matrix')
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.show()
```



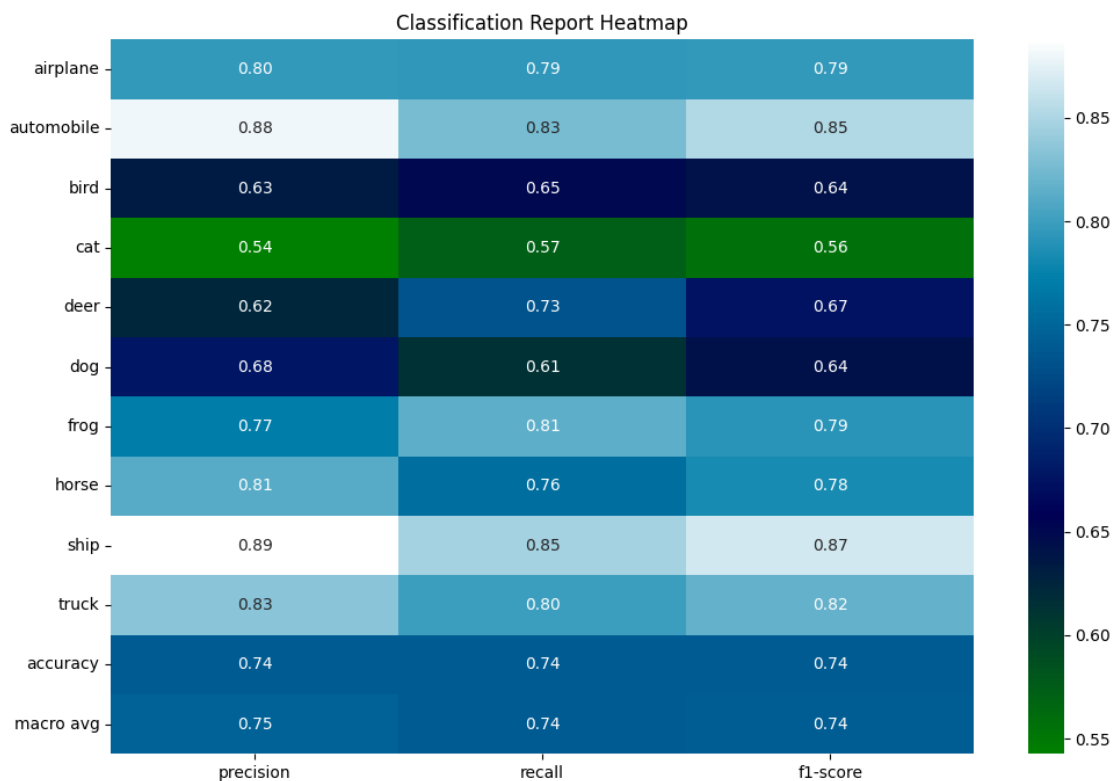
```
[25]: print(classification_report(y_test_labels,final_pred,target_names=label_name))
```

	precision	recall	f1-score	support
airplane	0.80	0.79	0.79	1000
automobile	0.88	0.83	0.85	1000
bird	0.63	0.65	0.64	1000
cat	0.54	0.57	0.56	1000
deer	0.62	0.73	0.67	1000
dog	0.68	0.61	0.64	1000
frog	0.77	0.81	0.79	1000
horse	0.81	0.76	0.78	1000
ship	0.89	0.85	0.87	1000
truck	0.83	0.80	0.82	1000
accuracy			0.74	10000
macro avg	0.75	0.74	0.74	10000

weighted avg 0.75 0.74 0.74 10000

7 Classification report

```
[26]: report = classification_report(y_test_labels, final_pred,
    ↪target_names=label_name, output_dict=True)
report_df = pd.DataFrame(report).transpose()
plt.figure(figsize=(12, 8))
sns.heatmap(report_df.iloc[:-1, :-1], annot=True, cmap='ocean', fmt='.2f')
plt.title('Classification Report Heatmap')
plt.show()
```



8 Accuracy Score

```
[27]: plt.figure(figsize=(10,5),facecolor="white")
acc_score = accuracy_score(y_test_labels,final_pred)
plt.plot([])
plt.text(0,0, f'Accuracy Score Score: {acc_score:.4f}', fontsize=20,
    ↪ha='center', va='center',color="crimson")
```

```
plt.axis('off')

# Set the x-axis limits
plt.xlim(-1, 1)
plt.ylim(-1,1)

plt.show()
```

Accuracy Score Score: 0.7406

9 ROC AUC Score

```
[28]: plt.figure(figsize=(10,5),facecolor='white')
roc_auc = roc_auc_score(y_test_labels, pred,multi_class='ovr')
plt.plot([])
plt.text(0,0, f'ROC AUC Score: {roc_auc:.4f}', fontsize=20, ha='center',
↪va='center',color="crimson")
plt.axis('off')

# Set the x-axis limits
plt.xlim(-1, 1)
plt.ylim(-1,1)

plt.show()
```

ROC AUC Score: 0.9579

10 Cohen Kappa Score

```
[29]: from sklearn.metrics import cohen_kappa_score
plt.figure(figsize=(10,5),facecolor="white")
kappa = cohen_kappa_score(y_test_labels,final_pred)
plt.plot([])
plt.text(0,0, f'Cohen Kappa Score: {kappa:.4f}', fontsize=20, ha='center',
        ↪va='center',color="crimson")
plt.axis('off')

# Set the x-axis limits
plt.xlim(-1, 1)
plt.ylim(-1,1)

plt.show()
```

Cohen Kappa Score: 0.7118

11 Matthews Correlation Coefficient

```
[30]: from sklearn.metrics import matthews_corrcoef
plt.figure(figsize=(10,5),facecolor="white")
mcc = matthews_corrcoef(y_test_labels, final_pred)

# Create a plot and display the MCC value as text
plt.plot([])
plt.text(0,0, f'Matthews Correlation Coefficient: {mcc:.4f}', fontsize=20,
        ha='center', va='center',color="crimson")
plt.axis('off')

# Set the x-axis limits
plt.xlim(-1, 1)
plt.ylim(-1,1)

plt.show()
```

Matthews Correlation Coefficient: 0.7120

12 Brier Score Loss __1

```
[31]: import pandas as pd
      from sklearn.metrics import brier_score_loss

      pred = model.predict(x_test)
      final_pred = np.argmax(pred, axis=1)
      y_test_labels = np.argmax(y_test, axis=1)


      num_classes = len(label_name)  # Number of classes

      brier_scores = []
      for class_label in range(num_classes):
          class_mask = (y_test_labels == class_label)
          class_pred = (final_pred == class_label)
          brier_score = brier_score_loss(class_mask, class_pred)
          brier_scores.append(brier_score)

      df = pd.DataFrame({'Class': label_name, 'Brier Score Loss': brier_scores})
      plt.figure(figsize=(10, 5), facecolor="white")
      plt.bar(df['Class'], df['Brier Score Loss'], color='navy')
      plt.xlabel('Class')
      plt.ylabel('Brier Score Loss')
      plt.title('Brier Score Loss for Each Class')
      plt.xticks(rotation=45)
      plt.show()
```

313/313

1s 2ms/step



```
[32]: df
```

```
[32]:
```

	Class	Brier Score Loss
0	airplane	0.0410
1	automobile	0.0286
2	bird	0.0726
3	cat	0.0910
4	deer	0.0708
5	dog	0.0680
6	frog	0.0430
7	horse	0.0419
8	ship	0.0260
9	truck	0.0359

13 Brier Score Loss

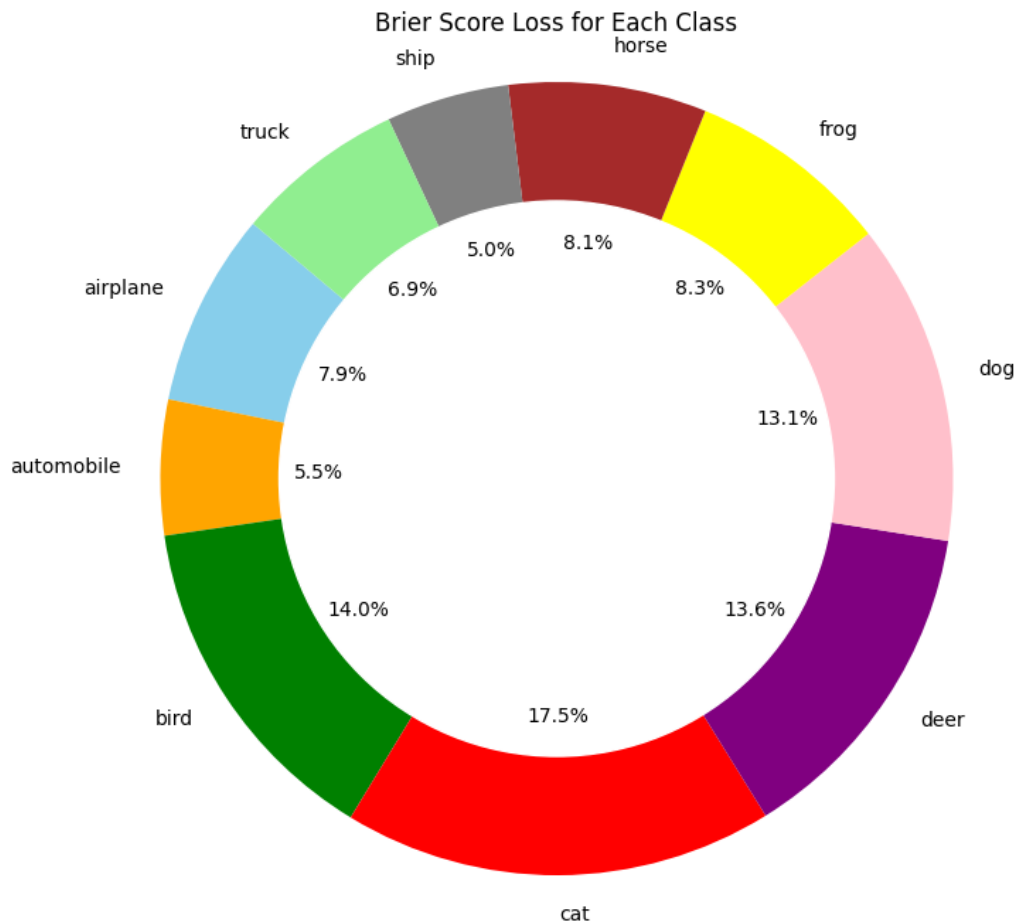
```
[33]: colors = ['skyblue', 'orange', 'green', 'red', 'purple', 'pink', 'yellow', 'brown', 'grey', 'lightgreen']

plt.figure(figsize=(10, 8))
ax = plt.gca()
```

```
wedges, texts, autotexts = ax.pie(df['Brier Score Loss'], labels=df['Class'],
    autopct='%1.1f%%', startangle=140, colors=colors)

circle = Circle((0, 0), 0.7, color='white')
ax.add_artist(circle)

ax.axis('equal')
plt.title('Brier Score Loss for Each Class')
plt.show()
```



14 Precision, Recall Score

```
[34]: from sklearn.metrics import precision_score, recall_score
```

```
precisions = []
```

```

recalls = []

for class_label in range(num_classes):
    class_mask = (y_test_labels == class_label)
    class_pred = (final_pred == class_label)

    precision = precision_score(class_mask, class_pred)
    recall = recall_score(class_mask, class_pred)


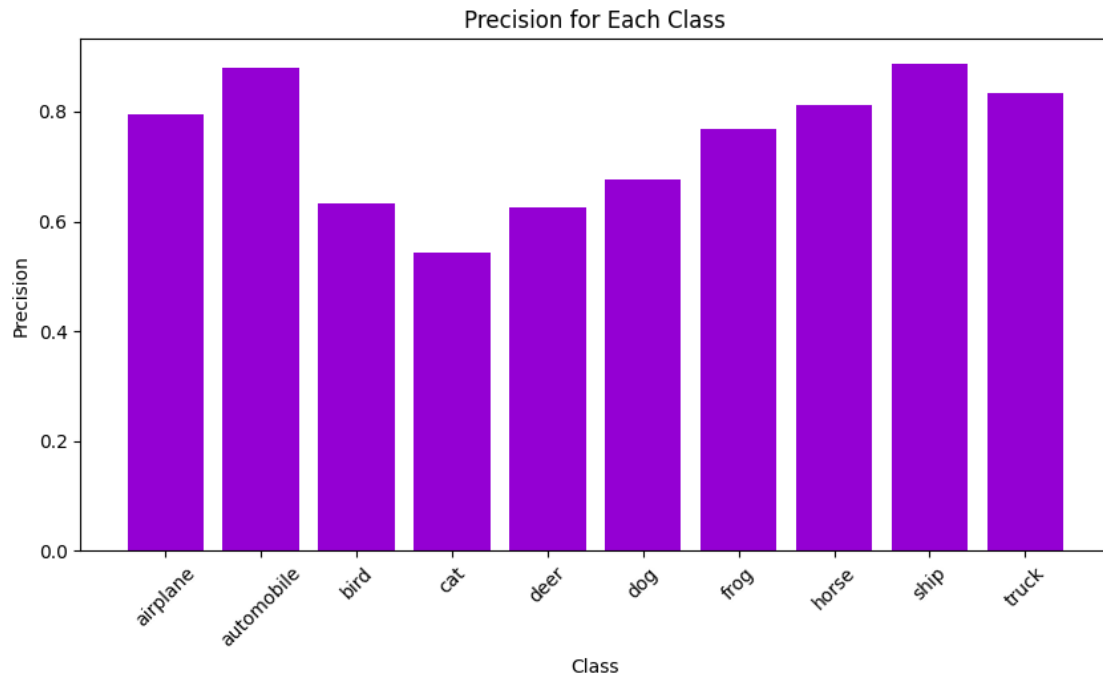
    precisions.append(precision)
    recalls.append(recall)

df = pd.DataFrame({'Class': label_name, 'Precision': precisions, 'Recall':
    ↪recalls})

# Plotting precision
plt.figure(figsize=(10, 5), facecolor="white")
plt.bar(df['Class'], df['Precision'], color='darkviolet')
plt.xlabel('Class')
plt.ylabel('Precision')
plt.title('Precision for Each Class')
plt.xticks(rotation=45)
plt.show()

# Plotting recall
plt.figure(figsize=(10, 5), facecolor="white")
plt.bar(df['Class'], df['Recall'], color='crimson')
plt.xlabel('Class')
plt.ylabel('Recall')
plt.title('Recall for Each Class')
plt.xticks(rotation=45)
plt.show()

```

```
[35]: colors = ['brown', 'coral', 'orchid', 'olive', 'yellowgreen', 'lime', 'yellow', 'teal', 'grey', 'purple']
```

```


plt.figure(figsize=(10, 8))
ax = plt.gca()

wedges, texts, autotexts = ax.pie(df['Recall'], labels=df['Class'], autopct='%1.
    ↳1f%%', startangle=140, colors=colors)

circle = Circle((0, 0), 0.7, color='white')
ax.add_artist(circle)

ax.axis('equal')
plt.title('Recall Score for Each Class')
plt.show()

```



```

[36]: colors = ['gold', 'maroon', 'peru', 'dodgerblue', 'powderblue', 'pink',
    ↳ 'yellow', 'brown', 'bisque', 'aqua']

plt.figure(figsize=(10, 8))

```

```


ax = plt.gca()

wedges, texts, autotexts = ax.pie(df['Precision'], labels=df['Class'],
    autopct='%1.1f%%', startangle=140, colors=colors)

circle = Circle((0, 0), 0.7, color='white')
ax.add_artist(circle)

ax.axis('equal')
plt.title('Precision Score for Each Class')
plt.show()

```



```

[37]: import matplotlib as mpl
import matplotlib.cm as cm

fpr = dict()
tpr = dict()
roc_auc = dict()

```

```

num_classes = len(label_name)

for class_index in range(num_classes):
    fpr[class_index], tpr[class_index], _ = roc_curve((y_test_labels ==
    ↪class_index).astype(int), pred[:, class_index])
    roc_auc[class_index] = auc(fpr[class_index], tpr[class_index])

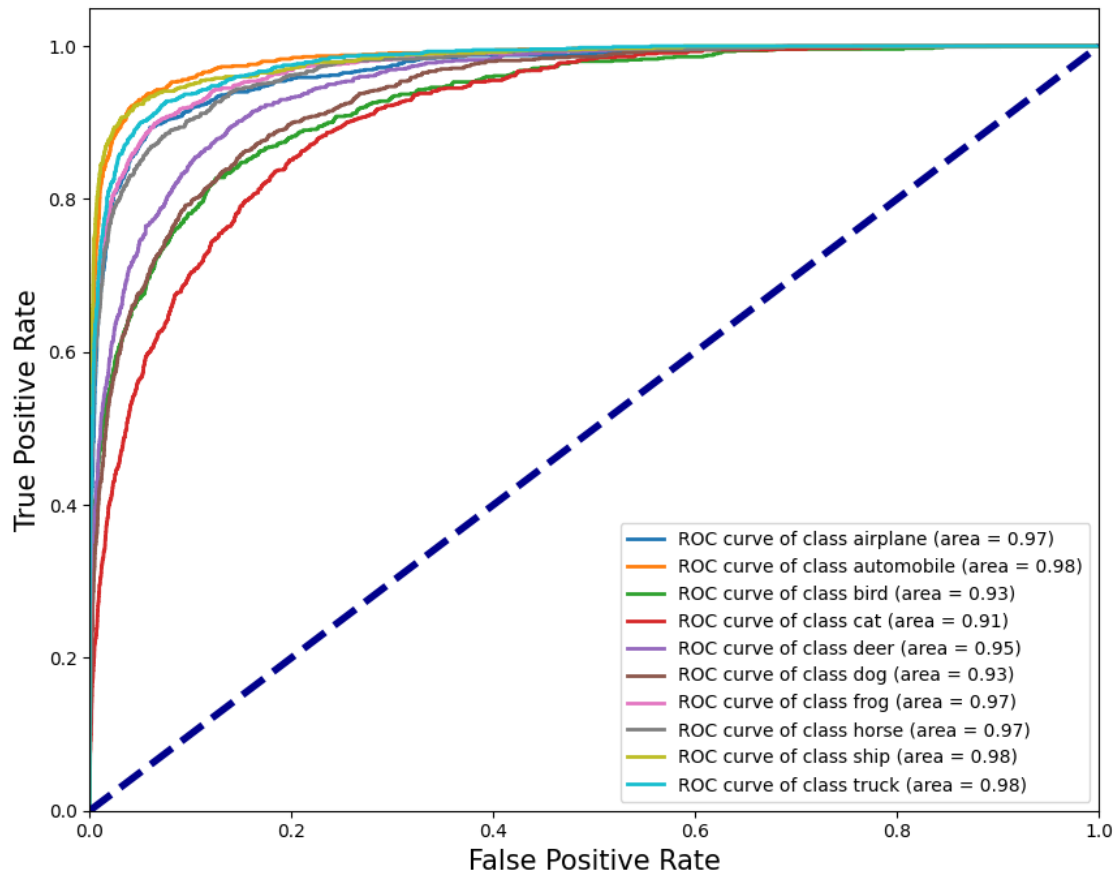
# Plot the ROC curves
plt.figure(figsize=(10, 8))
colors = cm.tab20(np.arange(num_classes) / num_classes)

for class_index in range(num_classes):
    plt.plot(fpr[class_index], tpr[class_index], color=colors[class_index],
    ↪lw=2, label='ROC curve of class {} (area = {:.2f})'.
    ↪format(label_name[class_index], roc_auc[class_index]))

plt.plot([0, 1], [0, 1], linestyle="--", lw=4, color="darkblue")
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate', fontsize=15, color="black")
plt.ylabel('True Positive Rate', fontsize=15, color="black")
plt.title('Receiver Operating Characteristic\n', fontsize=20, color="black")
plt.legend(loc="lower right")
plt.show()

```

Receiver Operating Characteristic

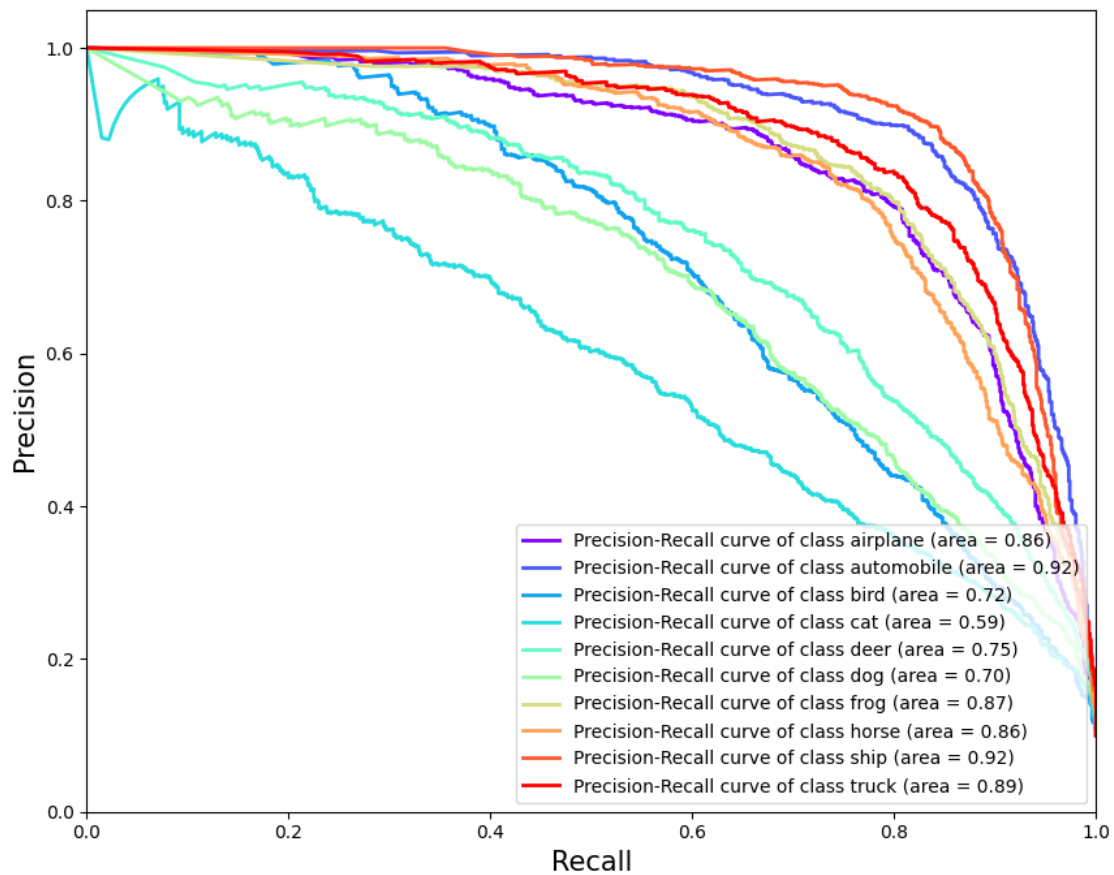


```
[38]: num_classes = len(label_name)
precision = dict()
recall = dict()
average_precision = dict()
for i in range(num_classes):
    precision[i], recall[i], _ = precision_recall_curve(y_test_labels == i,
    ↪pred[:, i])
    average_precision[i] = average_precision_score(y_test_labels == i, pred[:,
    ↪i])

plt.figure(figsize=(10, 8))
colors = cycle(cm.rainbow(np.linspace(0, 1, num_classes)))
for i, color in zip(range(num_classes), colors):
    plt.plot(recall[i], precision[i], color=color, lw=2,
    label='Precision-Recall curve of class {0} (area = {1:0.2f})'.
    ↪format(label_name[i], average_precision[i]))
```

```
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('Recall', fontsize=15, color="black")
plt.ylabel('Precision', fontsize=15, color="black")
plt.title('Precision-Recall curve\n', fontsize=20, color="black")
plt.legend(loc="lower right")
plt.show()
```

Precision-Recall curve



15 Prediction

```
[39]: plt.figure(figsize=(30, 30))
number_images = (5, 5)
num_samples = min(number_images[0] * number_images[1], len(x_test))
```

```


for i in range(num_samples):
    plt.subplot(number_images[0], number_images[1], i + 1)
    plt.axis("off")

    true_label = label_name[y_test_labels[i]]
    predicted_label = label_name[final_pred[i]]

    color = "darkgreen"
    if true_label != predicted_label:
        color = "red"

    plt.title(f"True: {true_label}\nPredicted: {predicted_label}", color=color)
    plt.imshow(x_test[i]) # Assuming x_test[i] is already an image array
plt.show()

```



16 Custom Data Prediction

```
[40]: plt.figure(figsize=(15, 15))
import urllib.request
import os
from PIL import Image
img_url = "https://t3.ftcdn.net/jpg/00/41/06/42/
↪360_F_41064239_IaGdGyf1vxHFaNDS5K1640F0wiMe1hC9.jpg"

# Retrieve the image from the URL
filename, headers = urllib.request.urlretrieve(img_url)
img_path = os.path.join(os.getcwd(), filename)
img = Image.open(img_path)
img = img.resize((32,32))
img = np.array(img) / 255.0
img = np.expand_dims(img, axis=0)

# Predict the class of the image using the model
probs = model.predict(img)[0]

# Get the predicted class index and name
pred_class_prob = np.argmax(probs)
pred_class_name = label_name[pred_class_prob]

max_prob = np.max(probs)
print(f'Predicted class: {pred_class_name}')
print(f'Maximum probability: {max_prob}')

# Display the image with the predicted class and probability
plt.imshow(img[0])
plt.axis('off')
plt.text(5, 15, f'Predicted class: {pred_class_name}\nMaximum probability:
↪{max_prob:.2f}', fontsize=20, color='red', bbox=dict(facecolor='white',
↪alpha=0.8))
plt.show()
```

```
1/1          1s 826ms/step
Predicted class: ship
Maximum probability: 1.0
```




[]: